

Cluster Analysis

Data Sets: Wholesale customers data and USArrests

Hierarchical Clustering

What is Hierarchical Clustering?

Let's say we have the below points and we want to cluster them into groups:

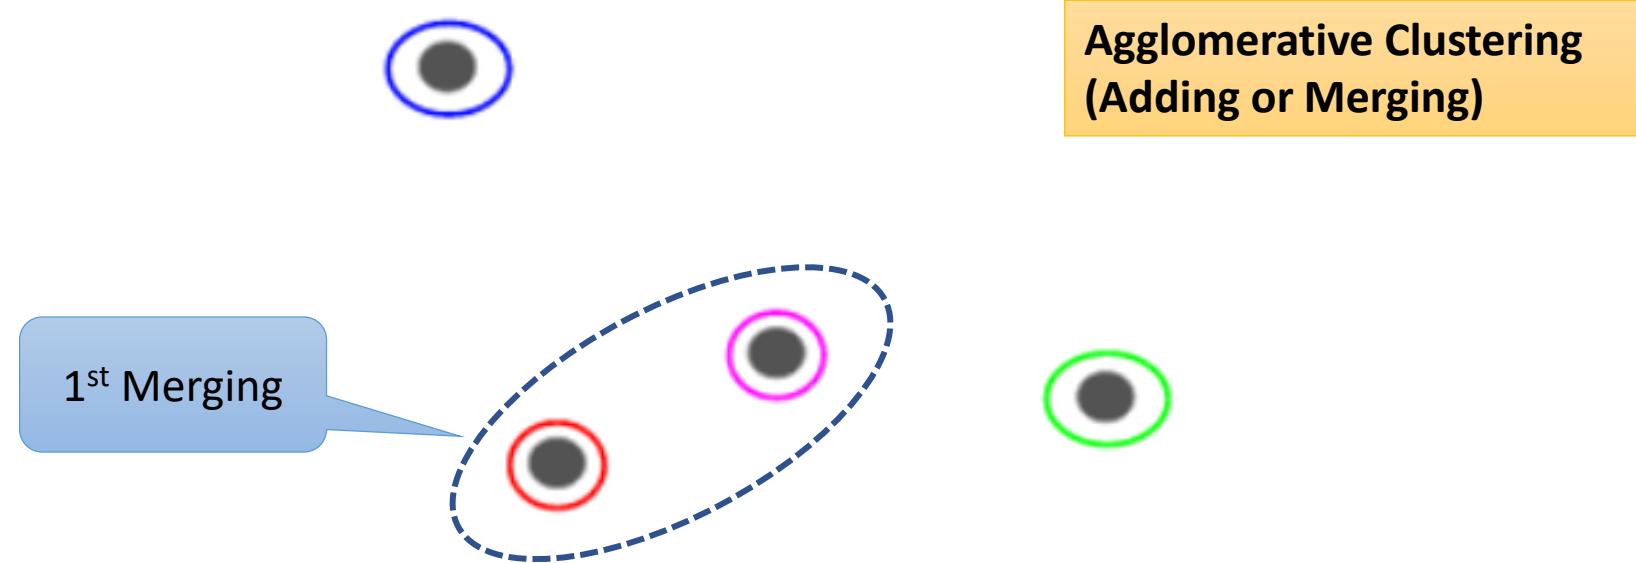


Each data point as one cluster

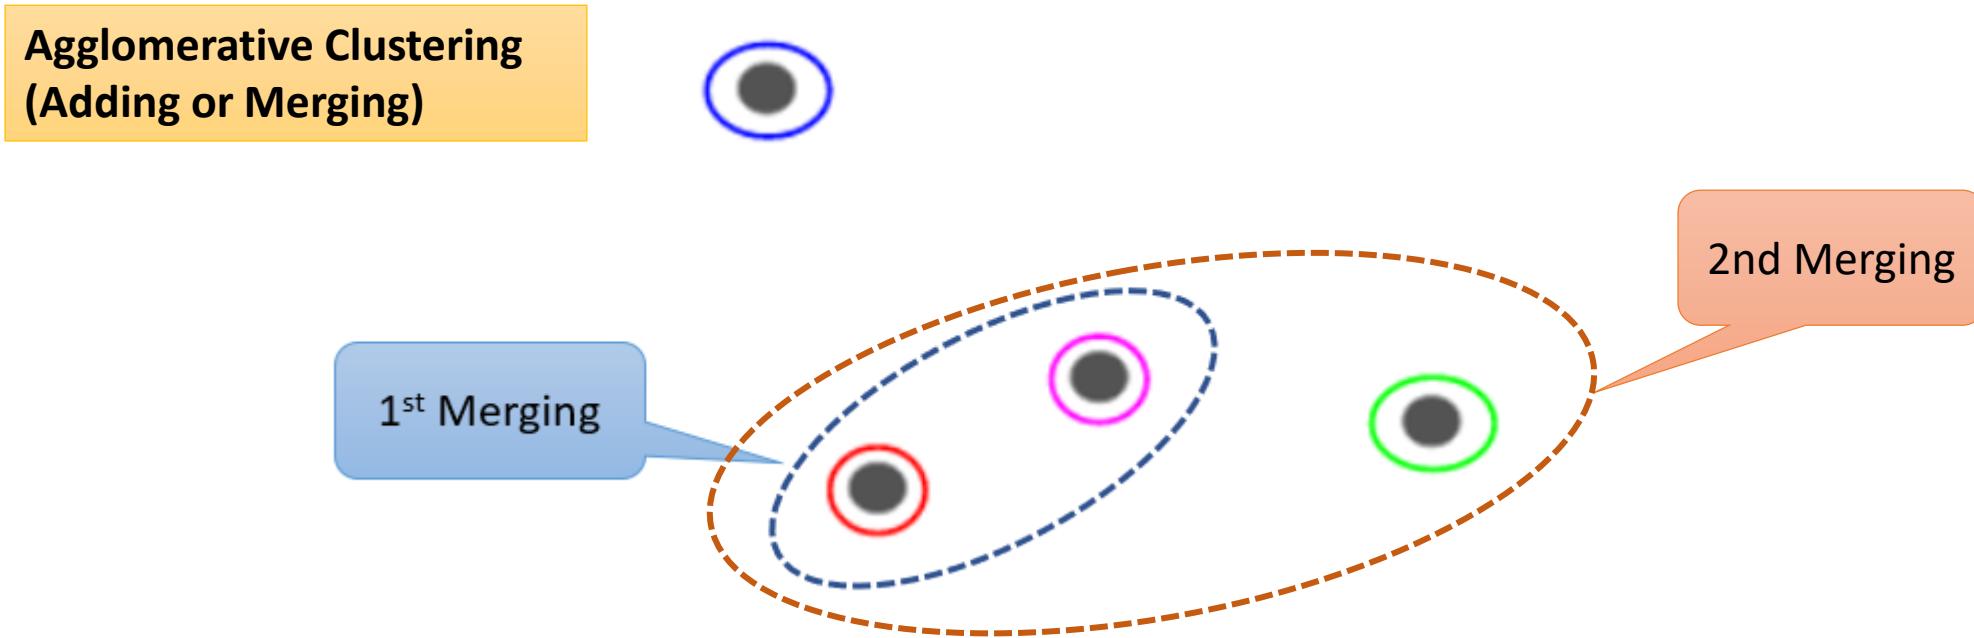
We can assign each of these points to a separate cluster:



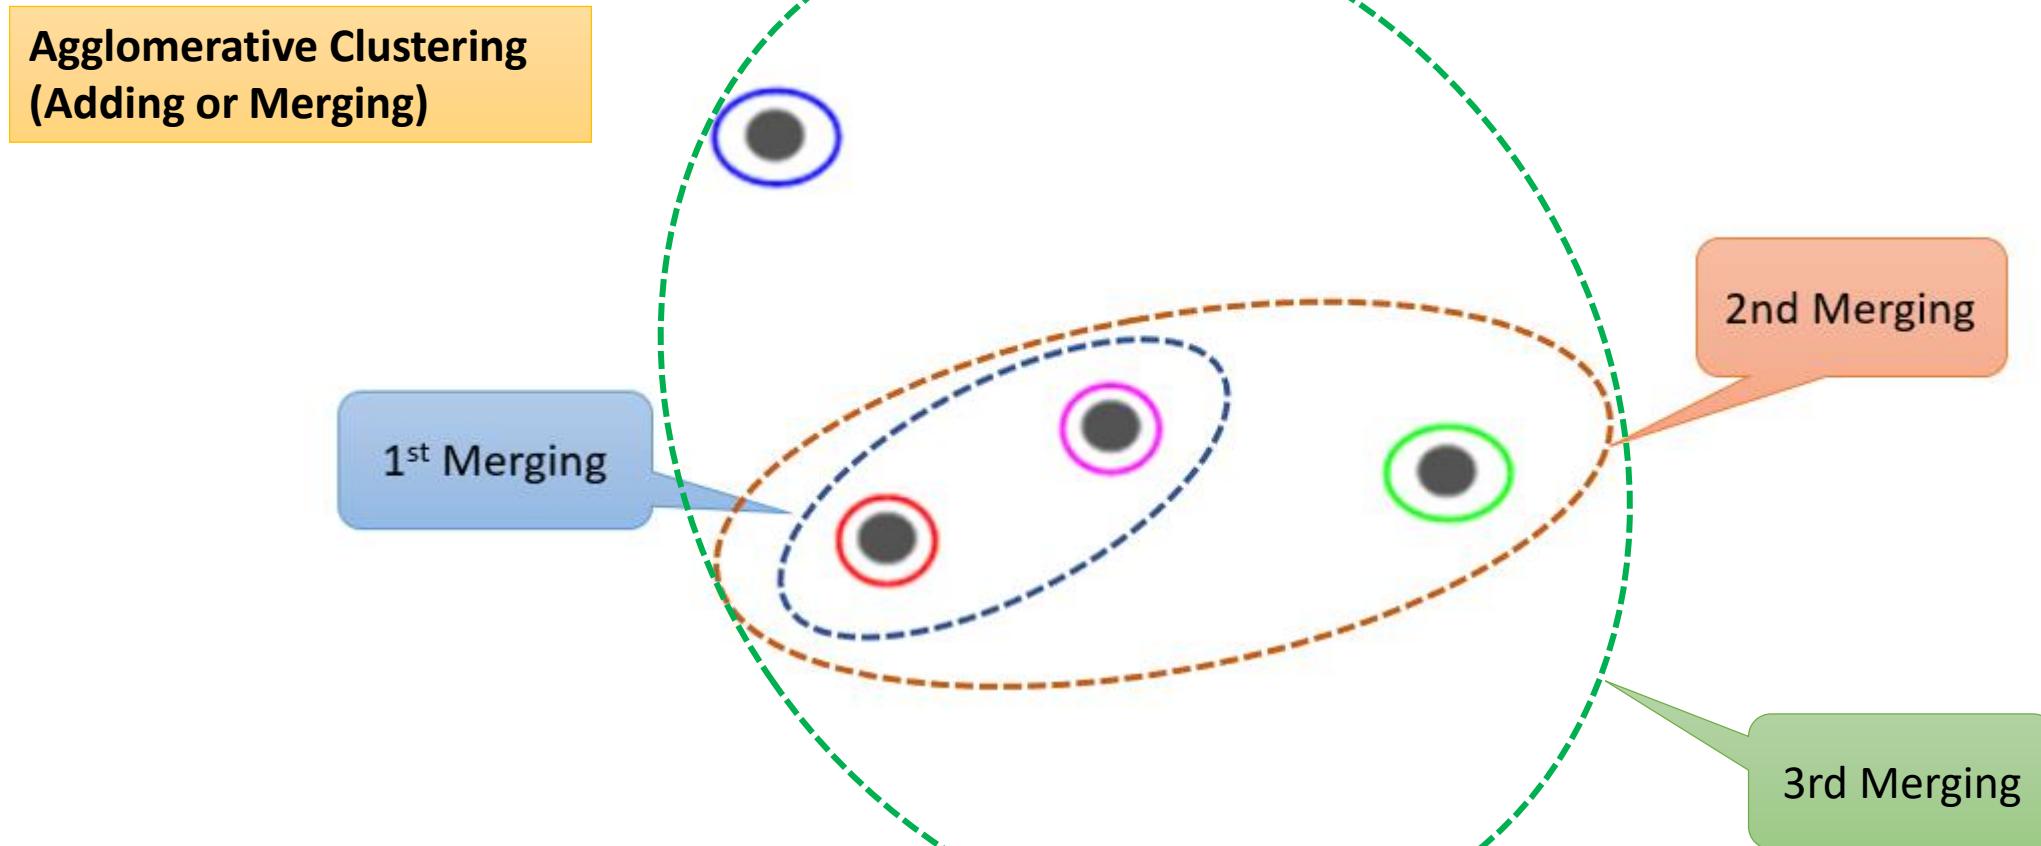
1st merging of two nearest clusters/(data points)



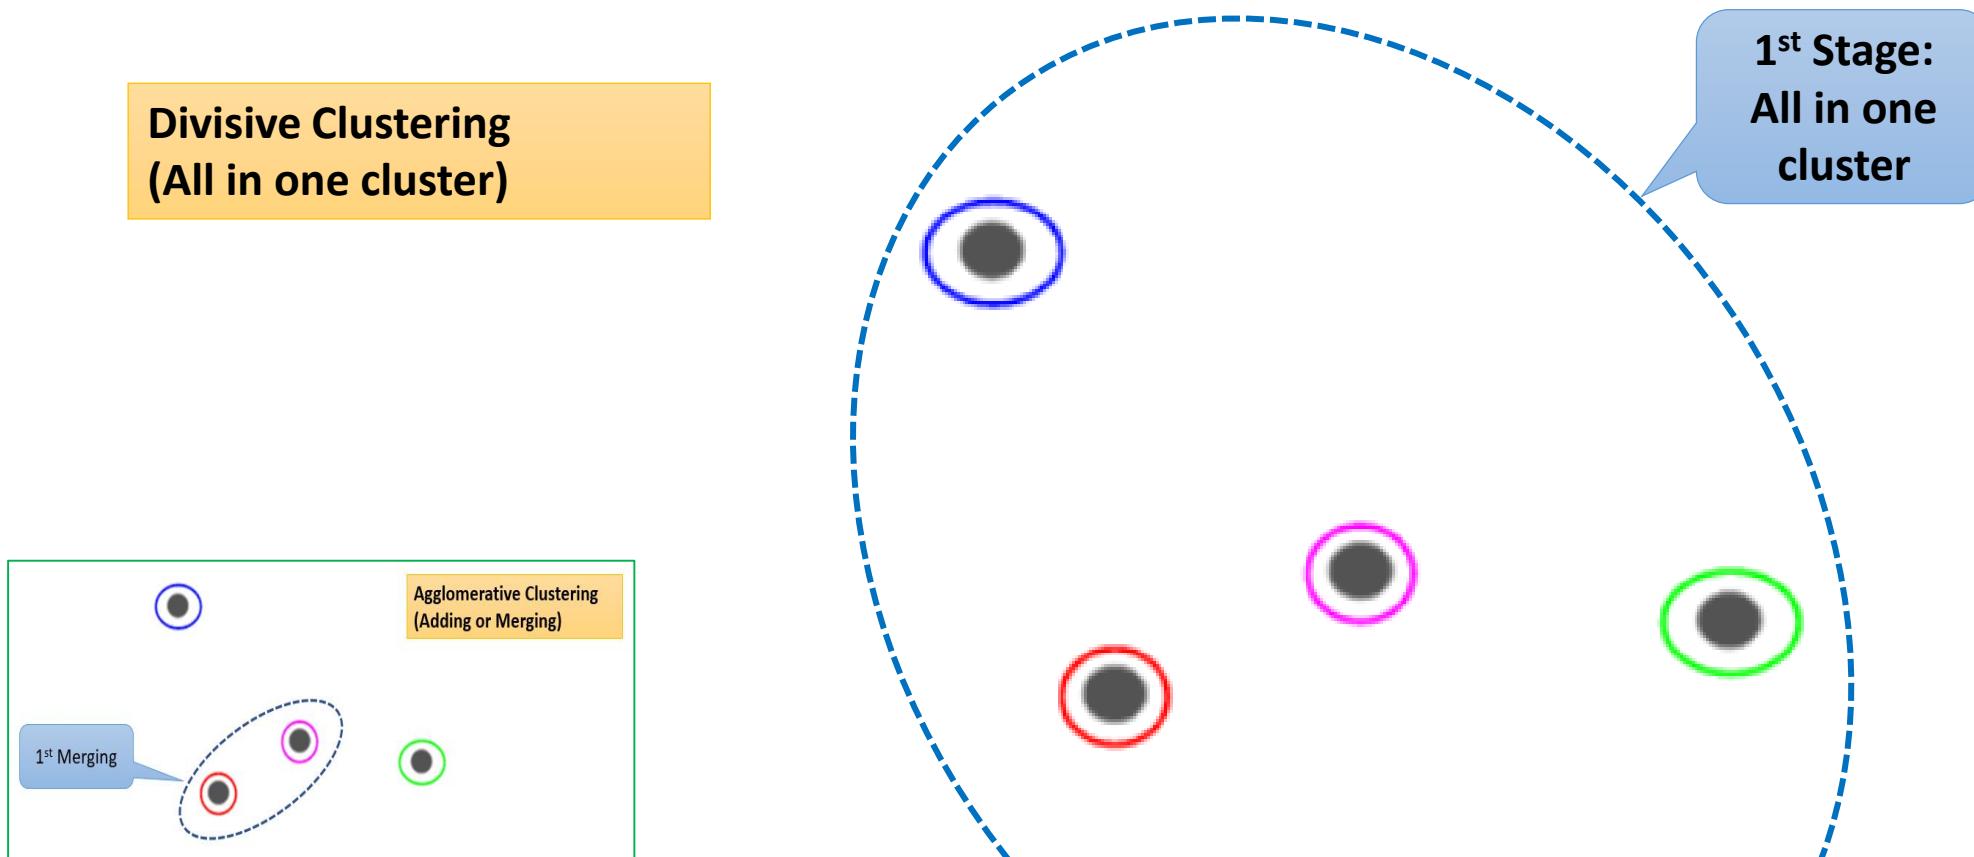
2nd merging of two nearest clusters/(data points)

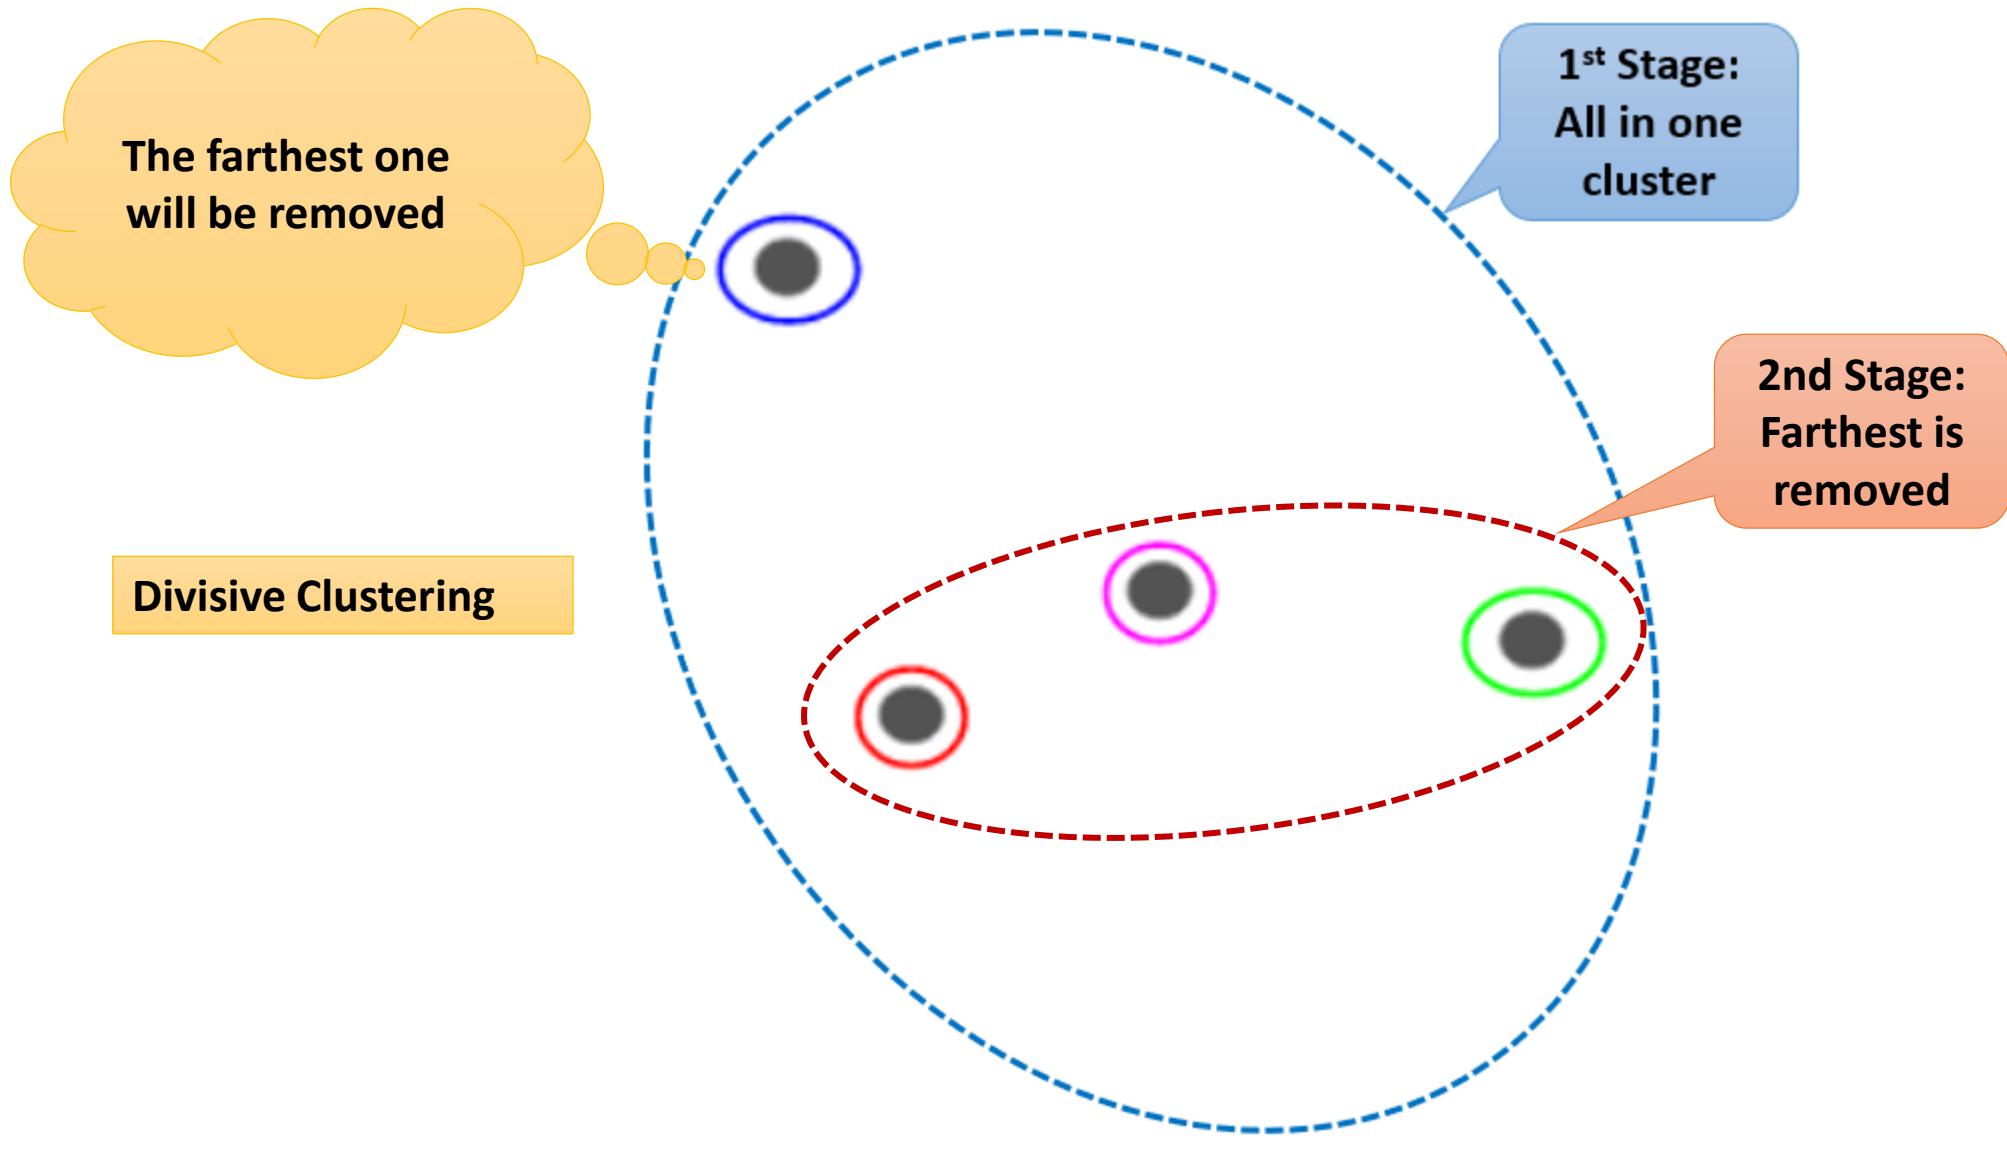


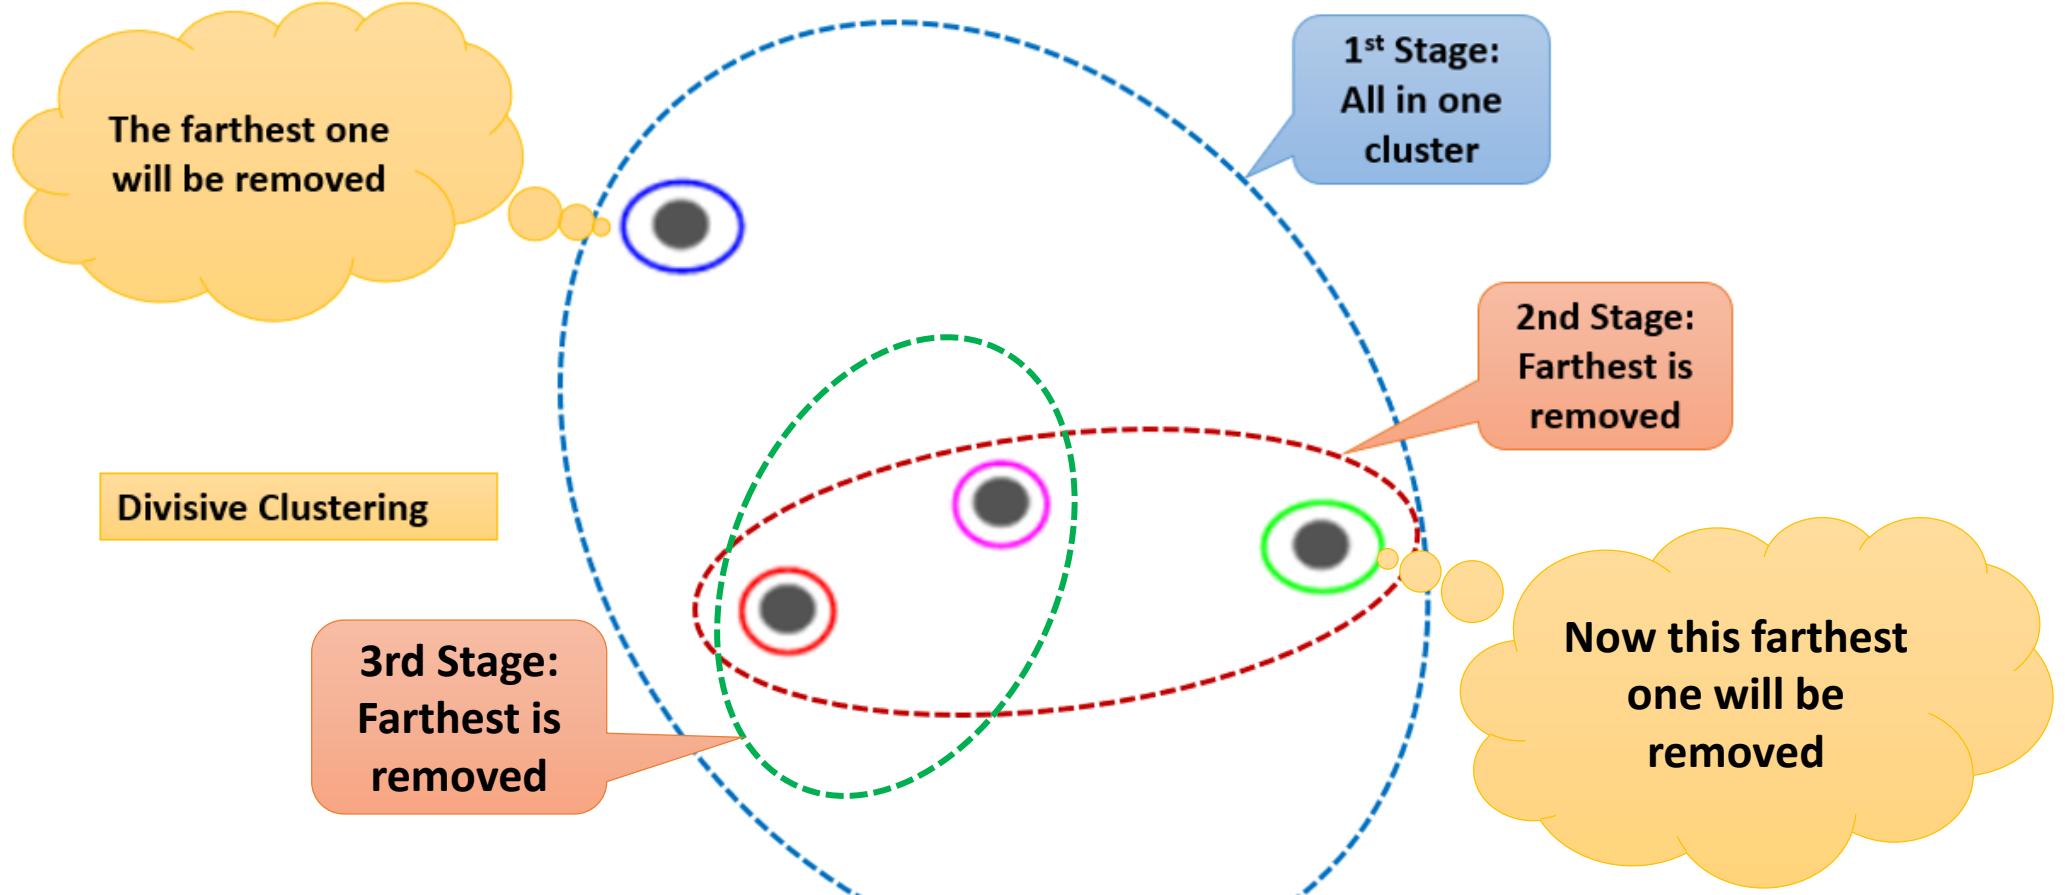
3rd merging of two nearest clusters/(data points)

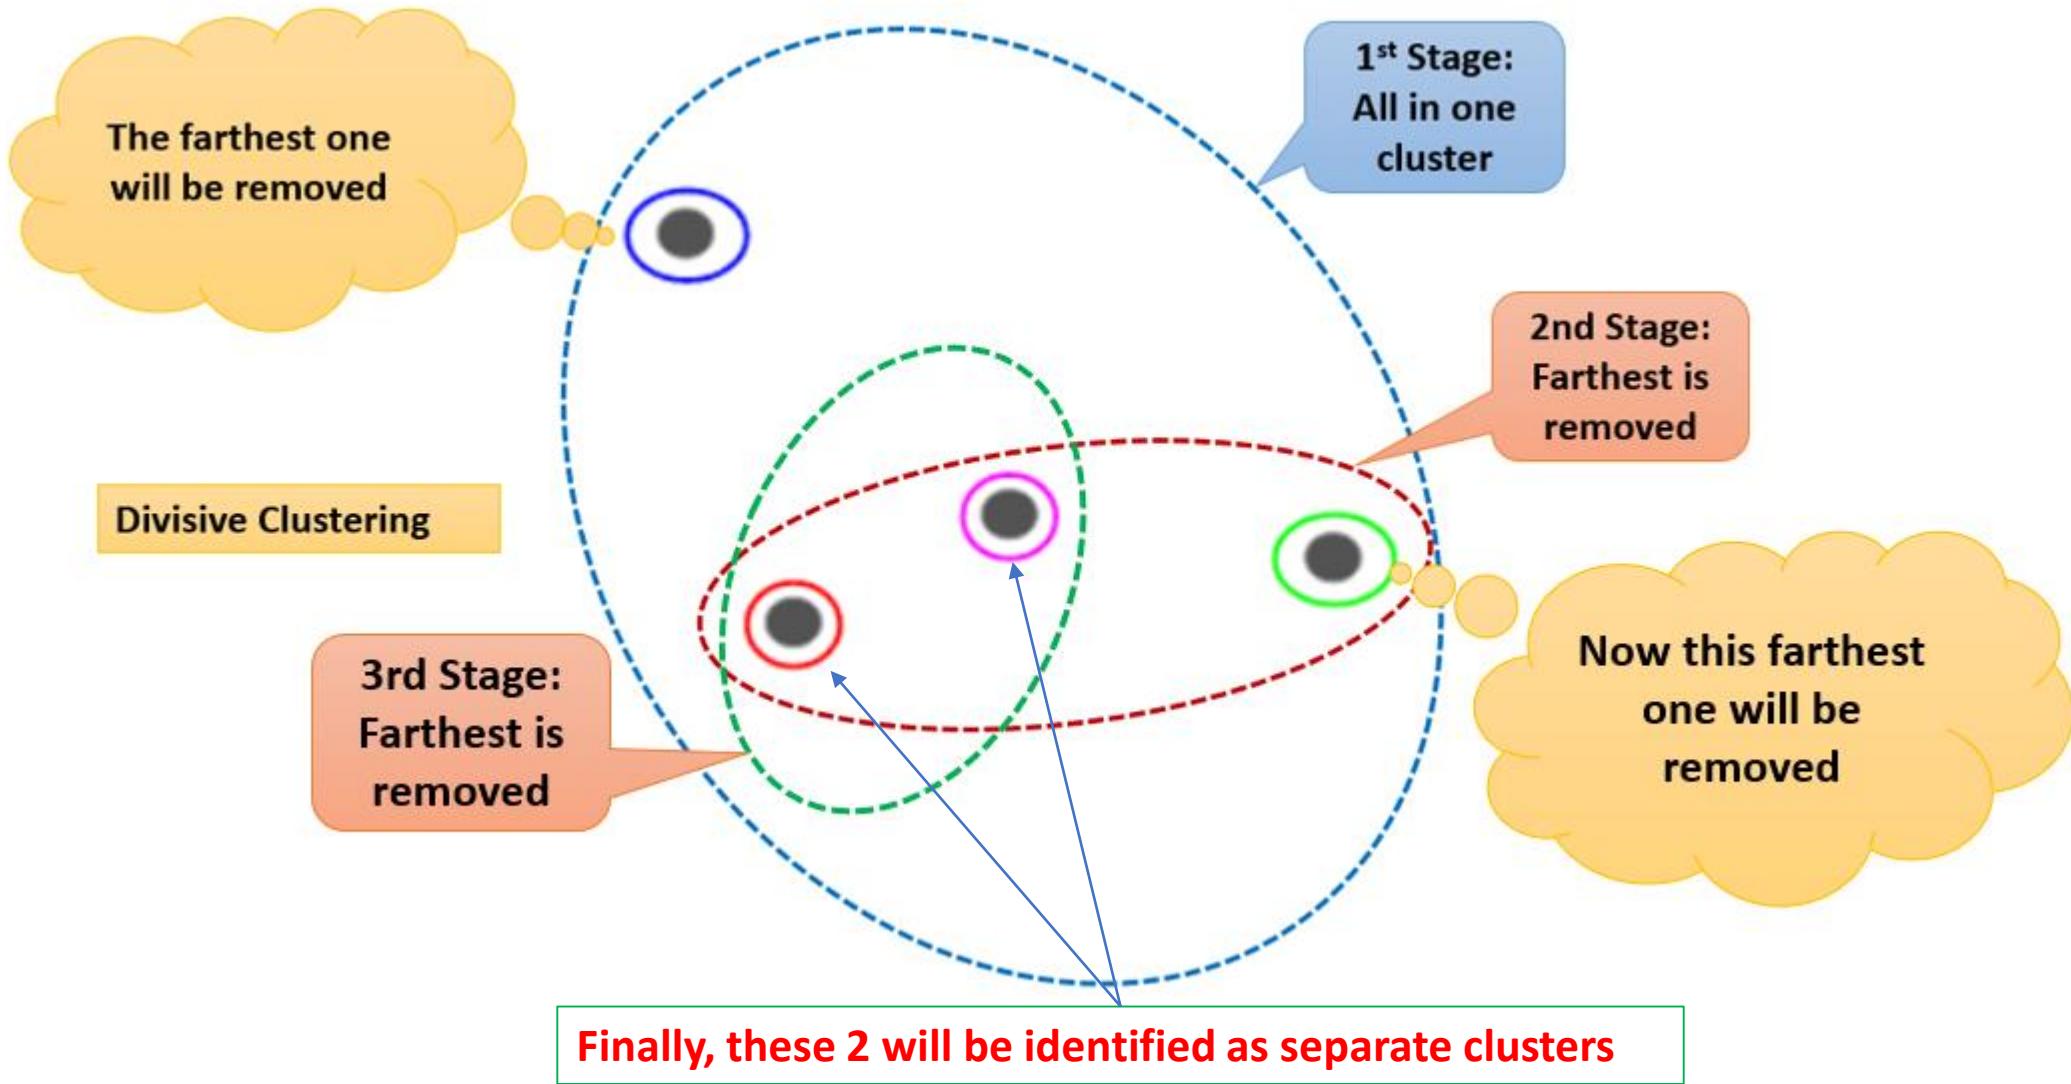


Divisive Hierarchical Clustering

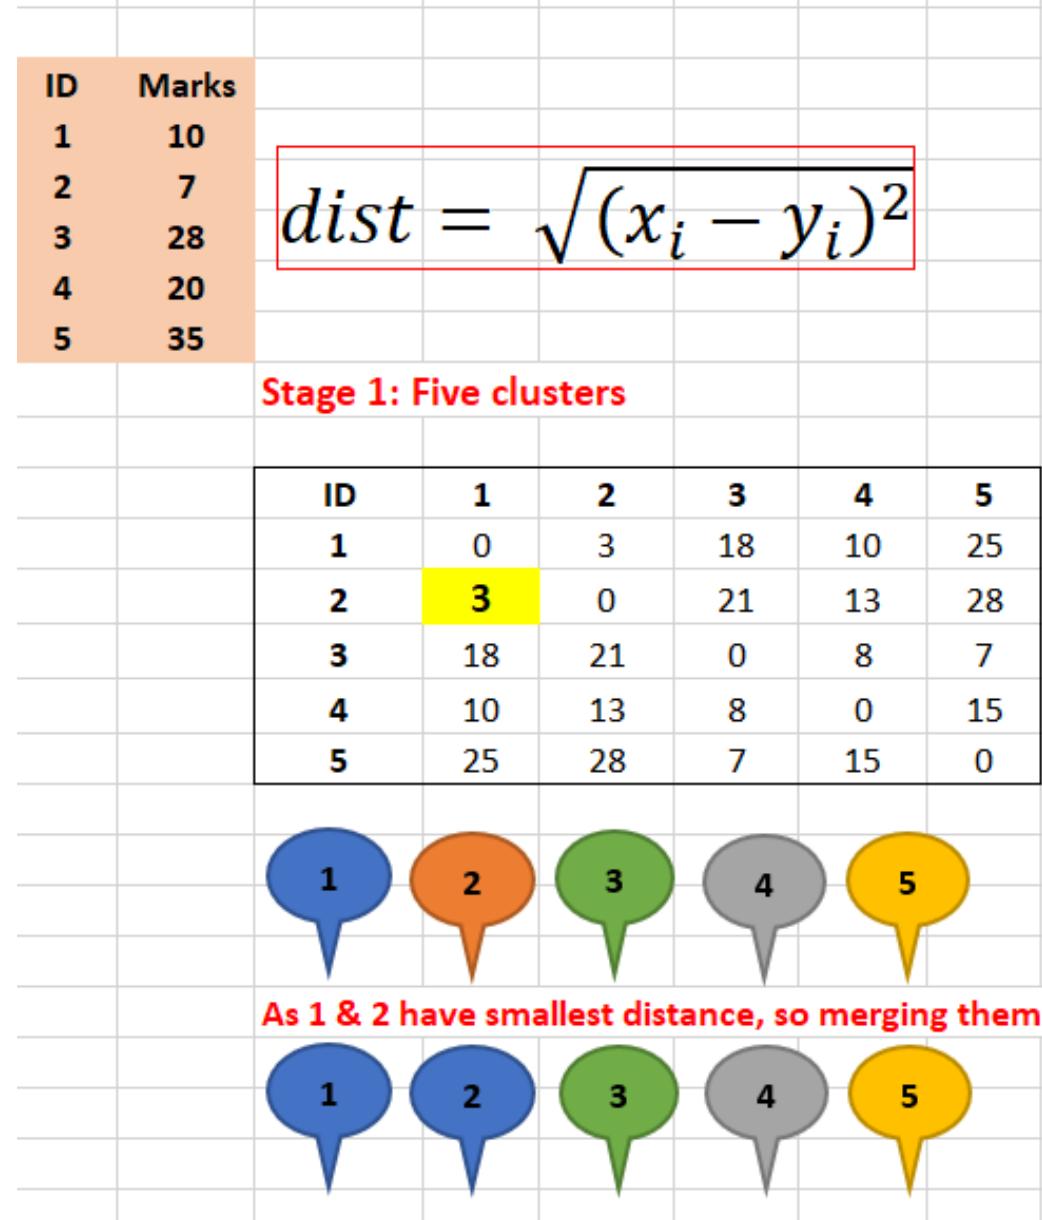








Agglomerative Clustering (Adding or Merging)



ID	Marks
1	10
2	7
3	28
4	20
5	35

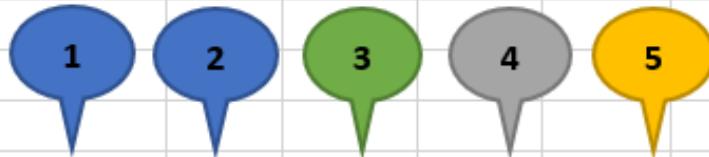
$$dist = \sqrt{(x_i - y_i)^2}$$

Stage 1: Five clusters

ID	1	2	3	4	5
1	0	3	18	10	25
2	3	0	21	13	28
3	18	21	0	8	7
4	10	13	8	0	15
5	25	28	7	15	0



As 1 & 2 have smallest distance, so merging them



ID	Marks
1-2	10
3	28
4	20
5	35

Stage 2: Four clusters

ID	1-2	3	4	5
1-2	0	18	10	25
3	18	0	8	7
4	10	8	0	15
5	25	7	15	0



As 3 & 5 have smallest distance, so merging them



ID	Marks
1-2	10
3	28
4	20
5	35

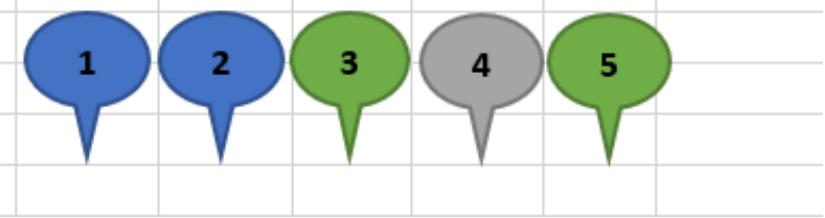
max of 10 & 7

Stage 2: Four clusters

ID	1-2	3	4	5
1-2	0	18	10	25
3	18	0	8	7
4	10	8	0	15
5	25	7	15	0



As 3 & 5 have smallest distance, so merging them



ID	Marks
1-2	10
3-5	35
4	20

max of 28 & 35

Stage 3: Three clusters

ID	1-2	3-5	4
1-2	0	15	10
3-5	15	0	15
4	10	15	0



As 1-2 & 4 have smallest distance, so merging them



ID	Marks
1-2	10
3-5	35
4	20

max of 28 & 35

Stage 3: Three clusters

ID	1-2	3-5	4
1-2	0	15	10
3-5	15	0	15
4	10	15	0



As 1-2 & 4 have smallest distance, so merging them



ID	Marks
1-2-4	20
3-5	35

max of 10 & 20
max of 28 & 35

Stage 4: Two clusters

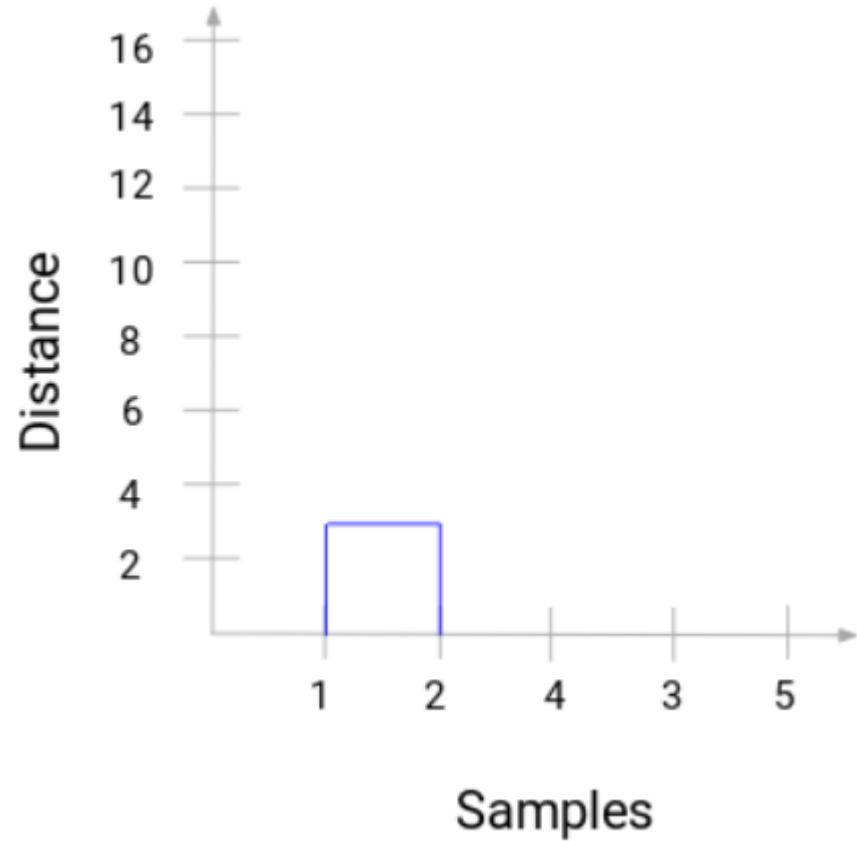
ID	1-2-4	3-5
1-2-4	0	15
3-5	15	0



As only 2 clusters remained, we merge them



Dendrogram



ID	Marks
1	10
2	7
3	28
4	20
5	35

$$dist = \sqrt{(x_i - y_i)^2}$$

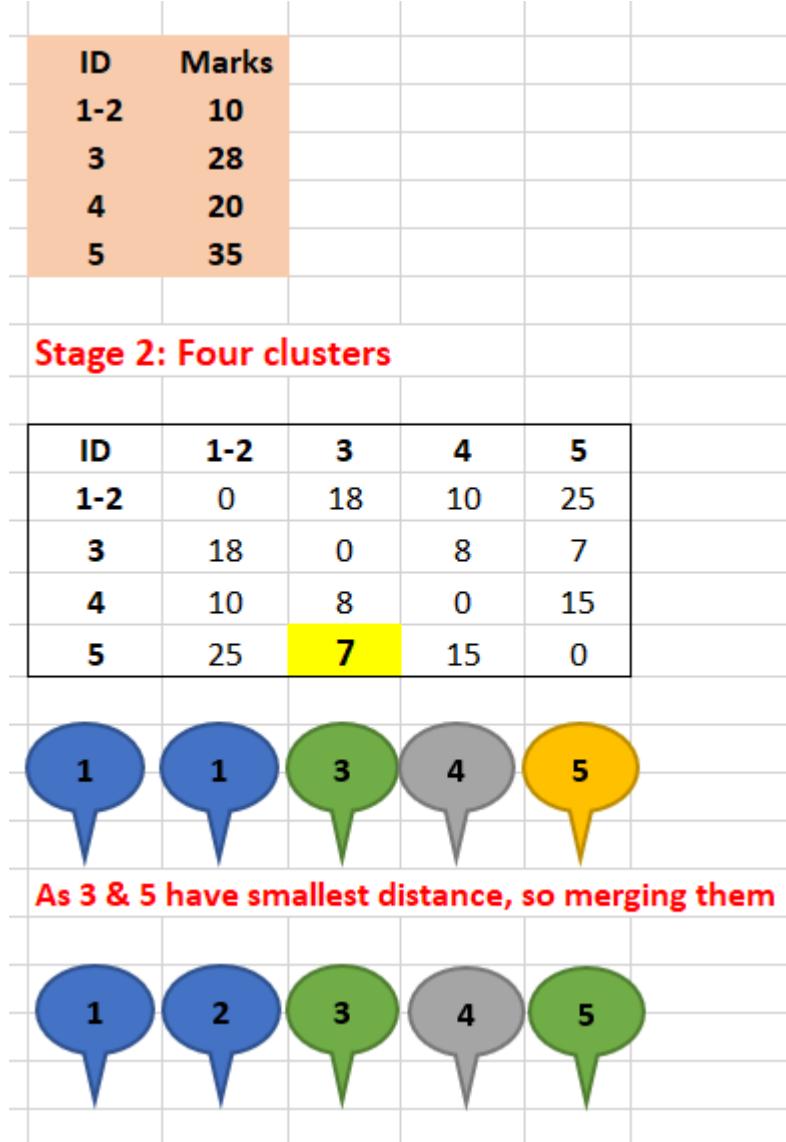
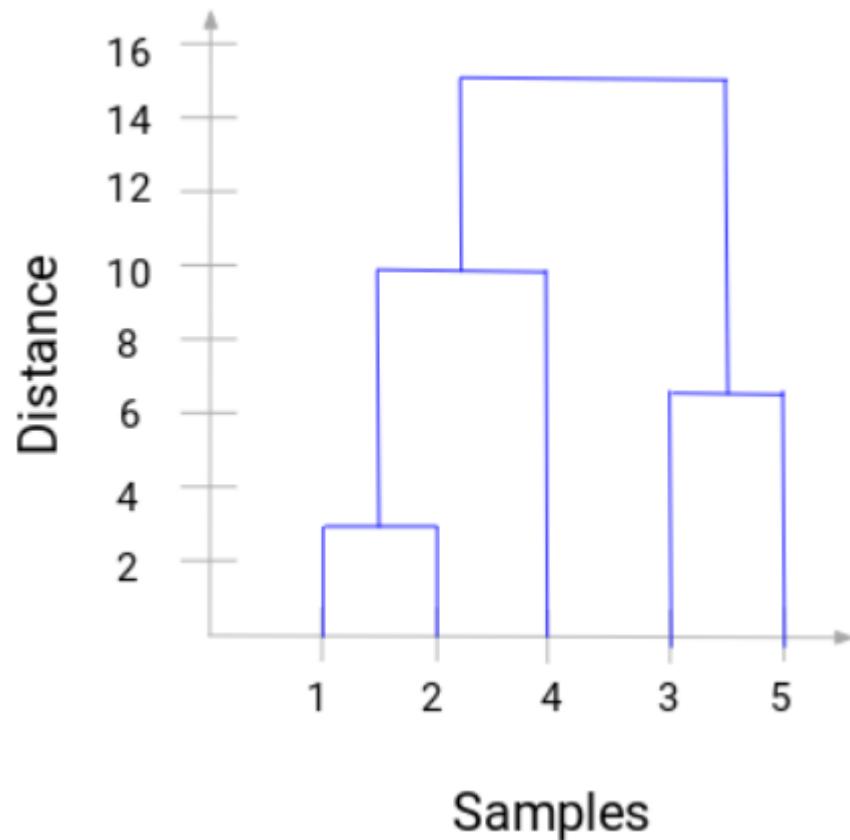
Stage 1: Five clusters

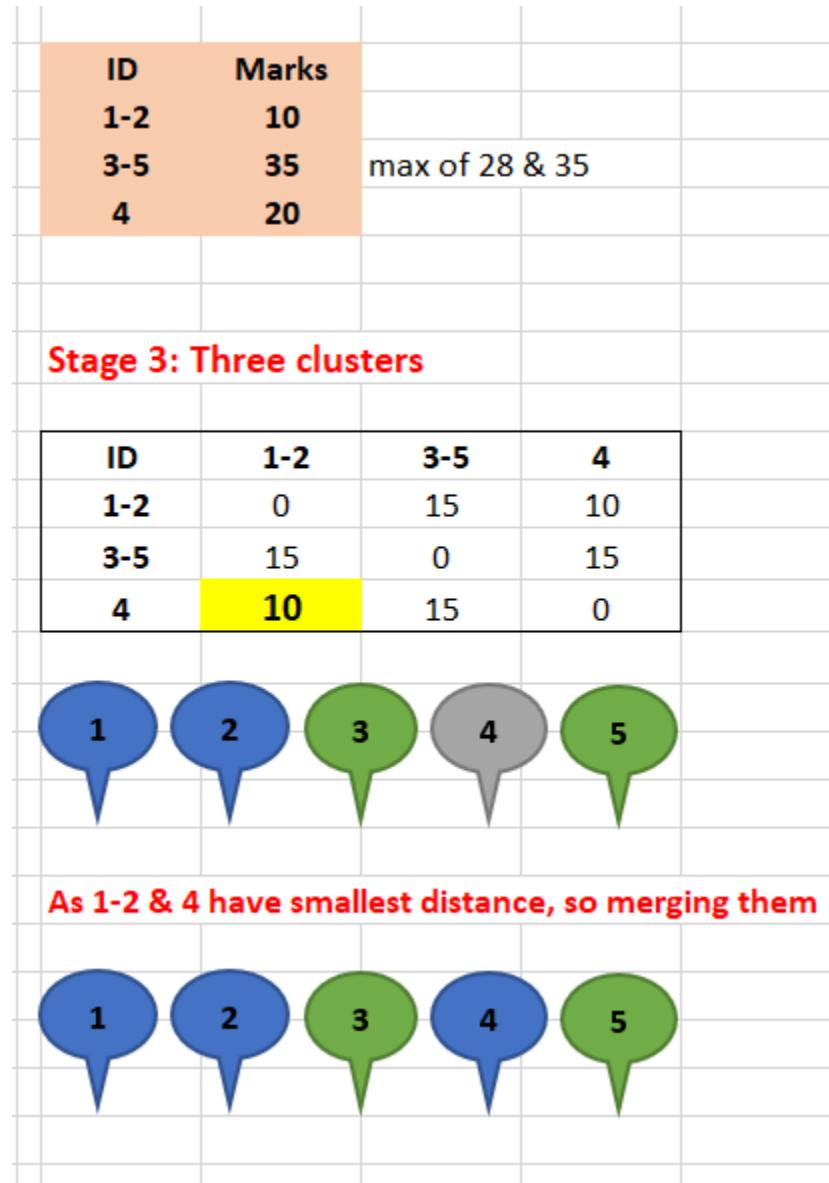
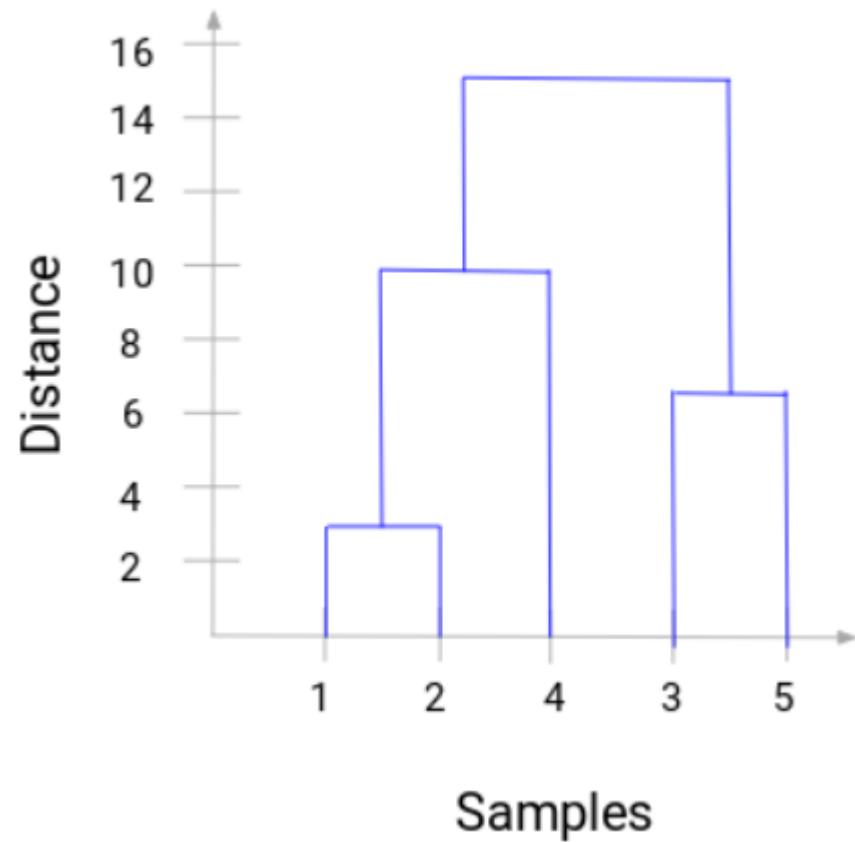
ID	1	2	3	4	5
1	0	3	18	10	25
2	3	0	21	13	28
3	18	21	0	8	7
4	10	13	8	0	15
5	25	28	7	15	0

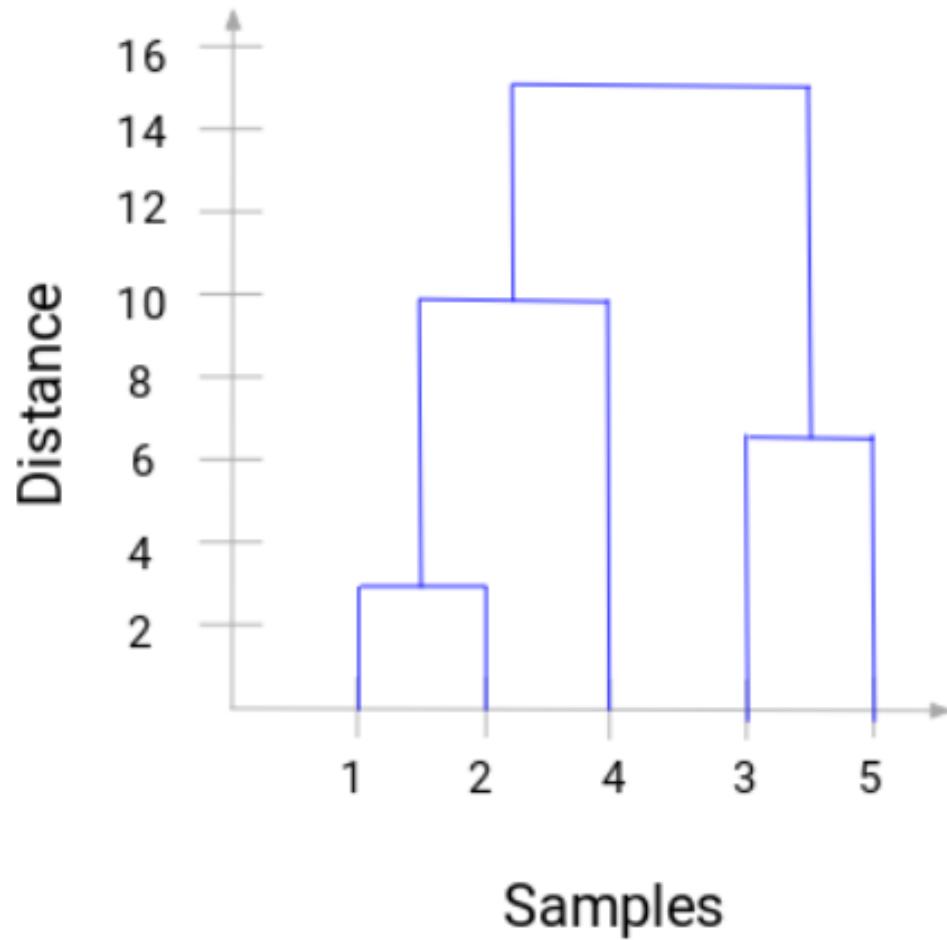


As 1 & 2 have smallest distance, so merging them









ID	Marks	
1-2-4	20	max of 10 & 20
3-5	35	max of 28 & 35

Stage 4: Two clusters

ID	1-2-4	3-5
1-2-4	0	15
3-5	15	0

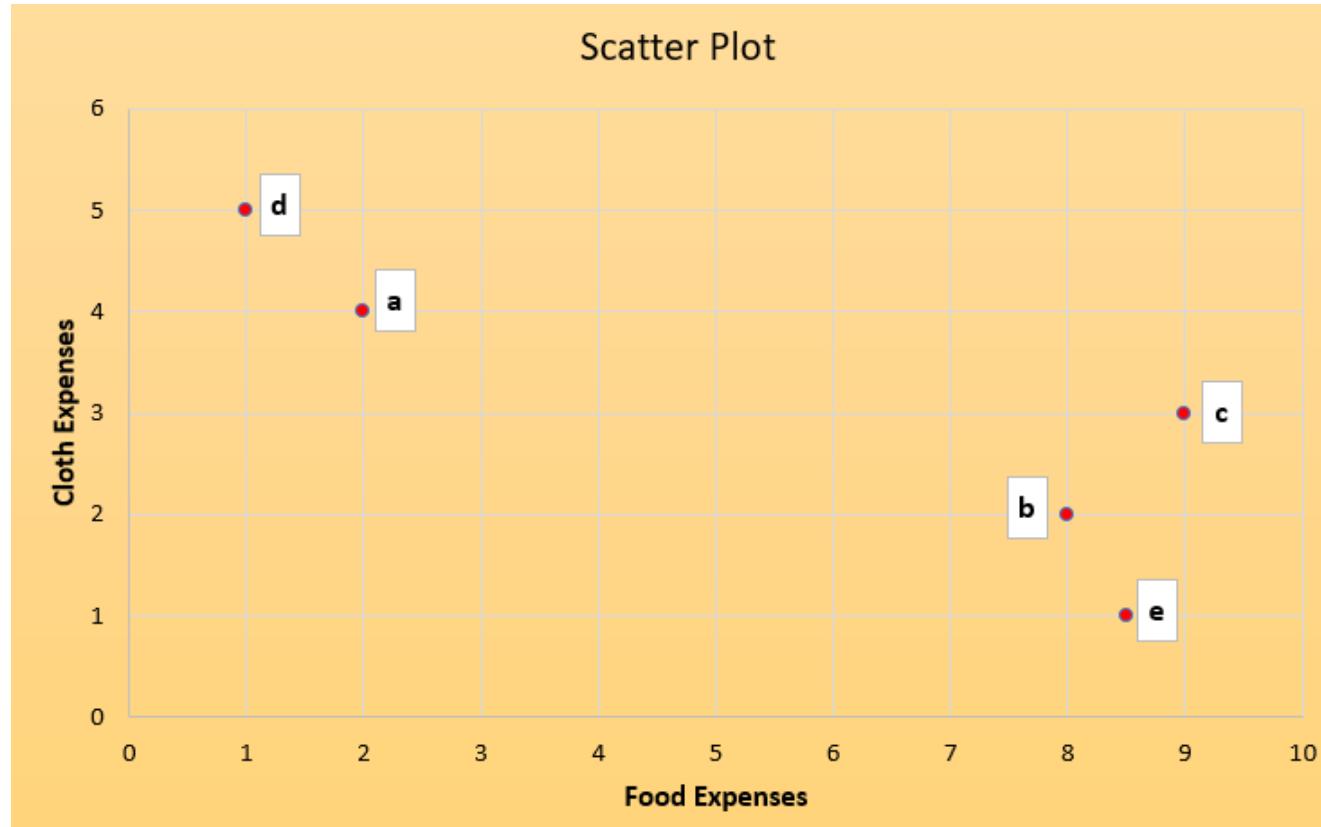


As only 2 clusters remained, we merge them

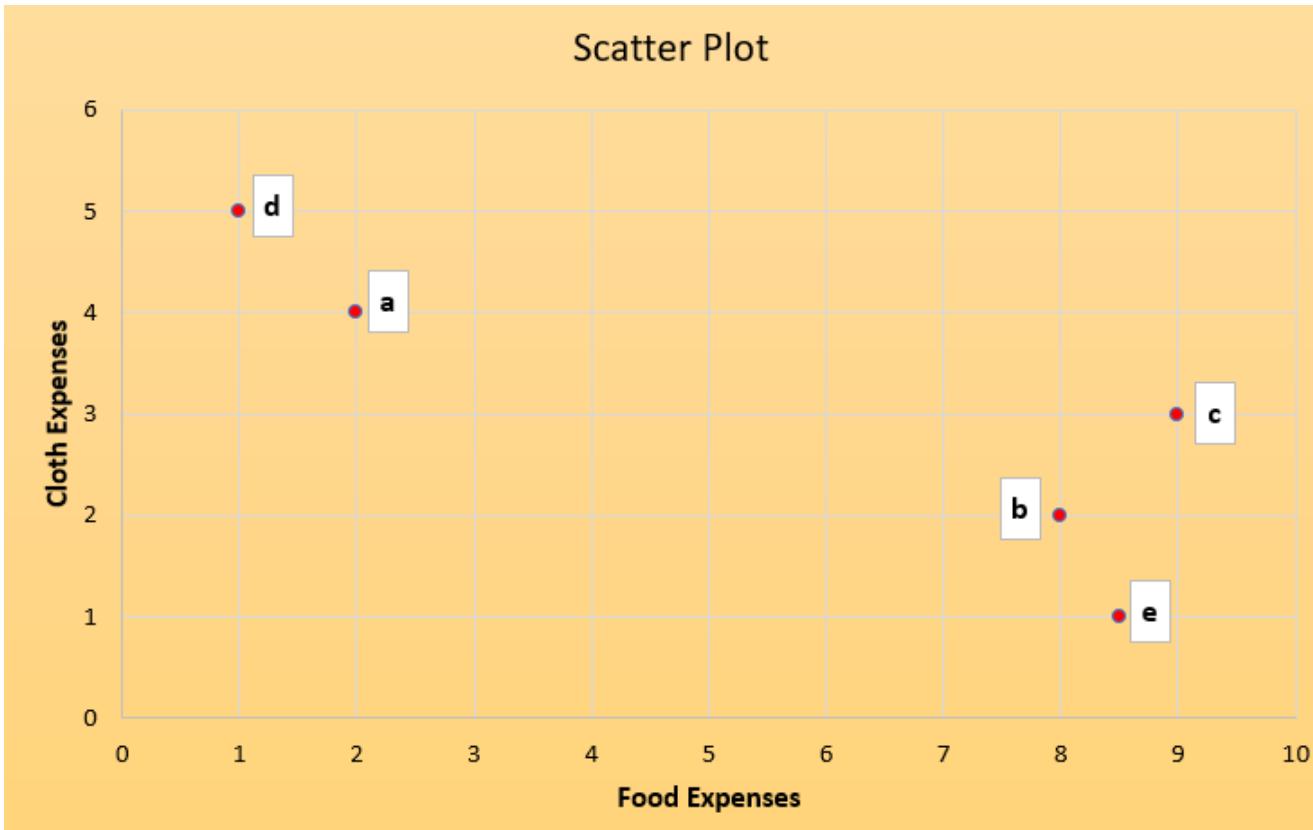


Data with X and Y: Single Linkage (Min)

	Food Exp	Cloth Exp
a	X	Y
b	2	4
c	8	2
d	9	3
e	1	5
	8.5	1



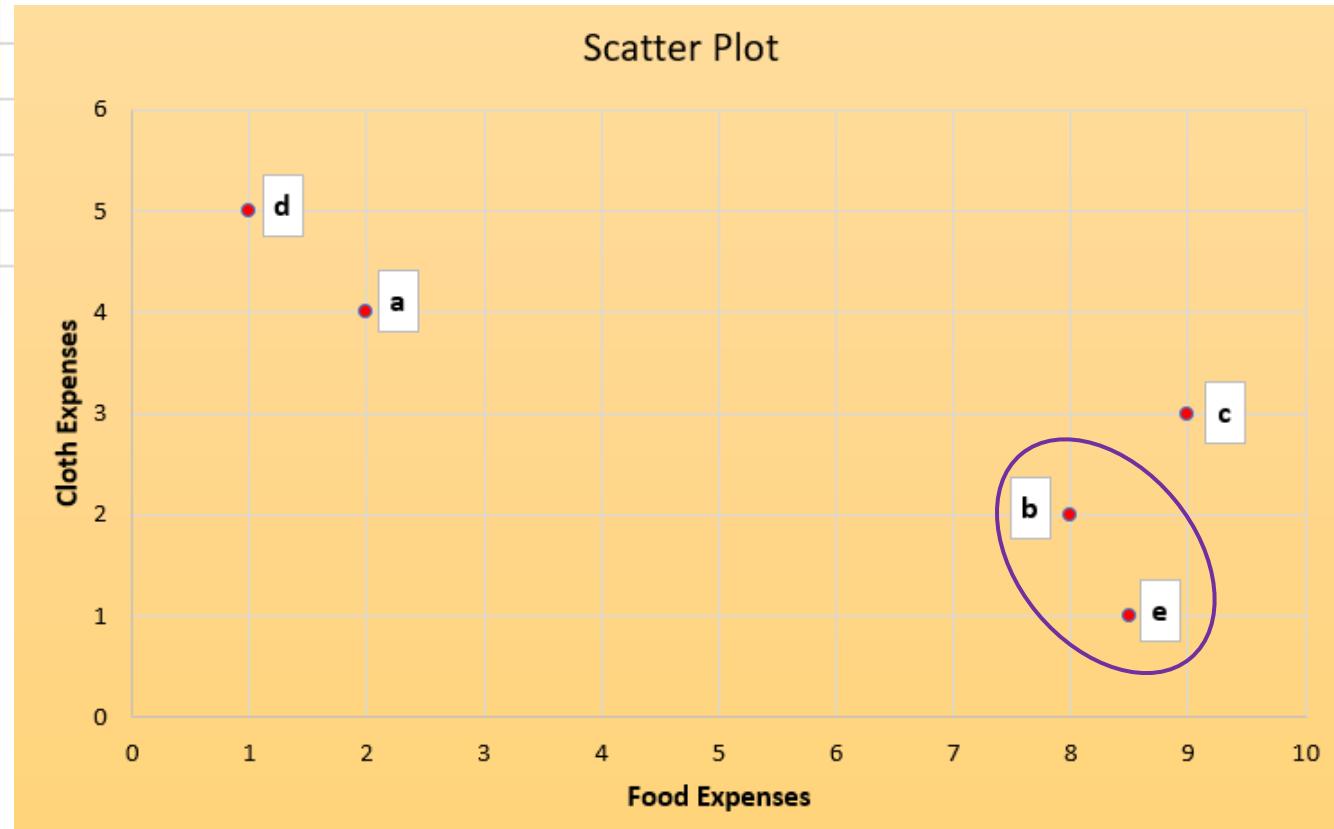
Give a try...make agglomerative process intuitively!



Stage 1: b and e can be clubbed

	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00

Now **b** and **e** are in one cluster. How their values will be treated for finding DIST vis-à-vis another point/s?



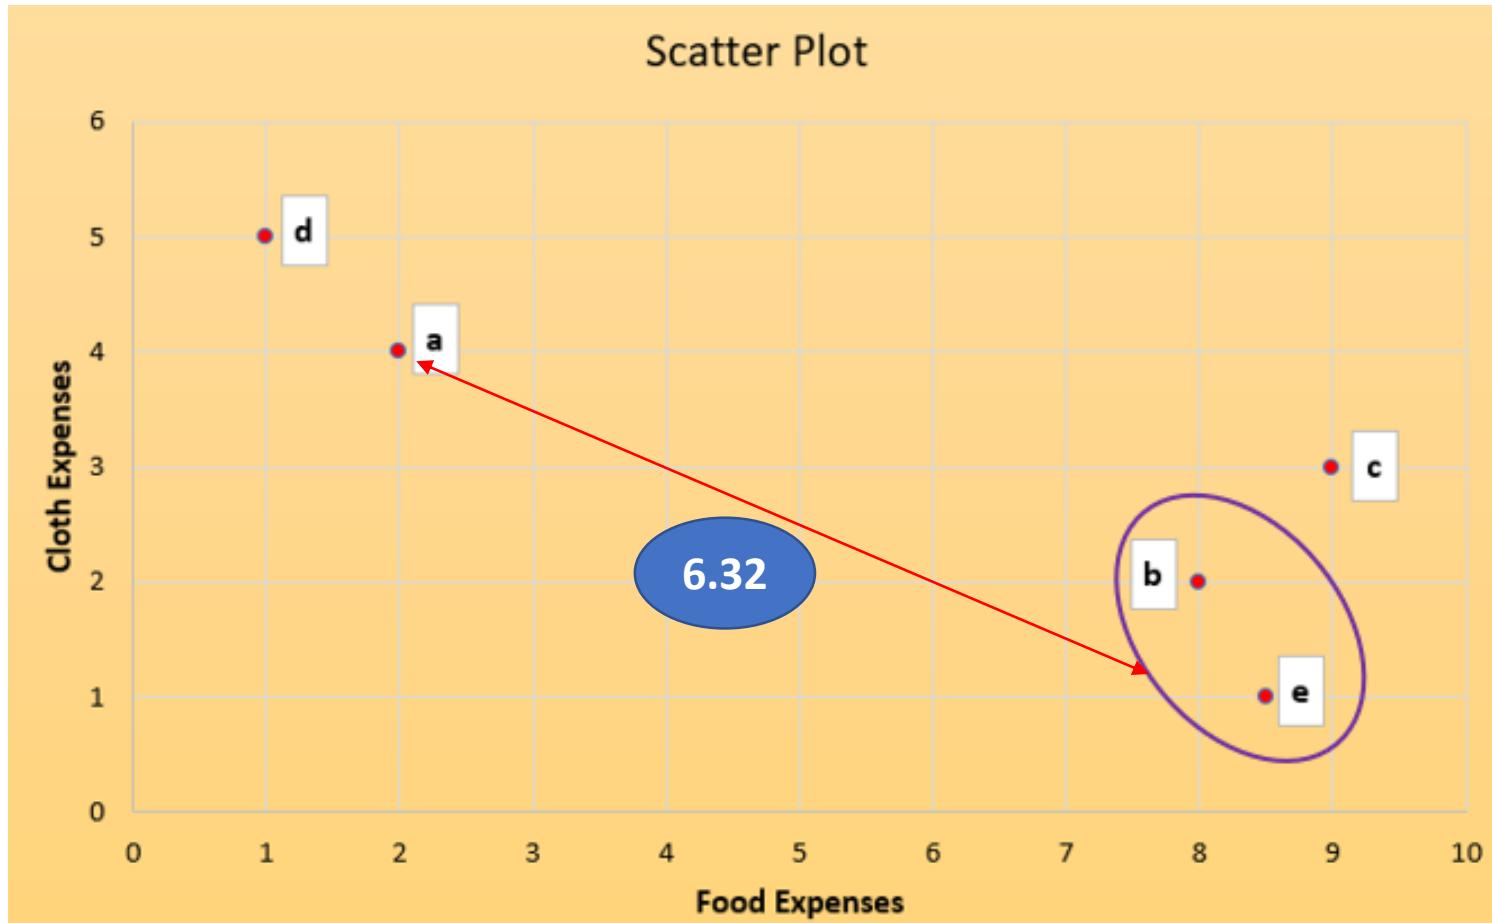
Distance between b, e and a

	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00

$$D(be, a) = \min\{D(b, a), D(e, a)\}$$

$$D(be, a) = \min\{6.32, 7.16\}$$

$$D(be, a) = 6.32$$



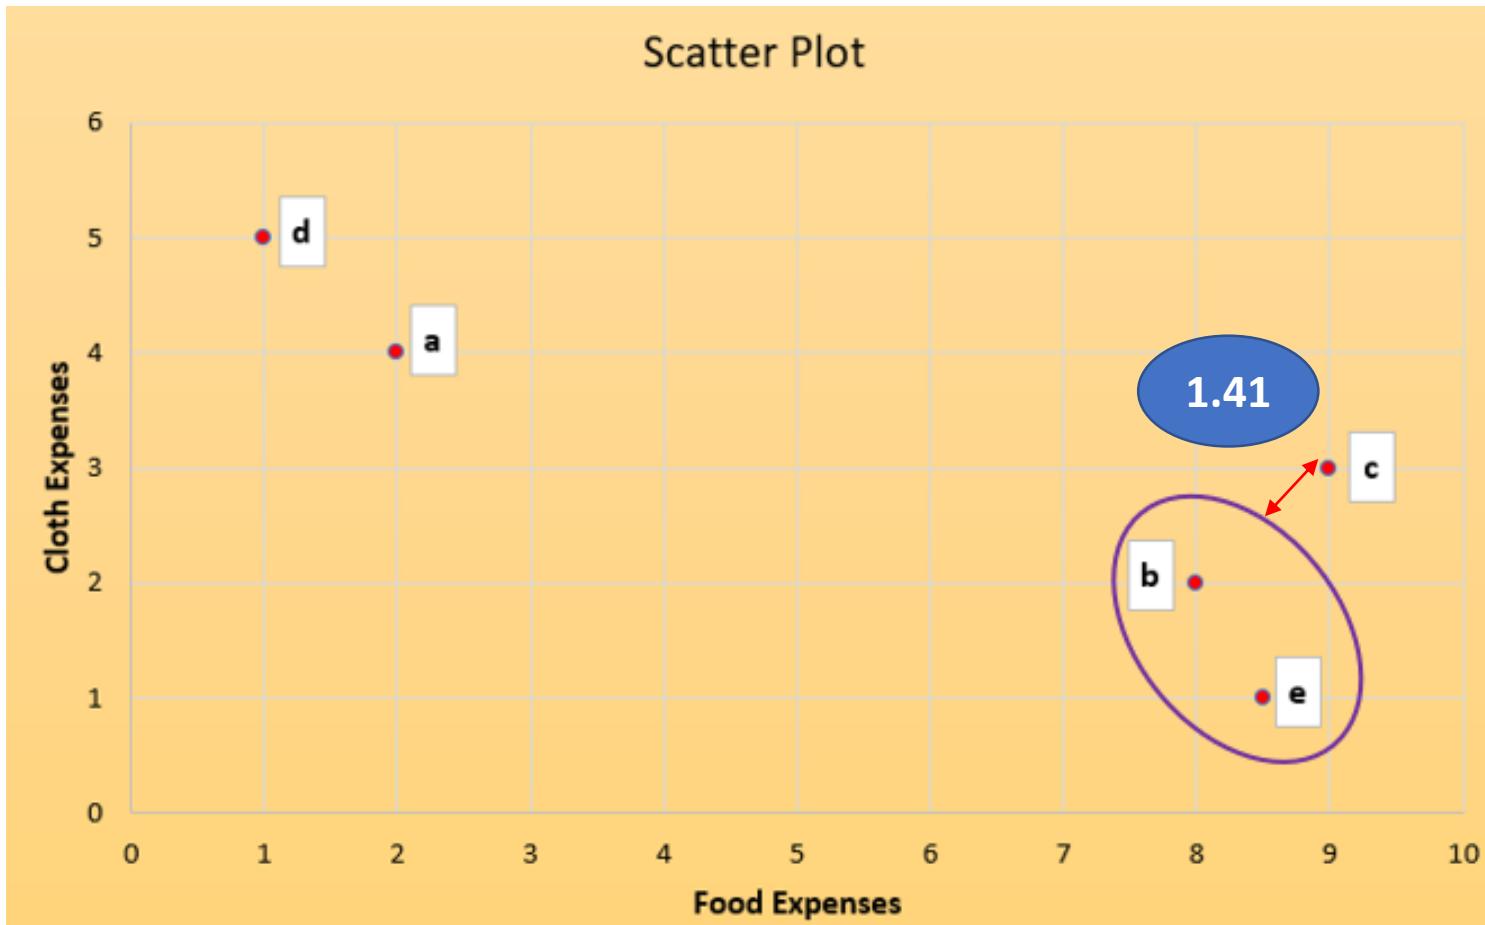
Distance between b, e and c

	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00

$$D(be, c) = \min\{D(b, c), D(e, c)\}$$

$$D(be, c) = \min\{1.41, 2.06\}$$

$$D(be, c) = 1.41$$



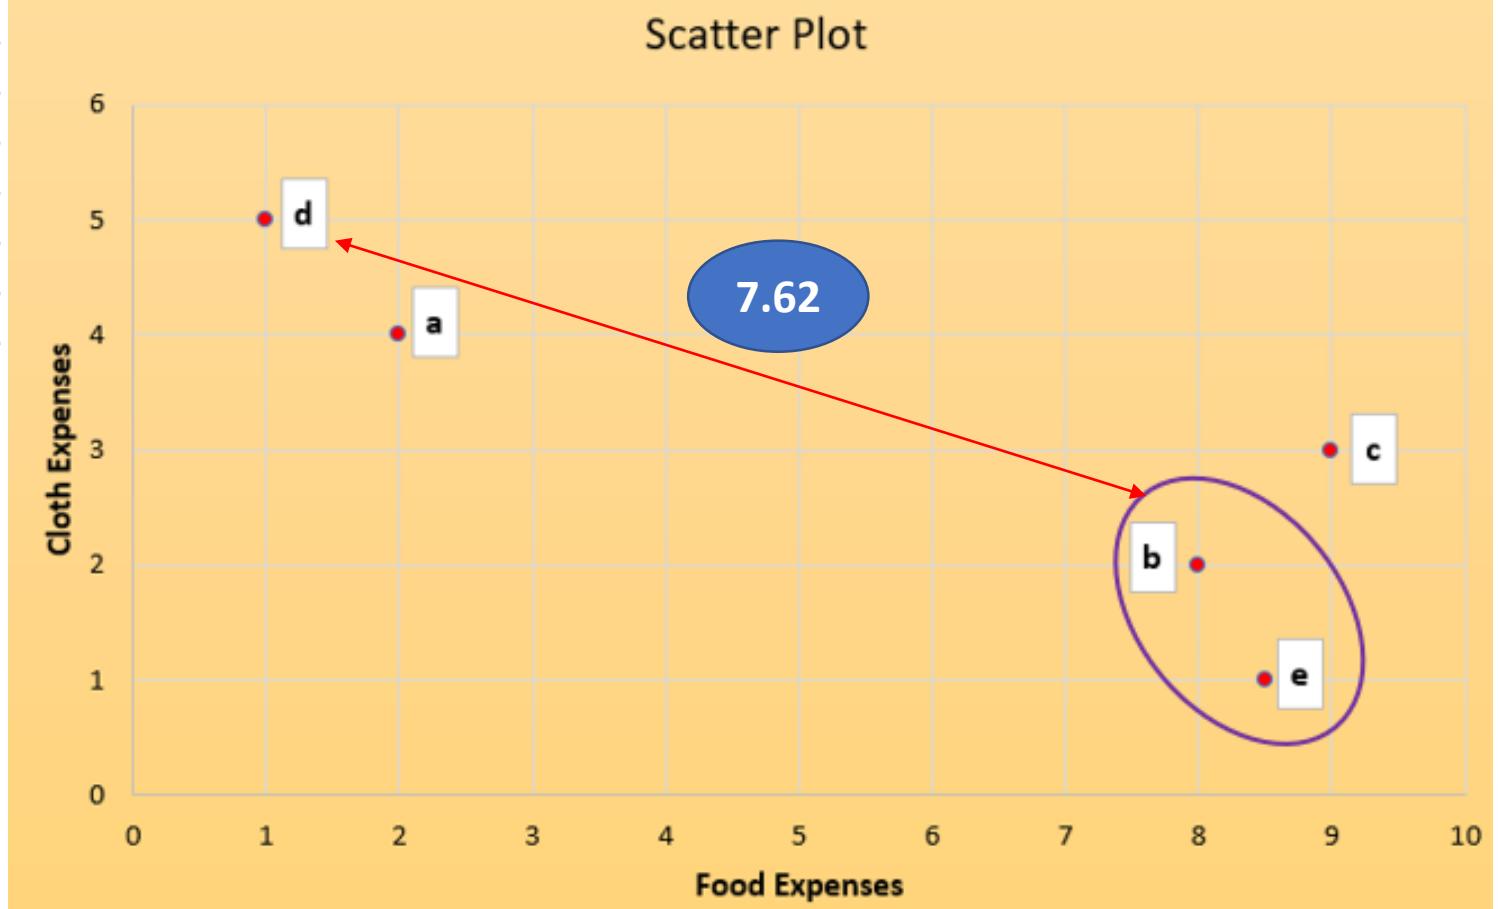
Distance between b, e and d

	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00

$$D(be, d) = \min\{D(b, d), D(e, d)\}$$

$$D(be, d) = \min\{7.62, 8.50\}$$

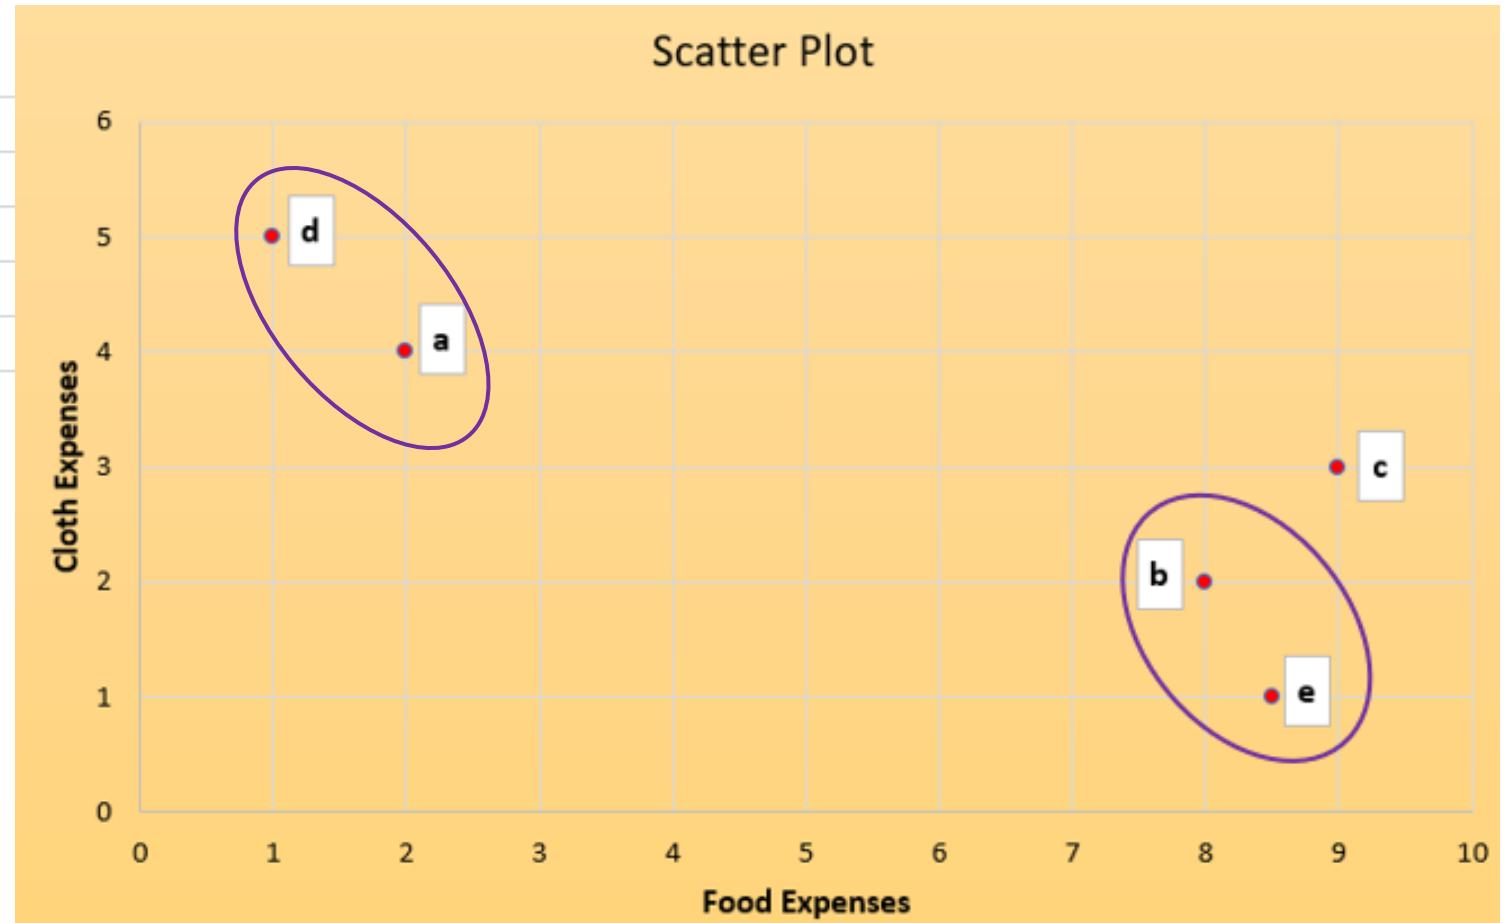
$$D(be, d) = 7.62$$



Stage 2: d and a can be clubbed

	b	e	a	c	d
b	e	0	6.32	1.41	7.62
a		0	7.07	1.41	
c			0	8.25	
d				0	

Distance
between
clusters (b, e)
and (a, d)?



Distance between clusters (b, e) and (a, d)?

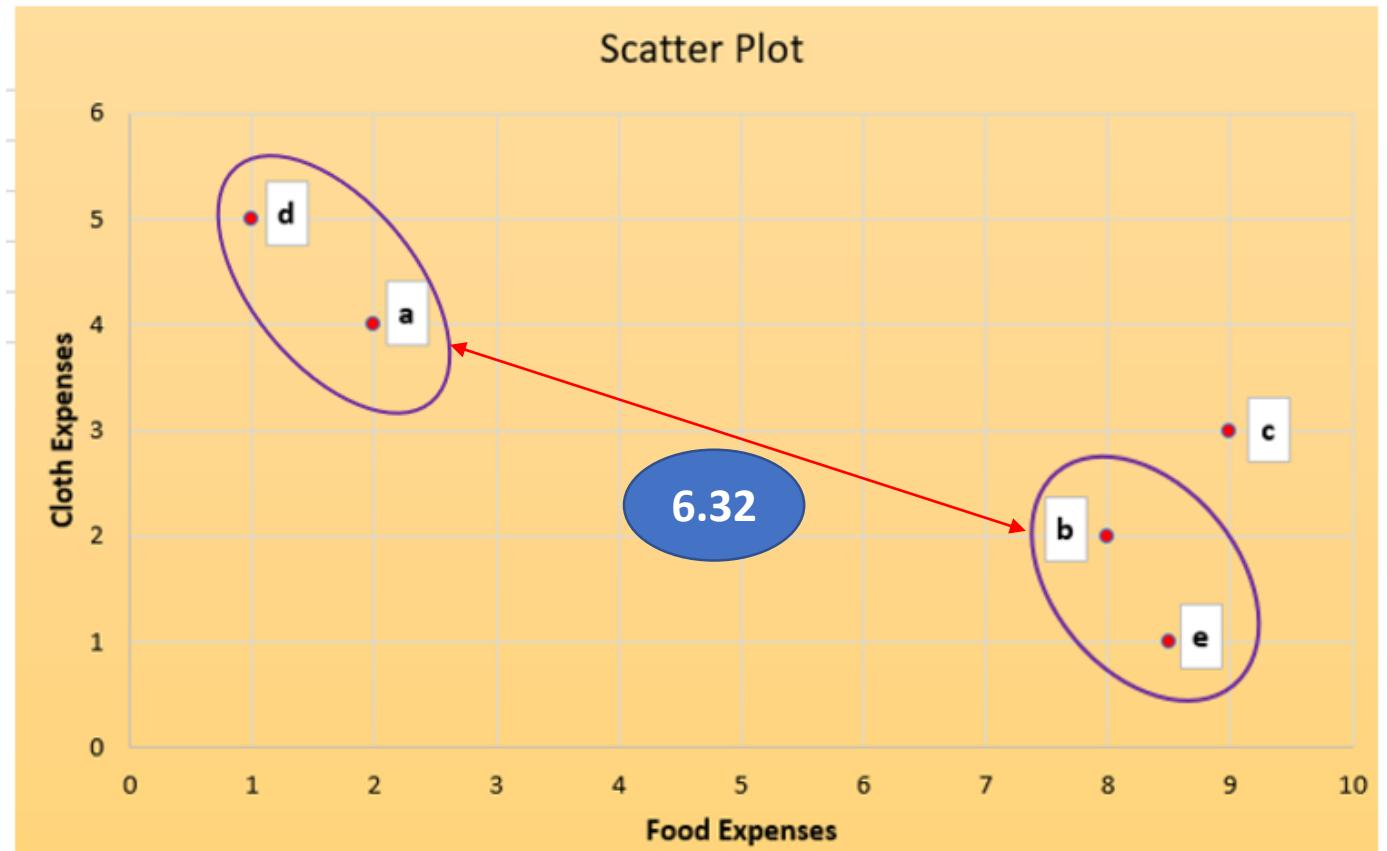
	b e	a	c	d
b e	0	6.32	1.41	7.62
a		0	7.07	1.41
c			0	8.25
d				0

$$D(b e, a d) = \min\{D(b e, a), D(b e, d)\}$$

$$D(b e, a d) = \min\{6.32, 7.62\}$$

$$D(b e, a d) = 6.32$$

	b e	a d	c
b e	0	6.32	1.41
a d		0	7.07
c			0



Distance between clusters (a, d) and (c)?

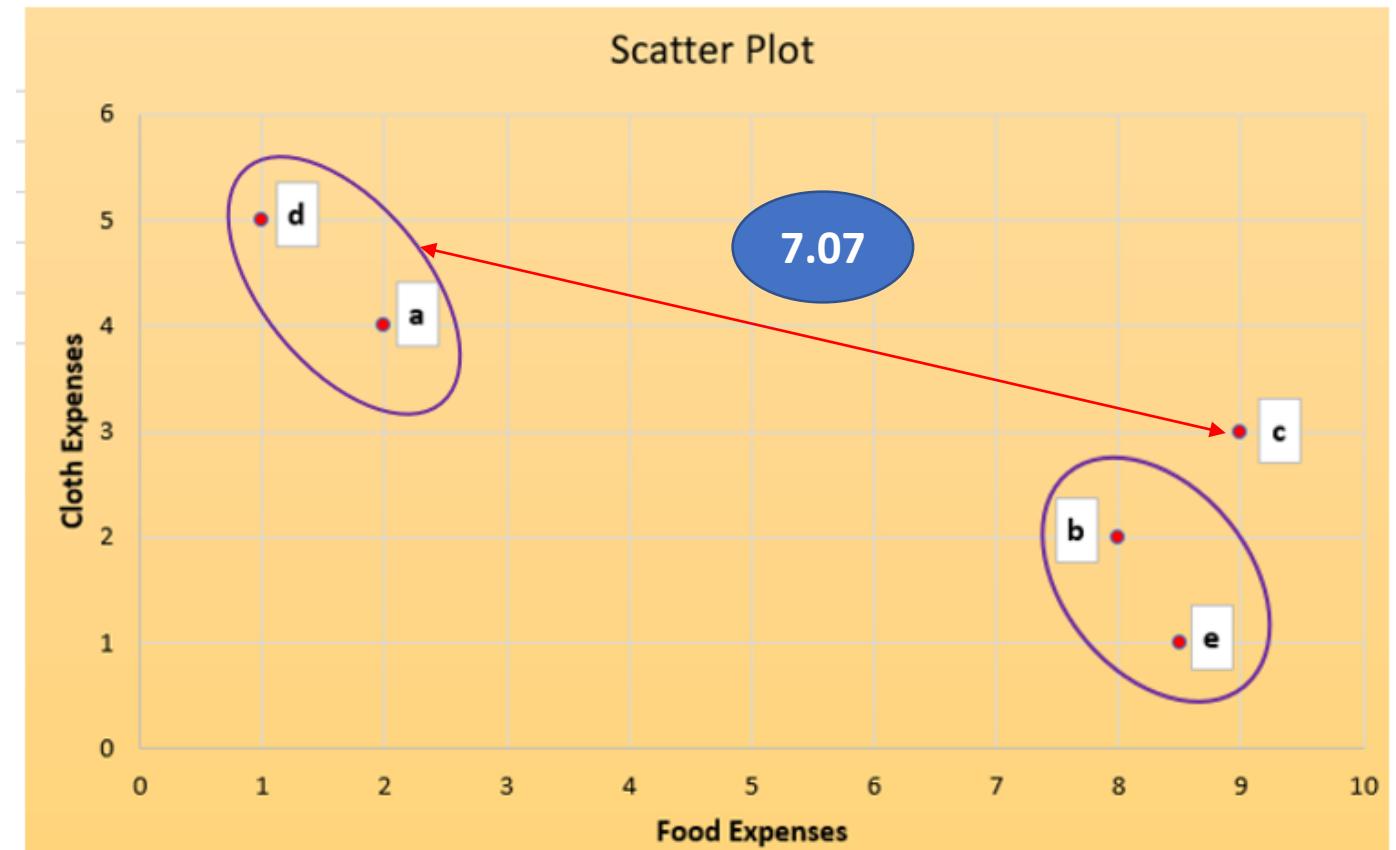
	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00

	b e	a d	c
b e	0	6.32	1.41
a d		0	7.07
c			0

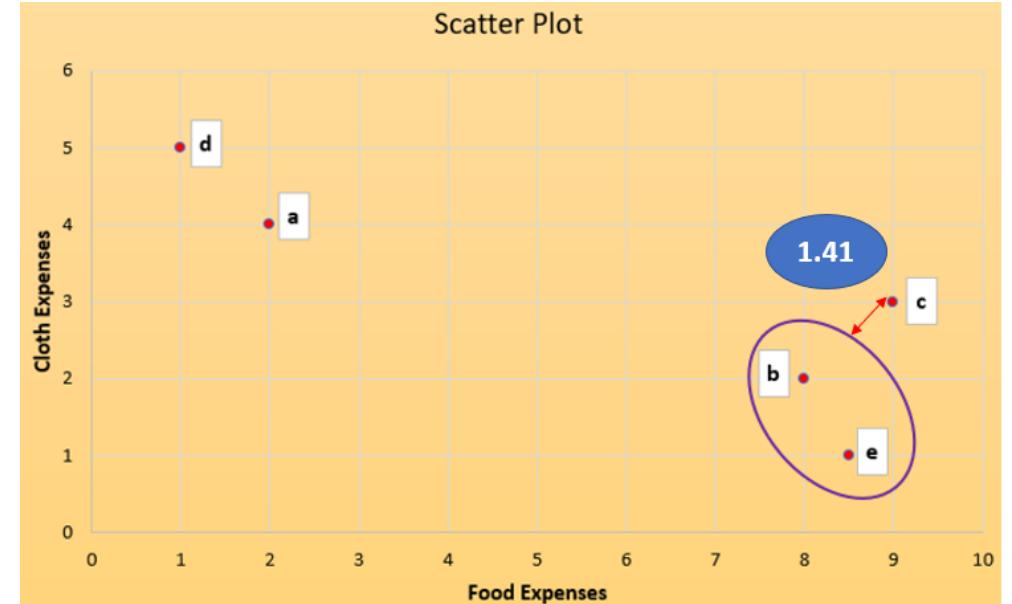
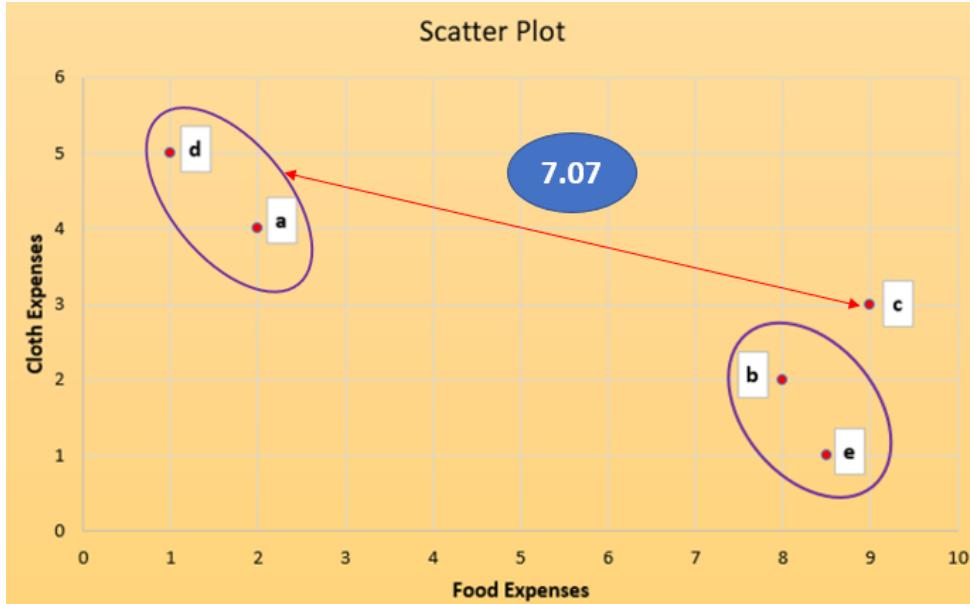
$$D(ad, c) = \min\{D(a, c), D(d, c)\}$$

$$D(ad, c) = \min\{7.07, 8.25\}$$

$$D(ad, c) = 7.07$$



Where c should go?



Stage 3: Cluster b e must envelop c

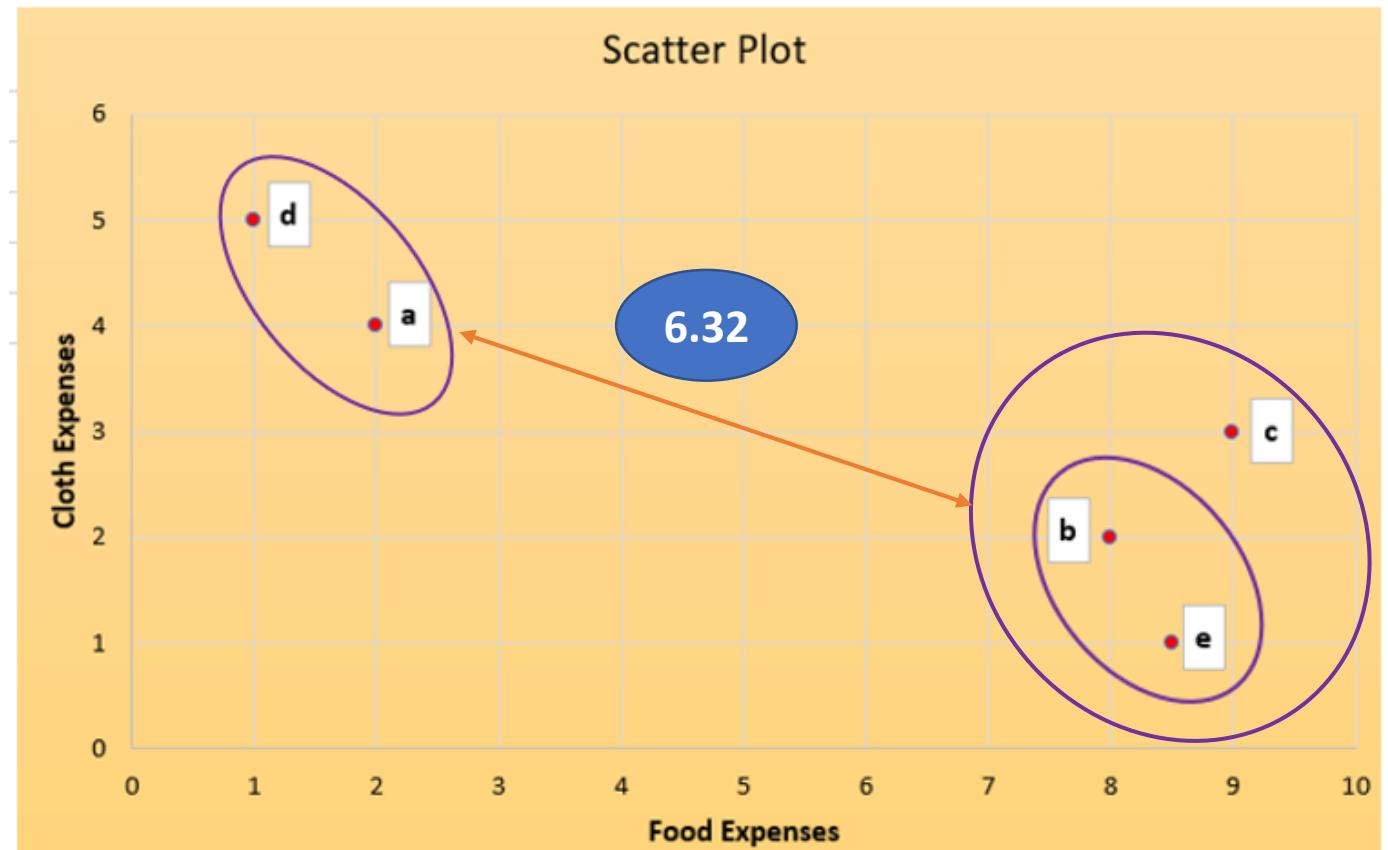
Stage 4

	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00

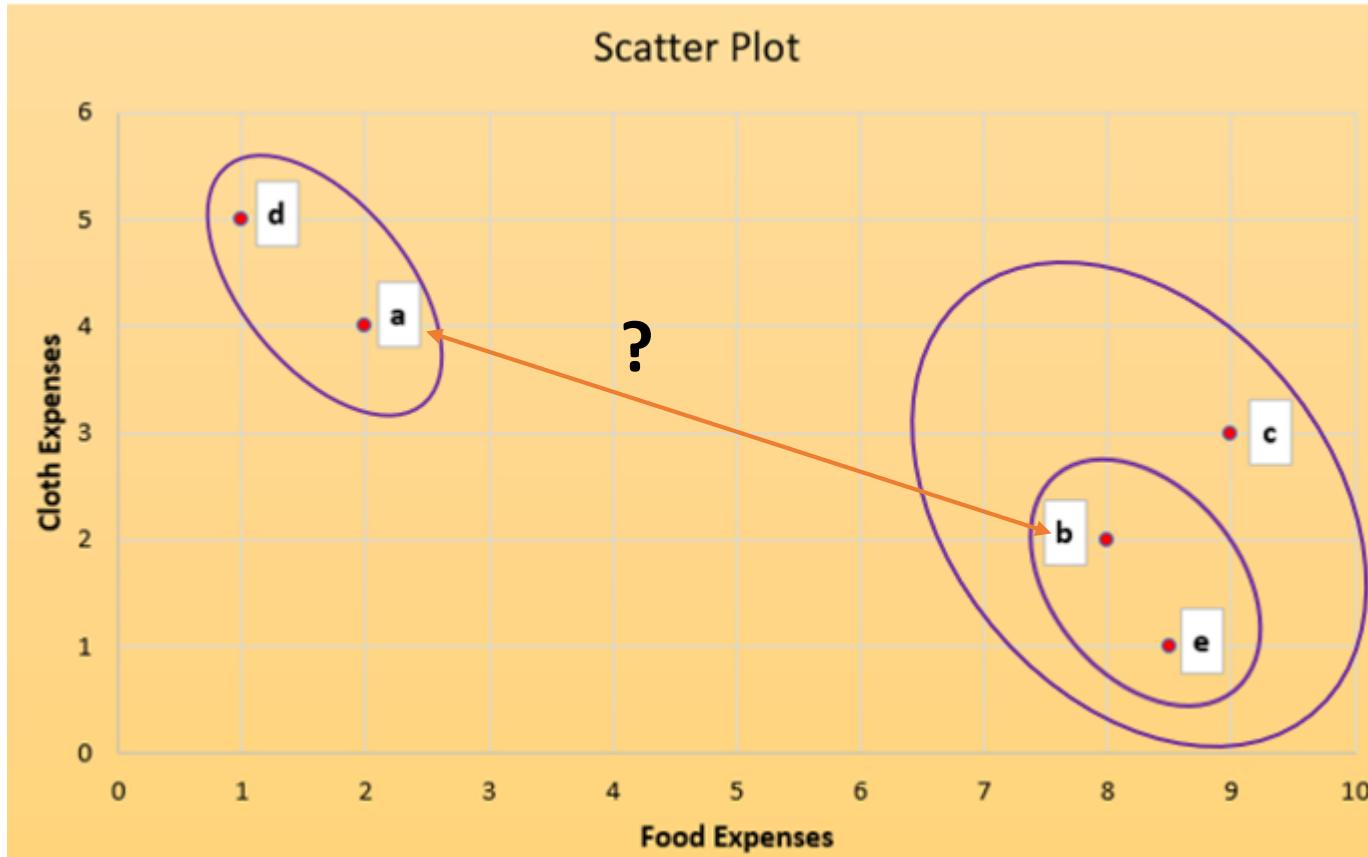
	b e c	a d
b e c	0	6.32
a d		0

$$D(bec, ad) = \min\{D(b, a), D(b, d), D(e, a), D(e, d), D(c, a), D(c, d)\}$$

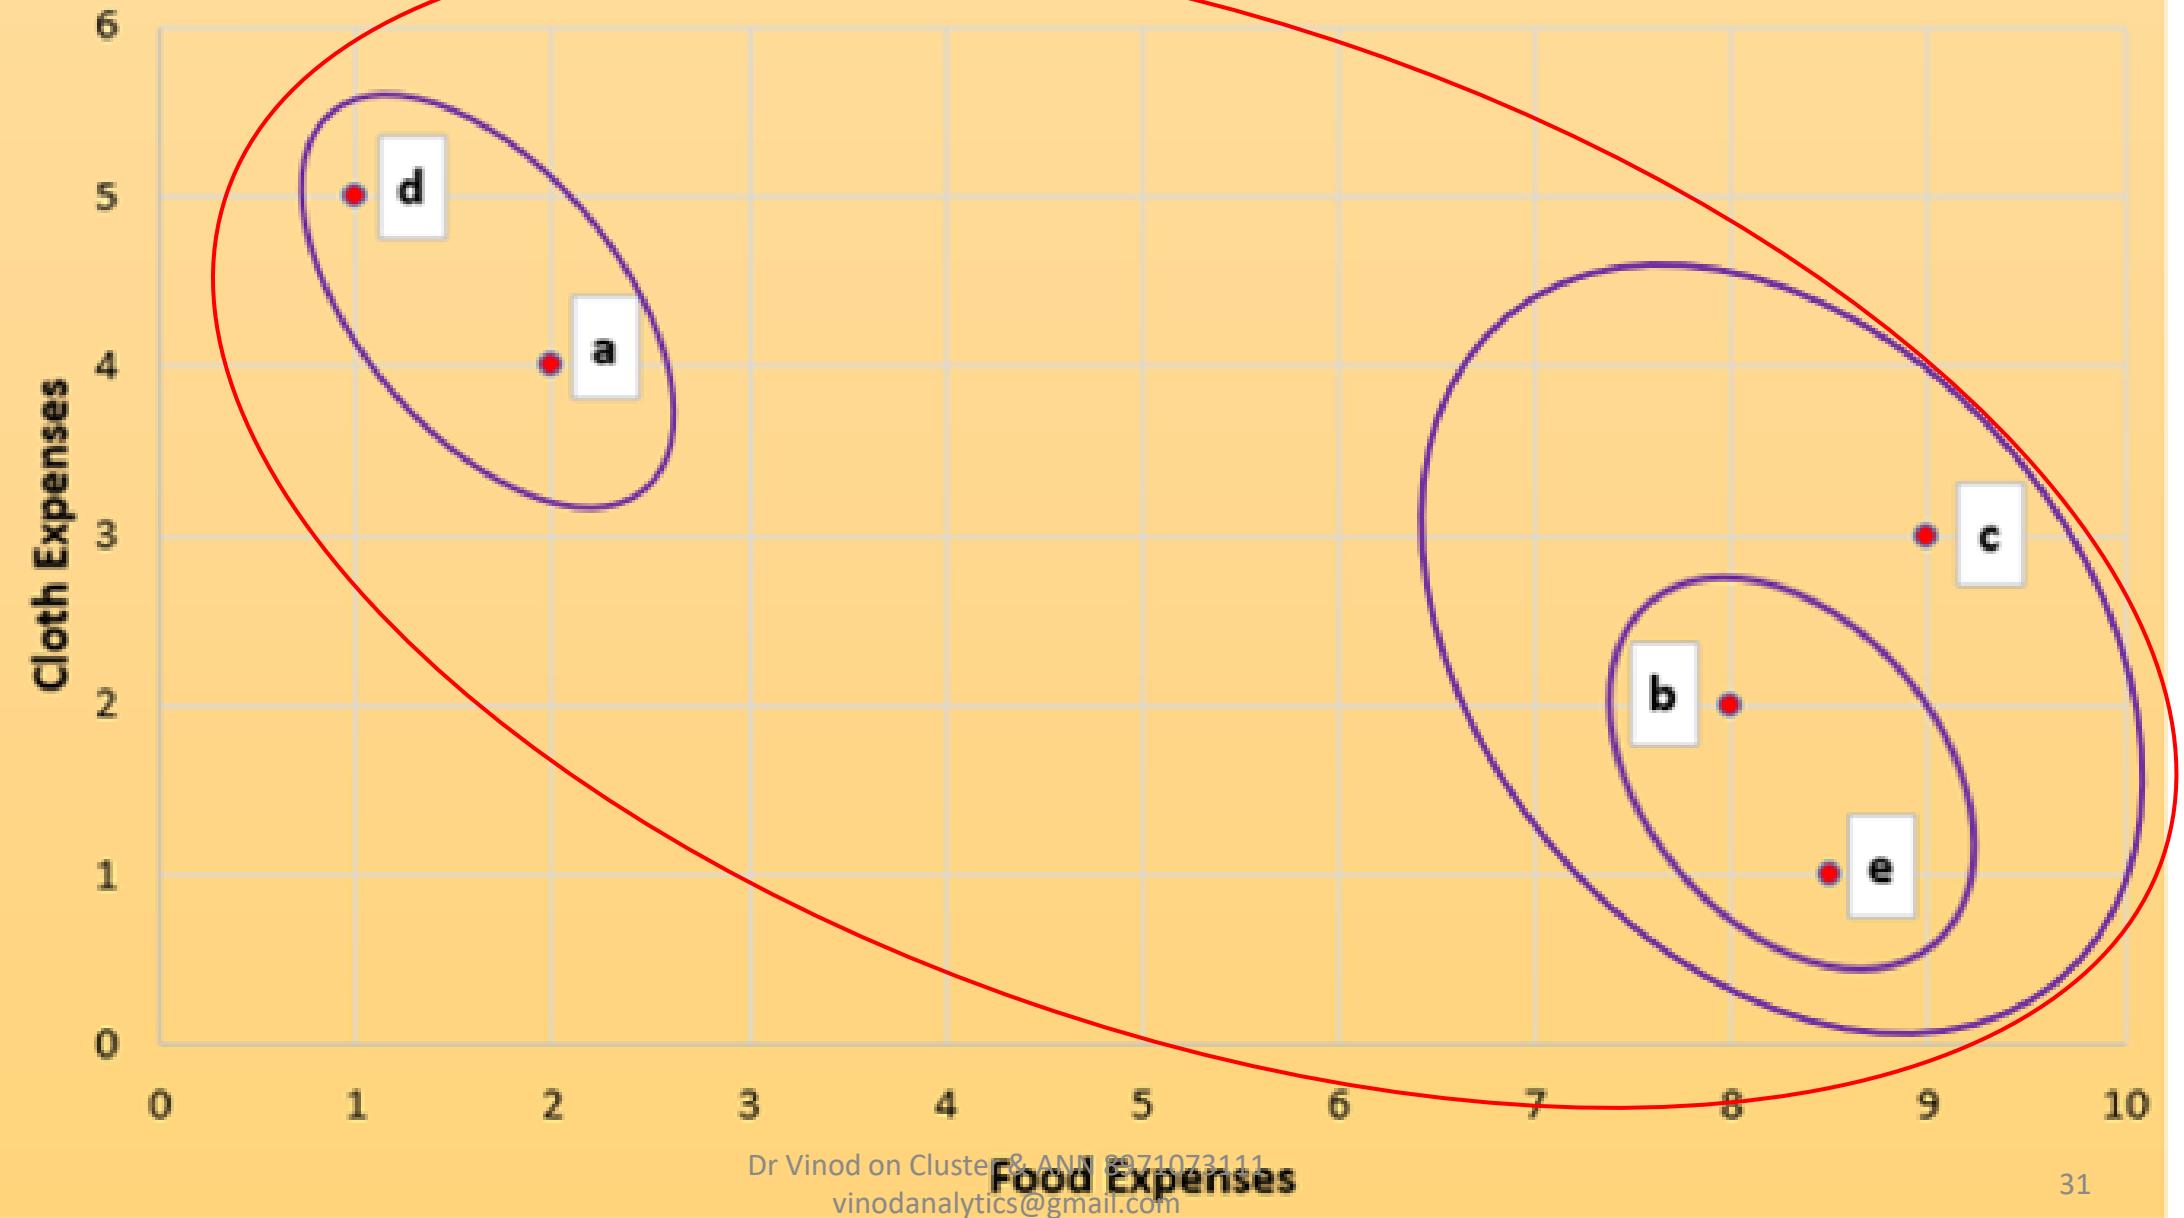
$$D(bec, ad) = \min\{6.32, 7.62, 7.16, 8.50, 7.07, 8.25\}$$



Final round: Distance between a d and b e c



Scatter Plot



Single Linkage: Nearest Neighbor

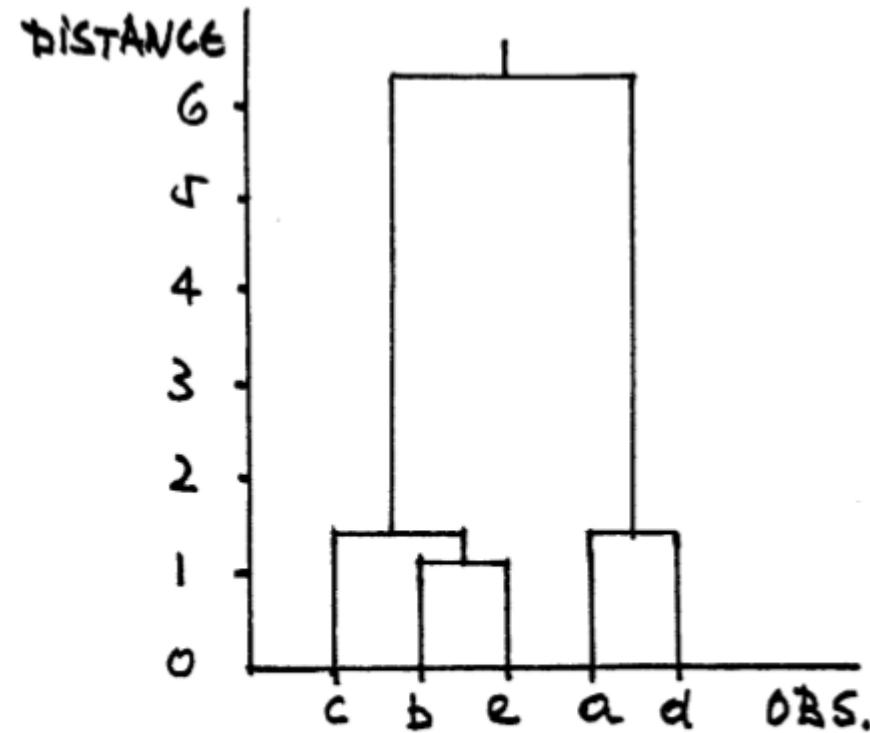


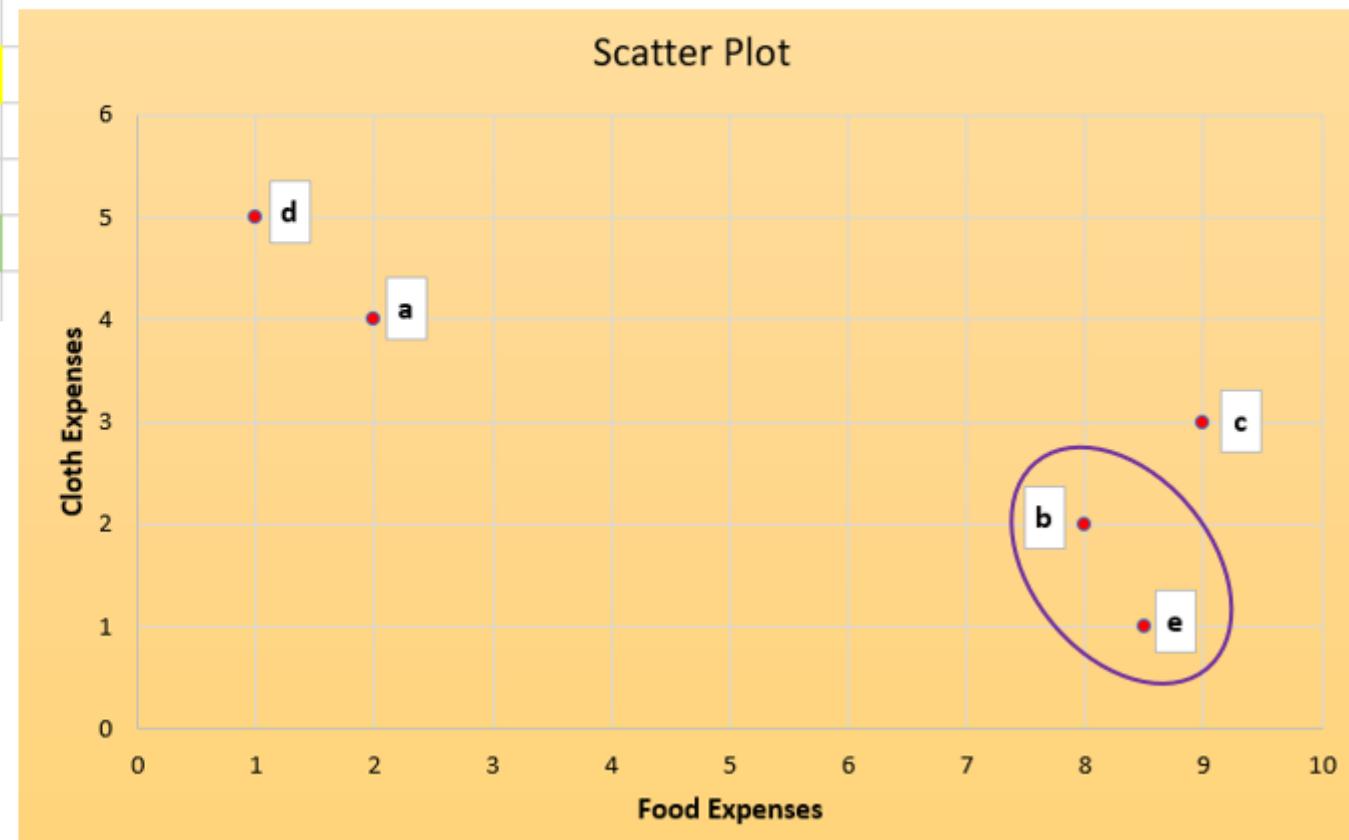
Figure 15.8
Nearest neighbor method, dendrogram

Complete Linkage: Farthest Neighbor

Stage 1: b and e can be clubbed

	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00

Now **b** and **e** are in one cluster. How their values will be treated for finding DIST vis-à-vis another point/s?



Stage 2: Complete Linkage

Points a and d can be clubbed (same as Single Linkage)

$$D(be, a) = \max\{D(b, a), D(e, a)\}$$

$$D(be, a) = \max\{6.32, 7.16\}$$

$$D(be, a) = 7.16$$

$$D(be, c) = \max\{D(b, c), D(e, c)\}$$

$$D(be, c) = \max\{1.41, 2.06\}$$

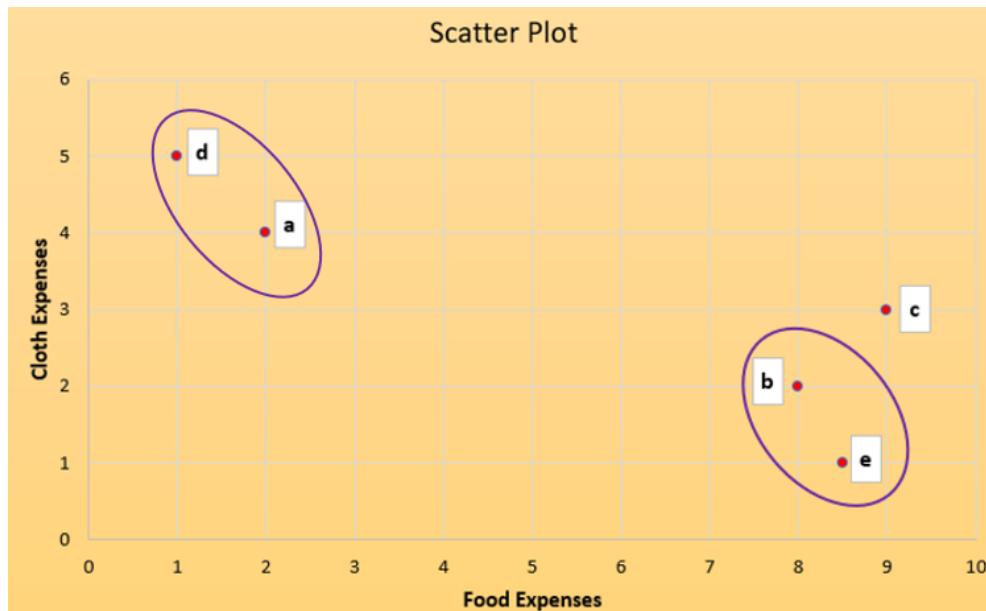
$$D(be, c) = 2.06$$

$$D(be, d) = \max\{D(b, d), D(e, d)\}$$

$$D(be, d) = \max\{7.62, 8.50\}$$

$$D(be, d) = 8.50$$

	b e	a	c	d
b e	0	7.16	2.06	8.5
a		0	7.07	1.41
c			0	8.25
d				0



Stage 3: Complete Linkage

Point c can be clubbed with b e (same as Single Linkage)

	b e	a	c	d
b e	0	7.16	2.06	8.5
a		0	7.07	1.41
c			0	8.25
d				0

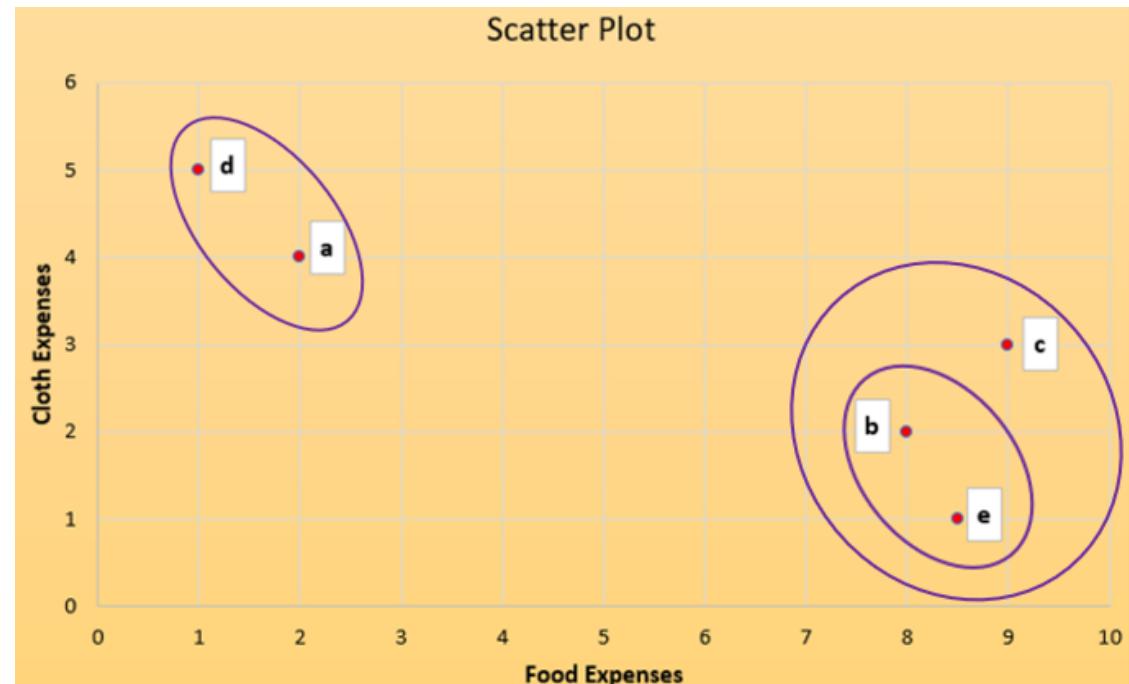
Only D(b e, ad) need to be calculated
Remaining are available in above table

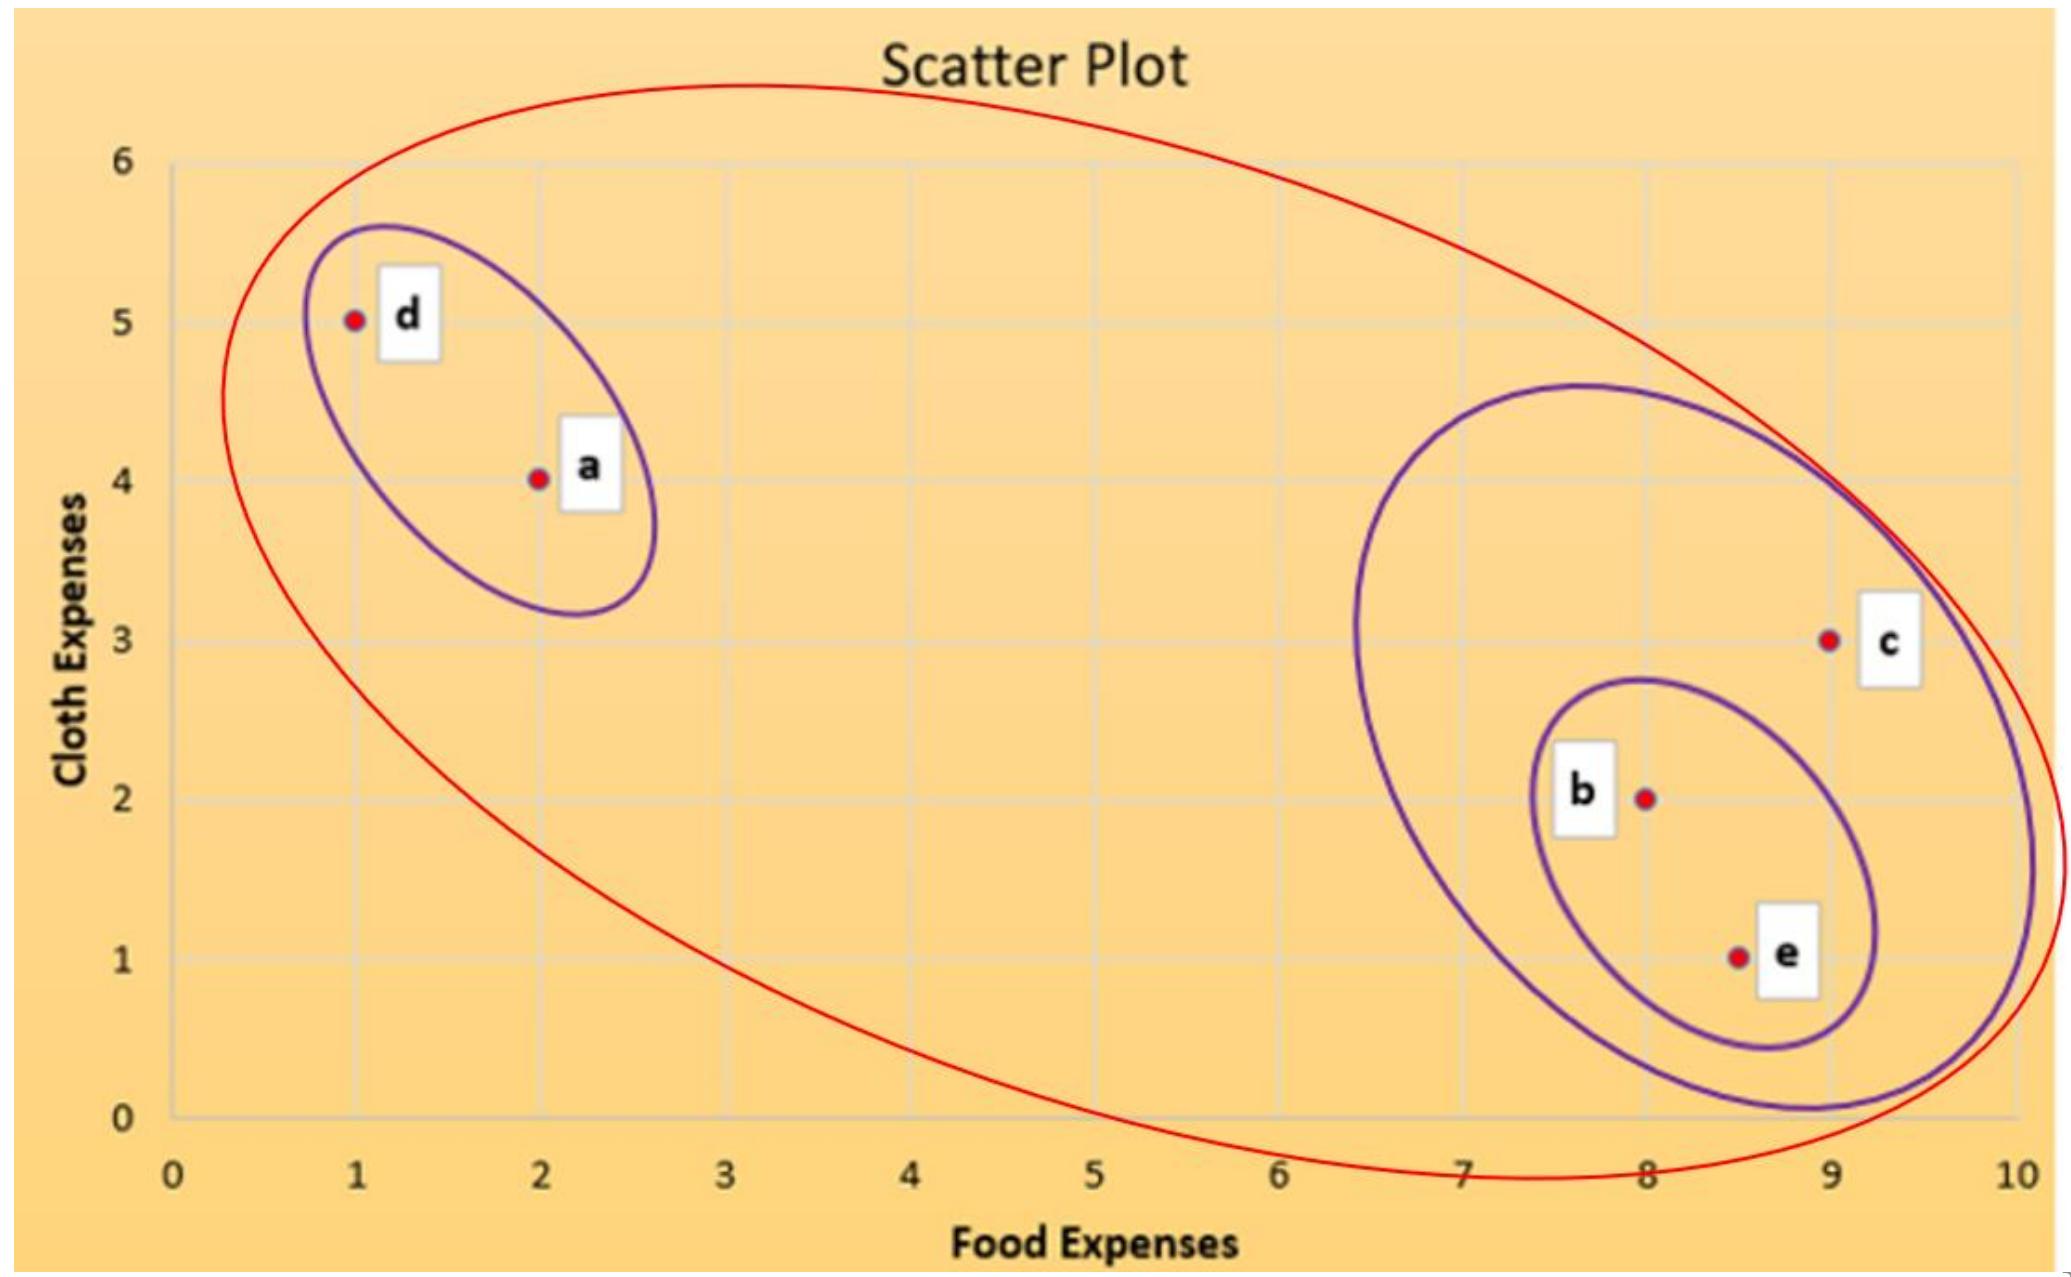
	b e	a d	c
b e	0	8.5	2.06
a d		0	8.25
c			0

$$D(b e, ad) = \max\{D(b e, a), D(b e, d)\}$$

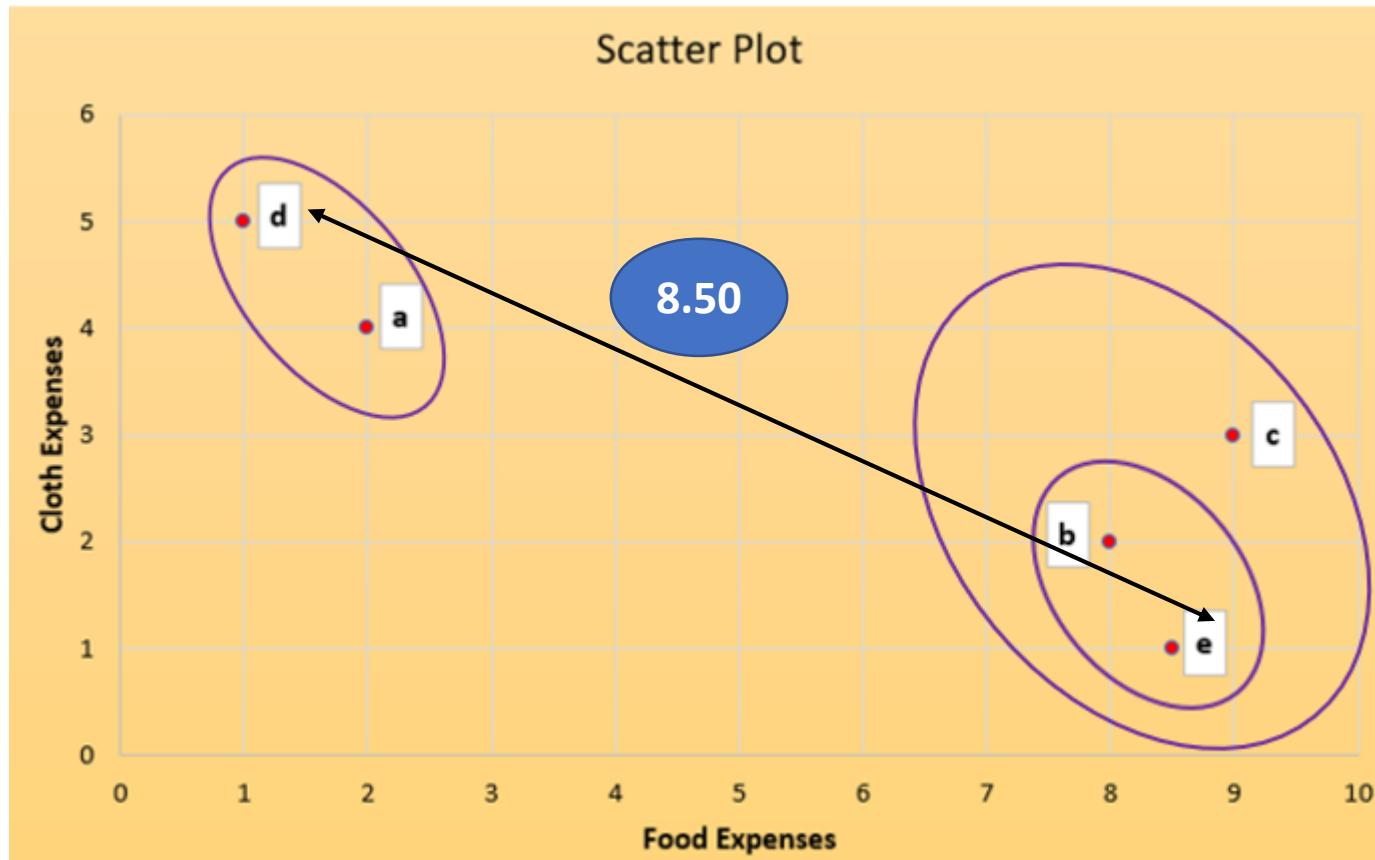
$$D(b e, ad) = \max\{7.16, 8.50\}$$

$$D(b e, ad) = 8.50$$





Final round: Distance between a d and b e c



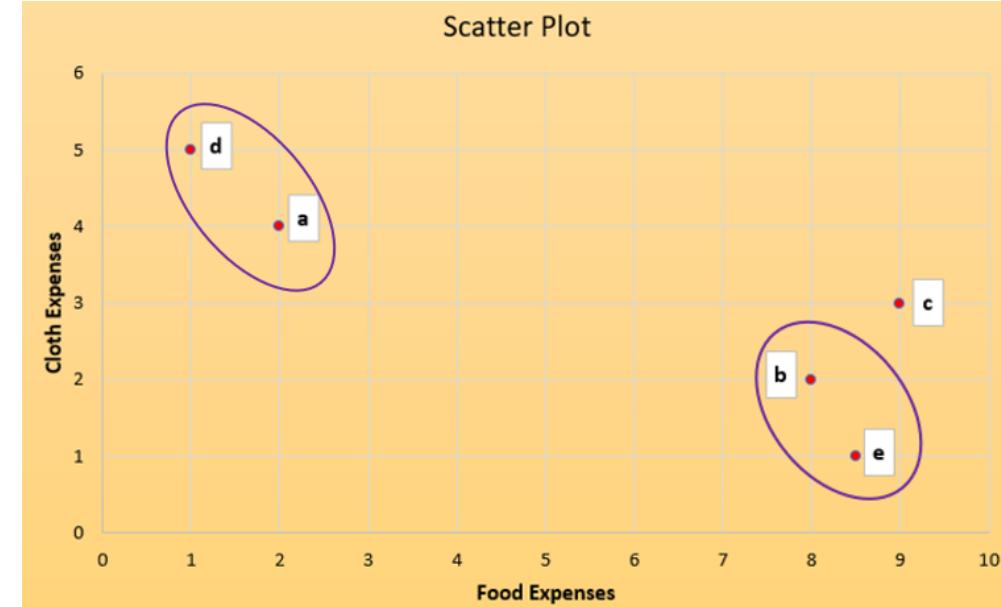
$$D(bec, ad) = \max\{D(b, a), D(b, d), D(e, a), D(e, d), D(c, a), D(c, d)\}$$

$$D(bec, ad) = \max\{6.32, 7.62, 7.16, 8.50, 7.07, 8.25\}$$

Average Linkage: Compromise (Min Max)

Stage 1 & 2

	a	b	c	d	e
a	0.00	6.32	7.07	1.41	7.16
b		0.00	1.41	7.62	1.12
c			0.00	8.25	2.06
d				0.00	8.50
e					0.00



First line only
to be calculated
(as averages),
Rest from
upper table

	b e	a	c	d
b e	0	6.74	1.74	8.06
a		0	7.07	1.41
c			0	8.25
d				0

$$D(be, a) = \{D(b, a) + D(e, a)\}/2$$

$$D(be, a) = \frac{6.32 + 7.16}{2} = \mathbf{6.74}$$

$$D(be, d) = \{D(b, d) + D(e, d)\}/2$$

$$D(be, d) = \frac{7.62 + 8.50}{2} = \mathbf{8.06}$$

$$D(be, c) = \{D(b, c) + D(e, c)\}/2$$

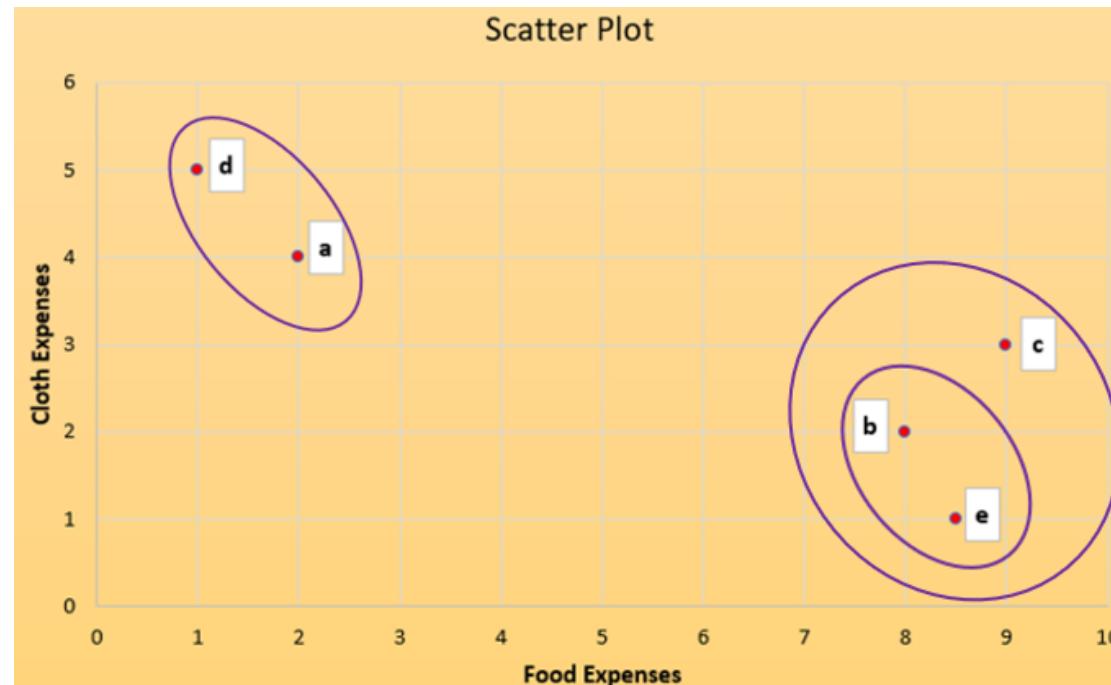
$$D(be, c) = \frac{1.41 + 2.06}{2} = \mathbf{1.74}$$

Stage 3

	b e	a	c	d
b e	0	6.74	1.74	8.06
a		0	7.07	1.41
c			0	8.25
d				0

	b e	a d	c
b e	0	7.4	1.74
a d		0	7.66
c			0

1.74 is from upper table, rest two were calculated



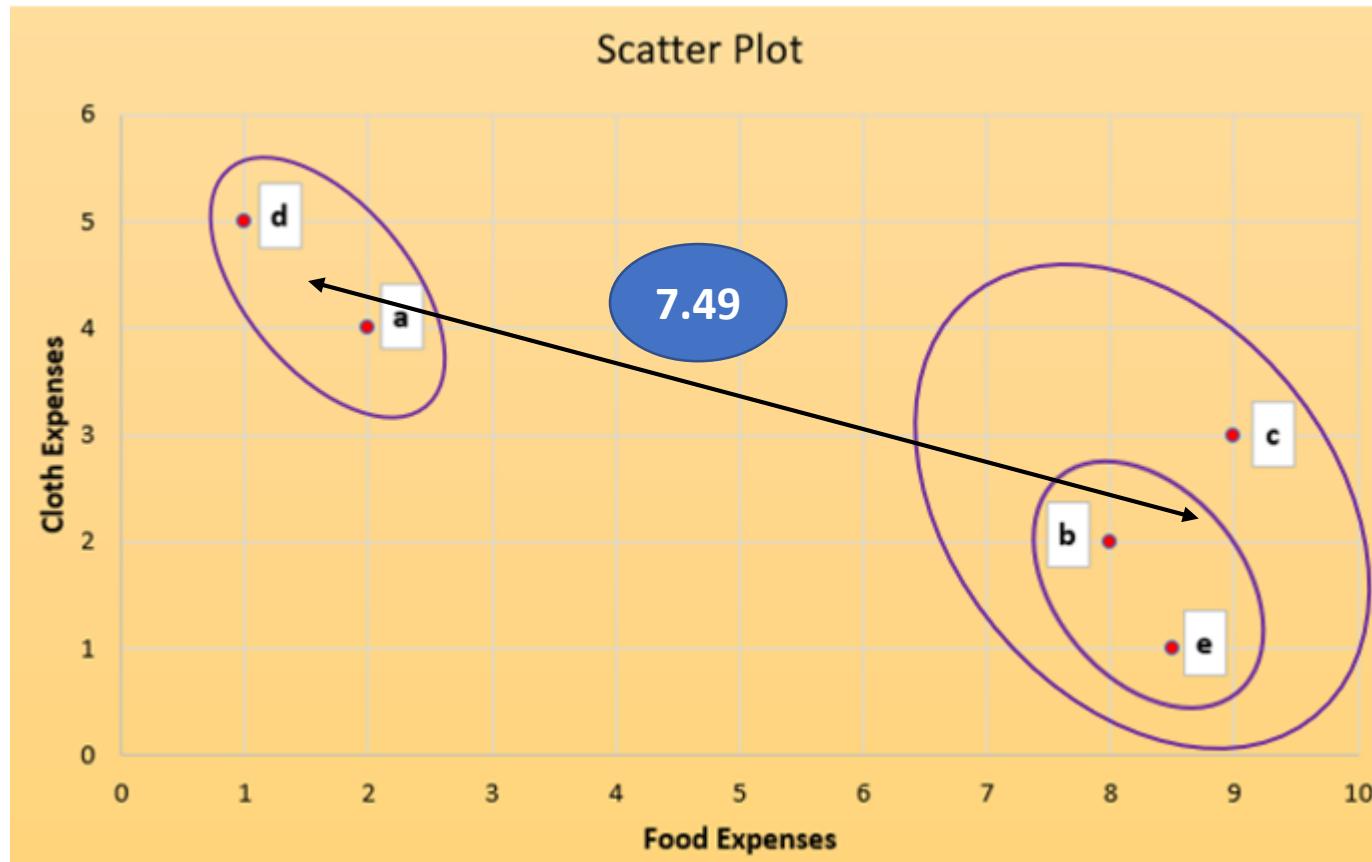
$$D(be, ad) = \{D(be, a) + D(be, d)\}/2$$

$$D(be, ad) = \frac{\{6.74 + 8.06\}}{2} = 7.4$$

$$D(ad, c) = \{D(a, c) + D(d, a)\}/2$$

$$D(ad, c) = \frac{\{7.07 + 8.25\}}{2} = 7.66$$

Final round: Distance between a d and b e c



$$D(bec, ad) = \text{average}\{D(b, a), D(b, d), D(e, a), D(e, d), D(c, a), D(c, d)\}$$

$$D(bec, ad) = \text{average}\{6.32, 7.62, 7.16, 8.50, 7.07, 8.25\} = 7.49$$

Ward Linkage: Finds centroids and then SSE

Distance Matrix

Ward's Method

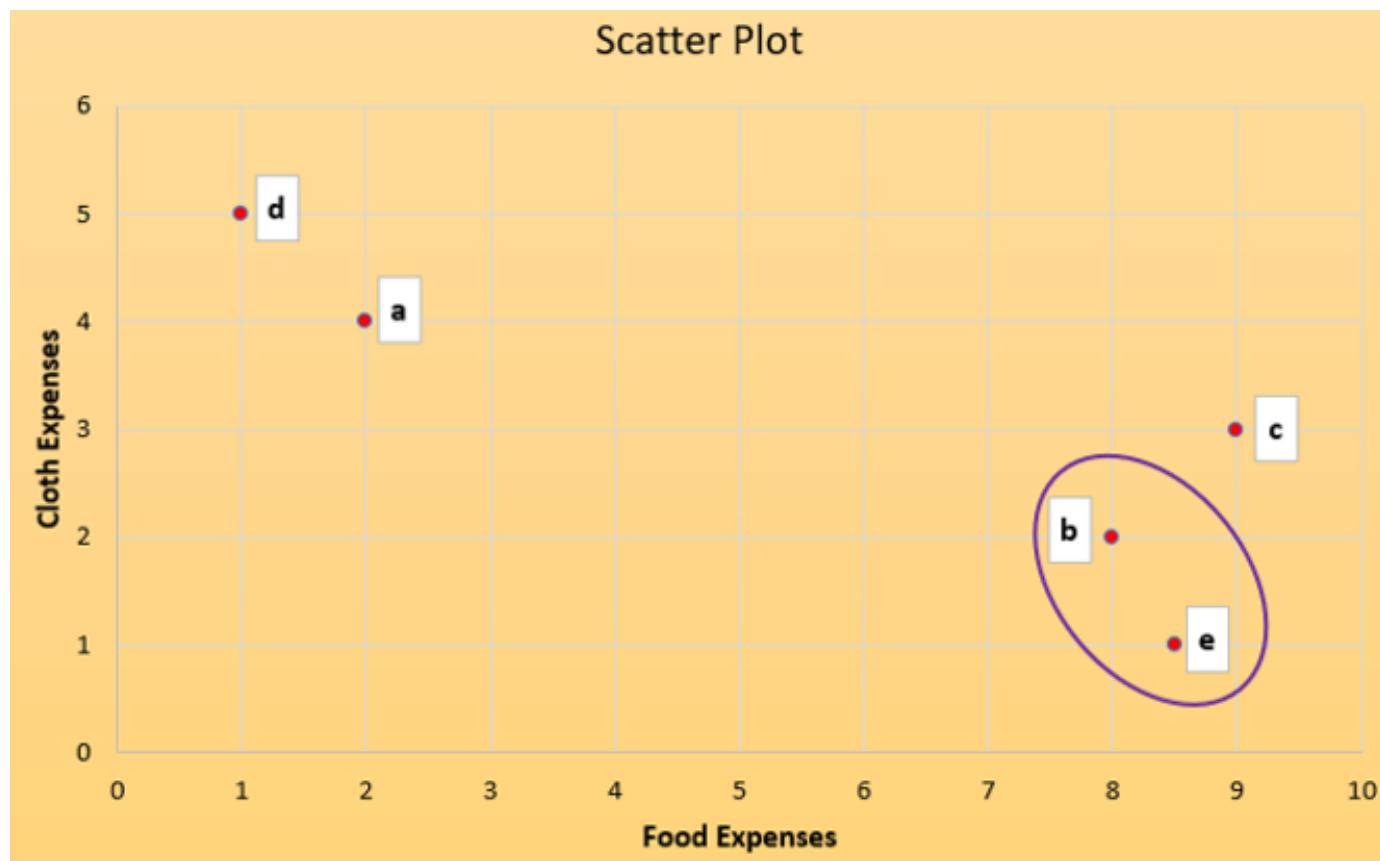
Ward's method means calculating the incremental sum of squares. *Half square Euclidean distance* is the only distance measure that can be used with this clustering method. Therefore, the distance measure is automatically set to *Half square Euclidean distance* when Ward's method is selected.

	a	b	c	d	e	stage 1
a	0.00	20.00	25.00	1.00	25.63	
b		0.00	1.00	29.00	0.63	
c			0.00	34.00	2.13	
d				0.00	30.13	
e					0.00	

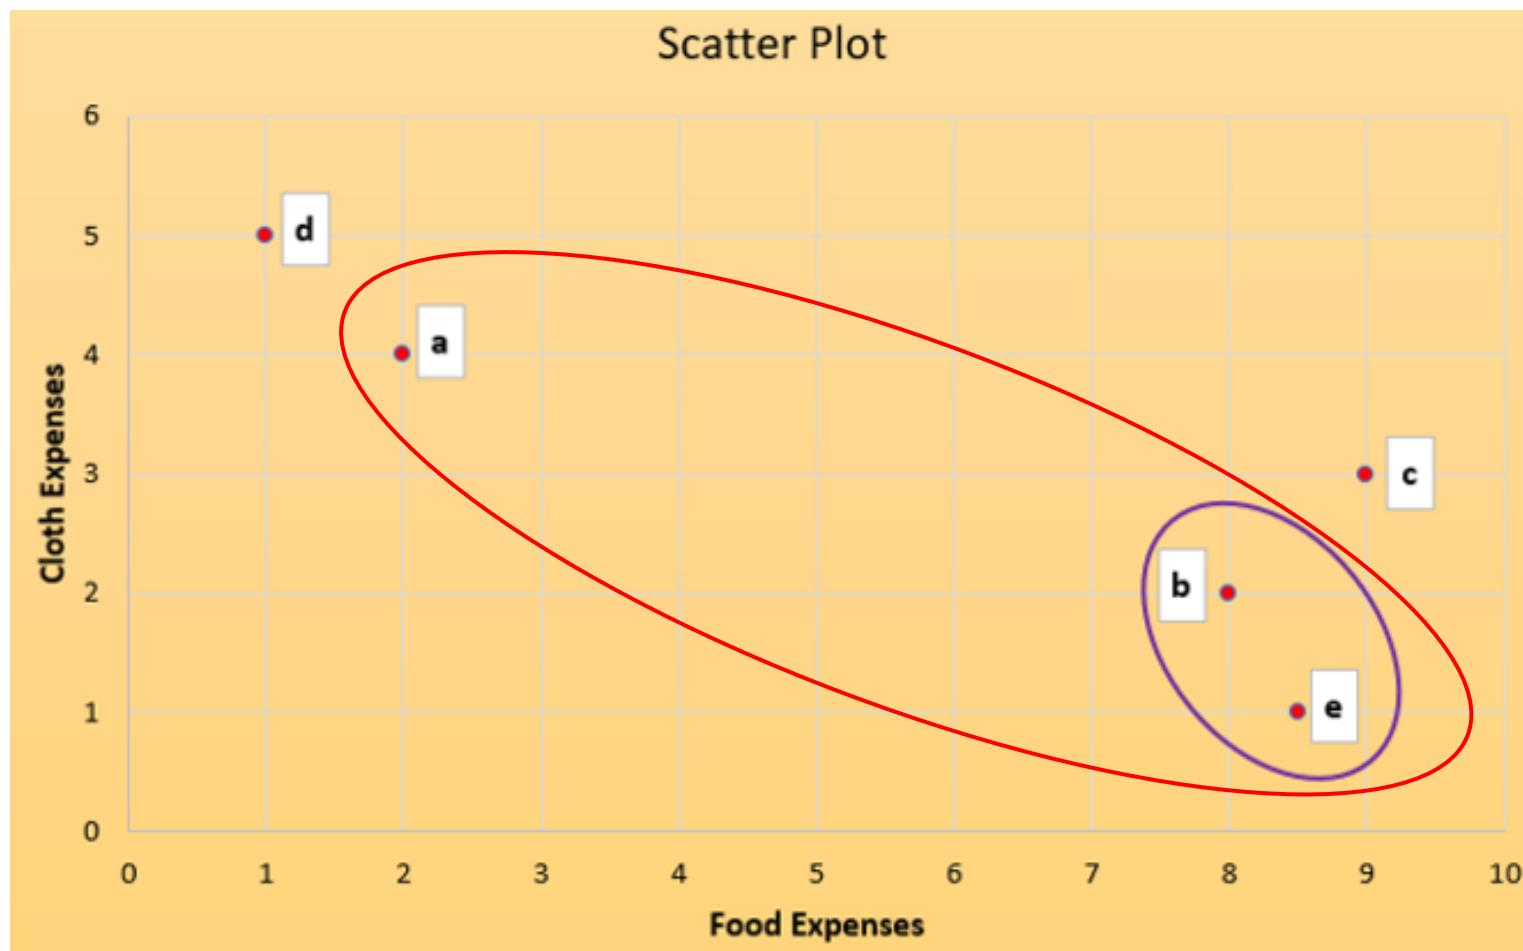
Euclidean
Distances

	b e	a	c	d
b e	0	6.74	1.74	8.06
a		0	7.07	1.41
c			0	8.25
d				0

Stage 1



Distance Between (be) and a

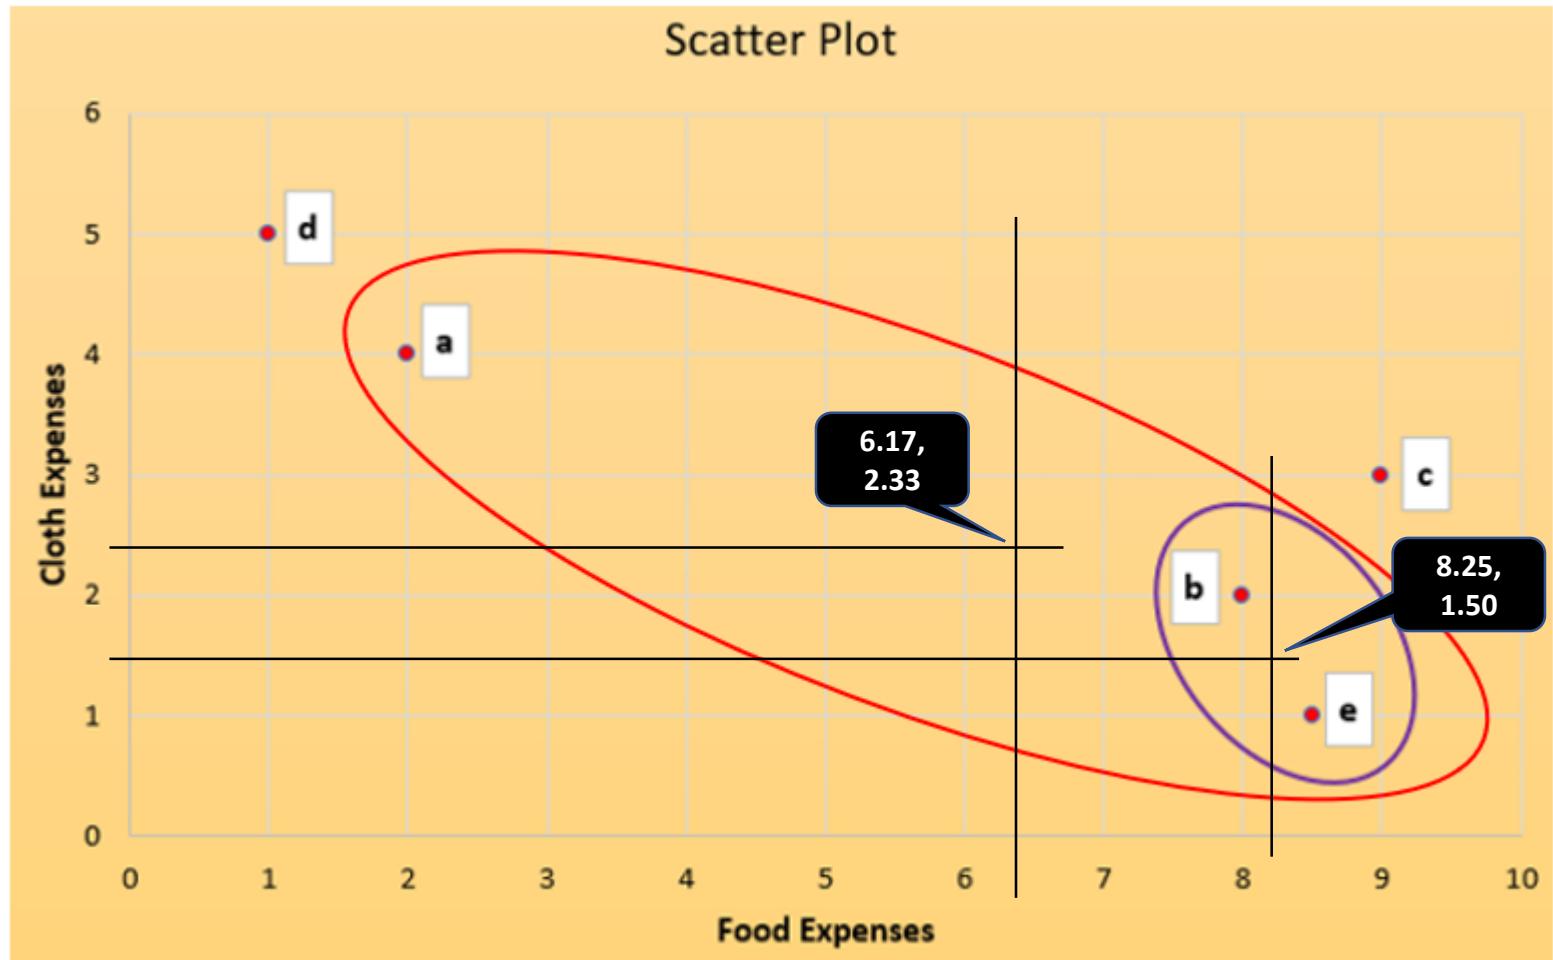


Distance Between (be) and a

First find centroids

Centroid for cluster (bea)		
	X	Y
b	8	2
e	8.5	1
a	2	4
	6.17	2.33

Centroid for cluster (be)		
	X	Y
b	8	2
e	8.5	1
	8.25	1.50



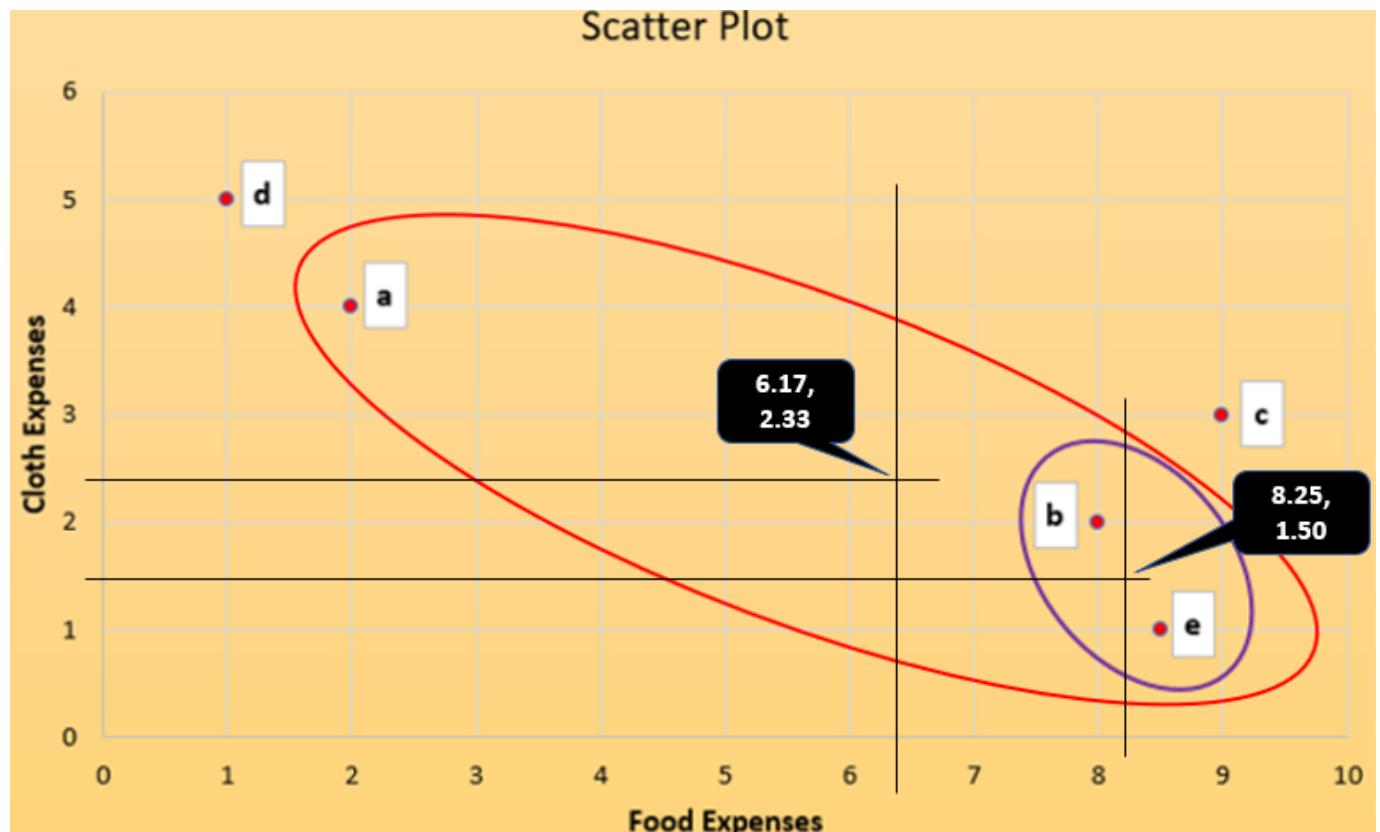
Distance Between (be) and a

Centroid for cluster (bea)		
	X	Y
b	8	2
e	8.5	1
a	2	4
	6.17	2.33
SSE(bea)	30.83	

Centroid for cluster (be)		
	X	Y
b	8	2
e	8.5	1
	8.25	1.50
SSE(be)	0.63	

$$dist_{(be,a)} = SSE_{(bea)} - SSE_{be} - SSE_a$$

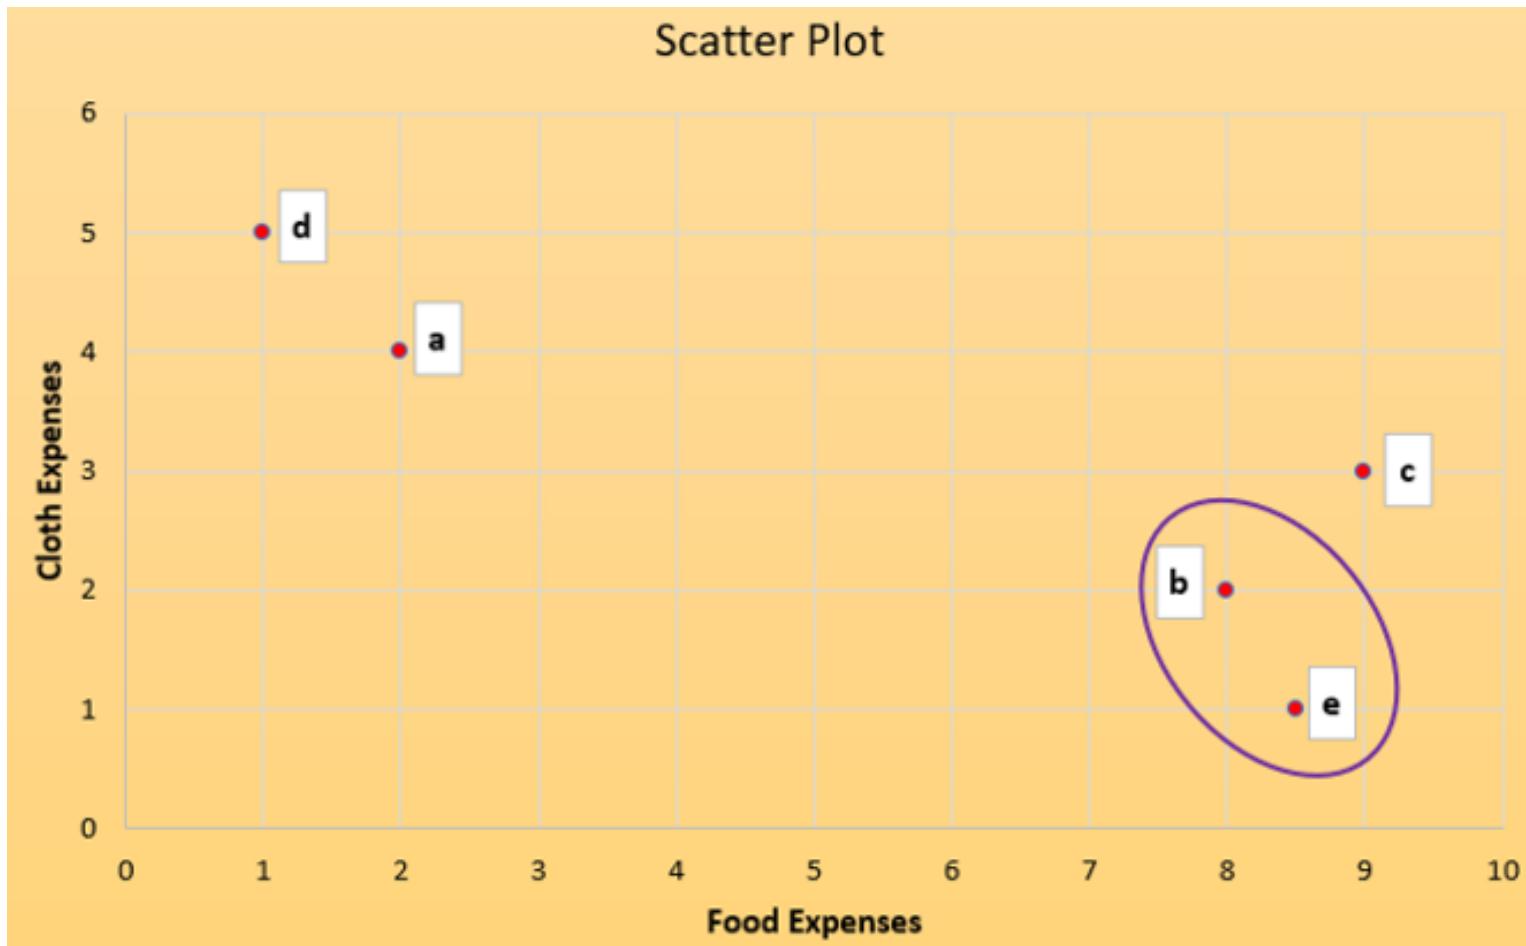
$$dist_{(be,a)} = 30.83 - 0.63 - 0 = 30.20$$



Distance Between (be) and d

First find centroids then find SSEs

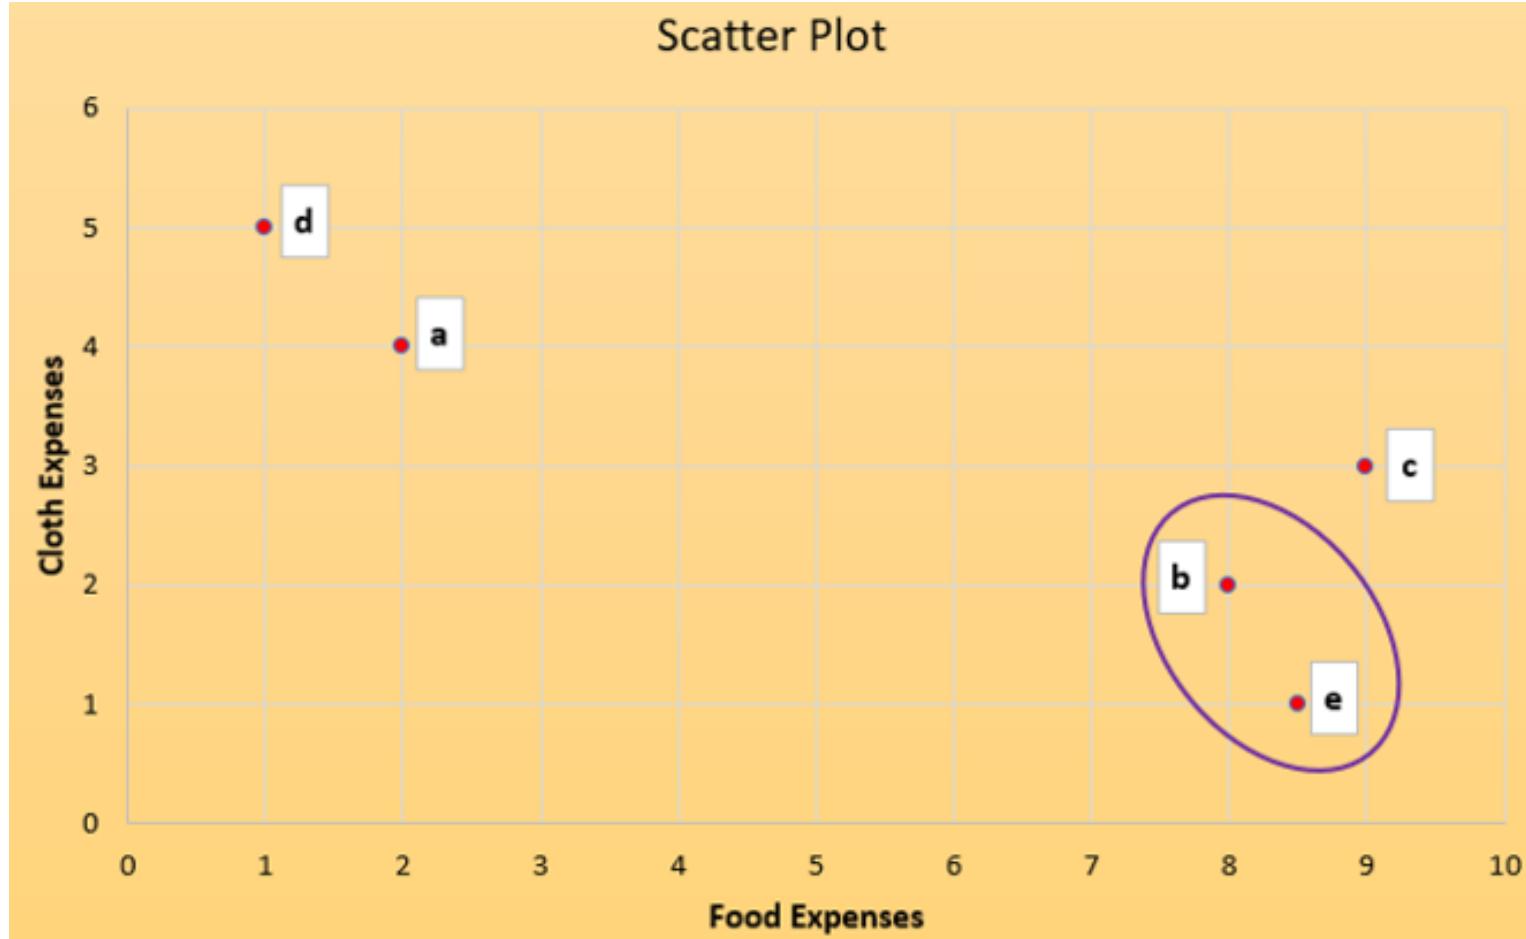
Centroid for cluster (bed)		
	X	Y
b	8	2
e	8.5	1
d	1	5
	5.83	2.67
SSE(bed)	43.83	
Centroid for cluster (be)		
	X	Y
b	8	2
e	8.5	1
	8.25	1.50
SSE(be)	0.63	
dist(be,d)	43.21	



Distance Between (be) and c

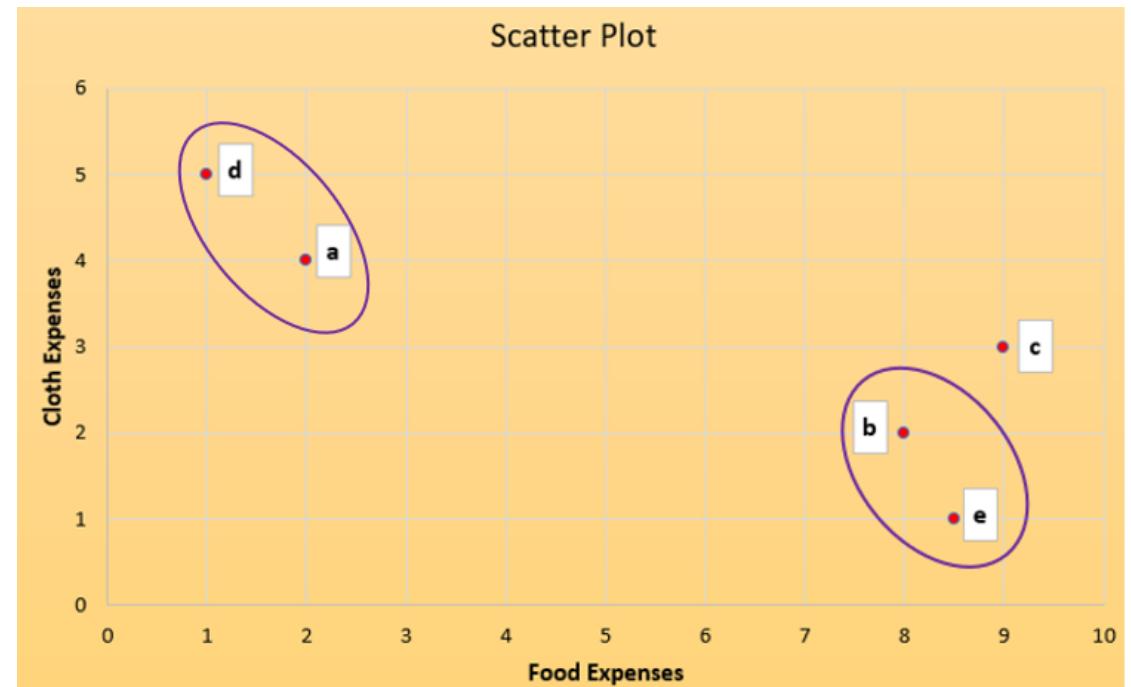
First find centroids then find SSEs

Centroid for cluster (bec)		
	X	Y
b	8	2
e	8.5	1
c	9	3
	8.50	2.00
SSE(bec)	2.50	
Centroid for cluster (be)		
	X	Y
b	8	2
e	8.5	1
	8.25	1.50
SSE(be)	0.63	
dist(be,c)	1.88	



Stage 2

	b e	a	c	d	stage 2
b e	0	30.21	1.88	43.21	
a		0	25.00	1.00	
c			0	34.00	
d				0	



Lets make it

```
In [1]: # Jesus is my Saviour!
```

```
In [2]: import os
```

```
In [3]: os.chdir('C:\\\\Users\\\\Dr Vinod\\\\Desktop\\\\WD_python')
```

```
In [4]: # our exported file will appear here
```

```
In [5]: import pandas as pd
```

```
In [6]: import numpy as np
```

```
In [7]: import matplotlib.pyplot as plt
```

```
In [8]: from sklearn.preprocessing import normalize
```

Data: Wholesale customers data

```
In [9]: data = pd.read_csv("C:/Users/Dr Vinod/Desktop/DataSets1/Wholesale customers data.csv")
```

```
In [10]: data = pd.DataFrame(data)
```

```
In [11]: data.head(3)
```

```
Out[11]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844

Scale the data

```
In [12]: data_scaled = normalize(data)
```

```
In [13]: data_scaled = pd.DataFrame(data_scaled, columns = data.columns)
```

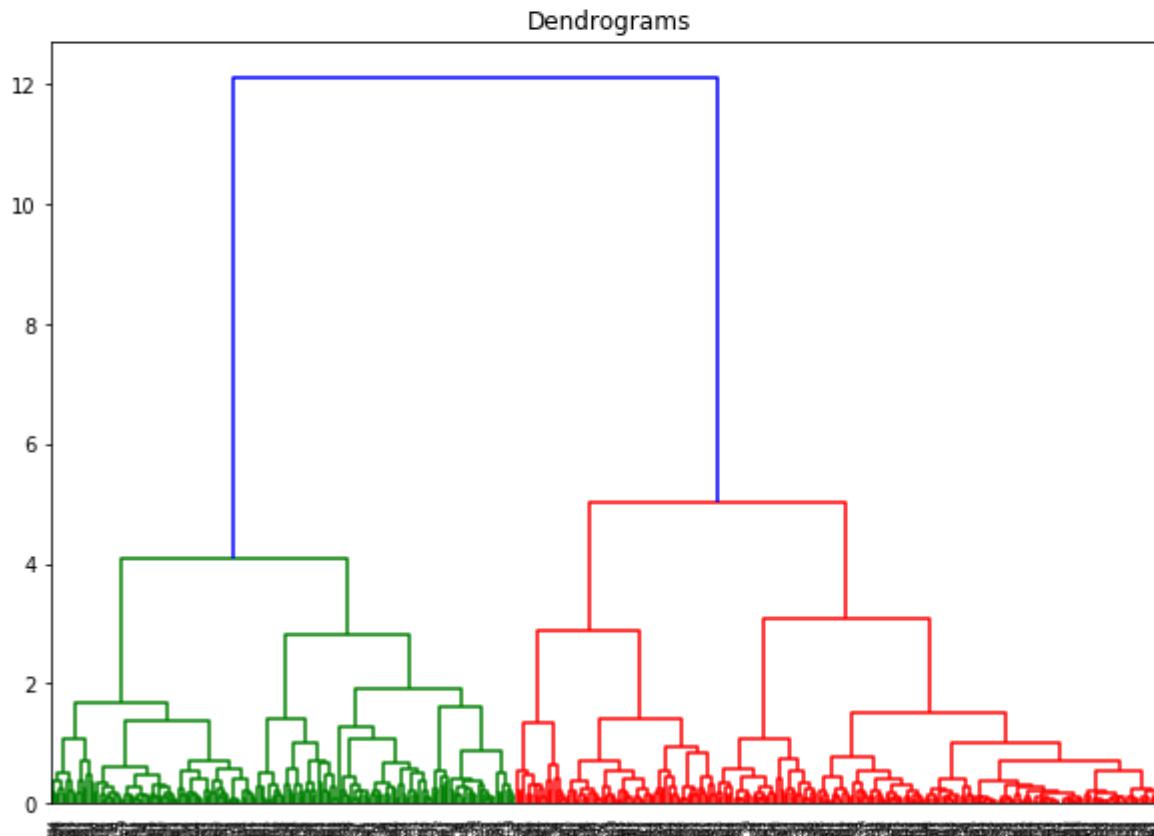
```
In [14]: data.head(data_scaled)
```

```
Traceback (most recent call last):
```

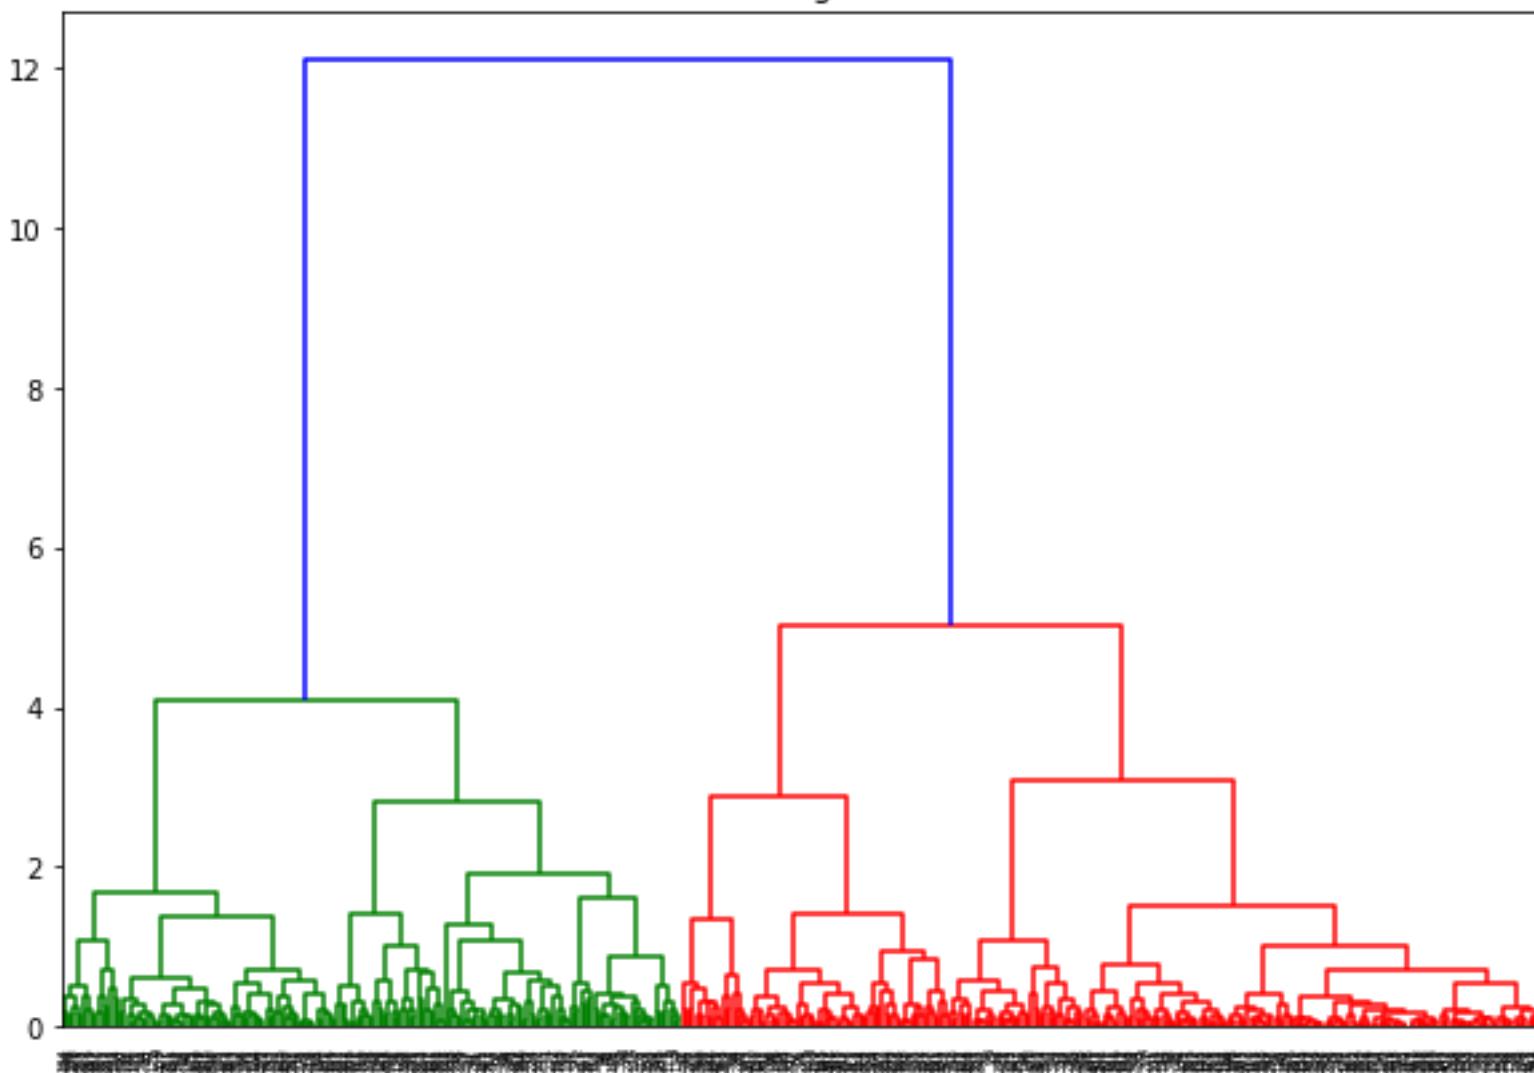
Dendrogram

```
In [15]: import scipy.cluster.hierarchy as shc
```

```
In [16]: plt.figure(figsize = (10,7))
....: plt.title("Dendrograms")
....: dend = shc.dendrogram(shc.linkage(data_scaled, method = 'ward'))
```

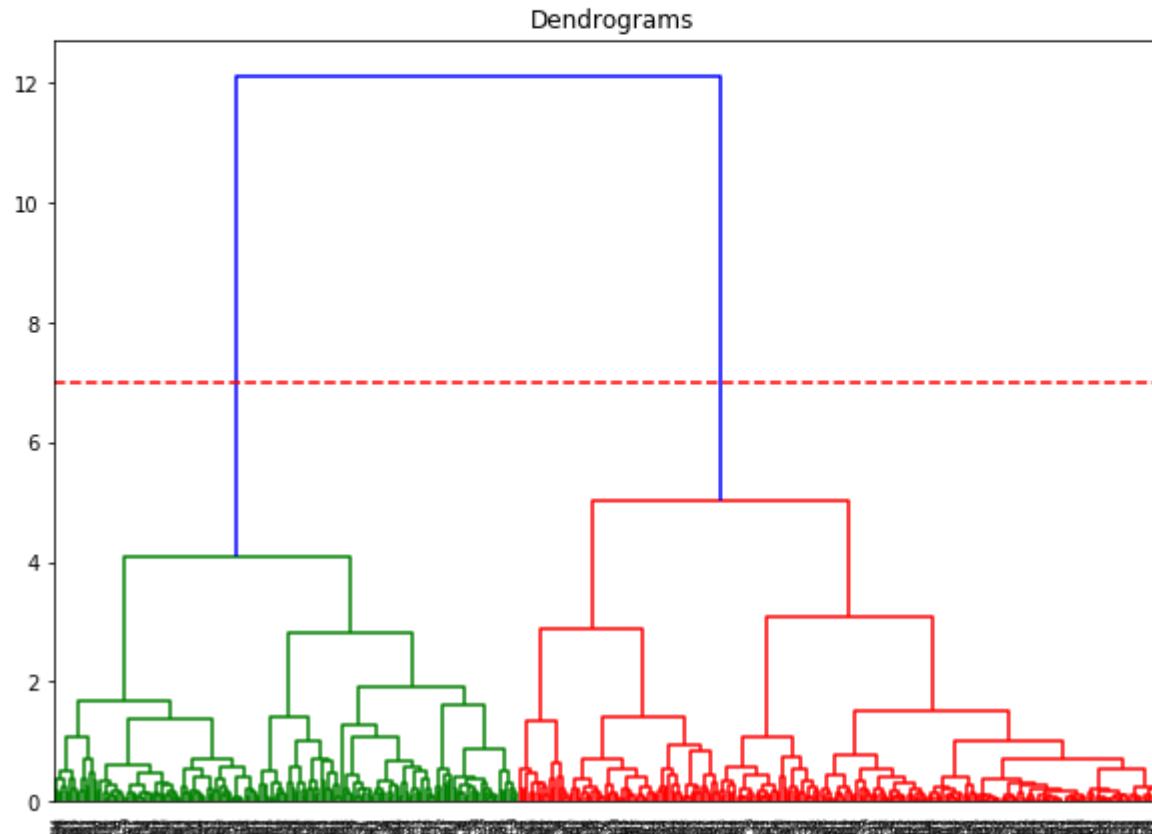


Dendograms

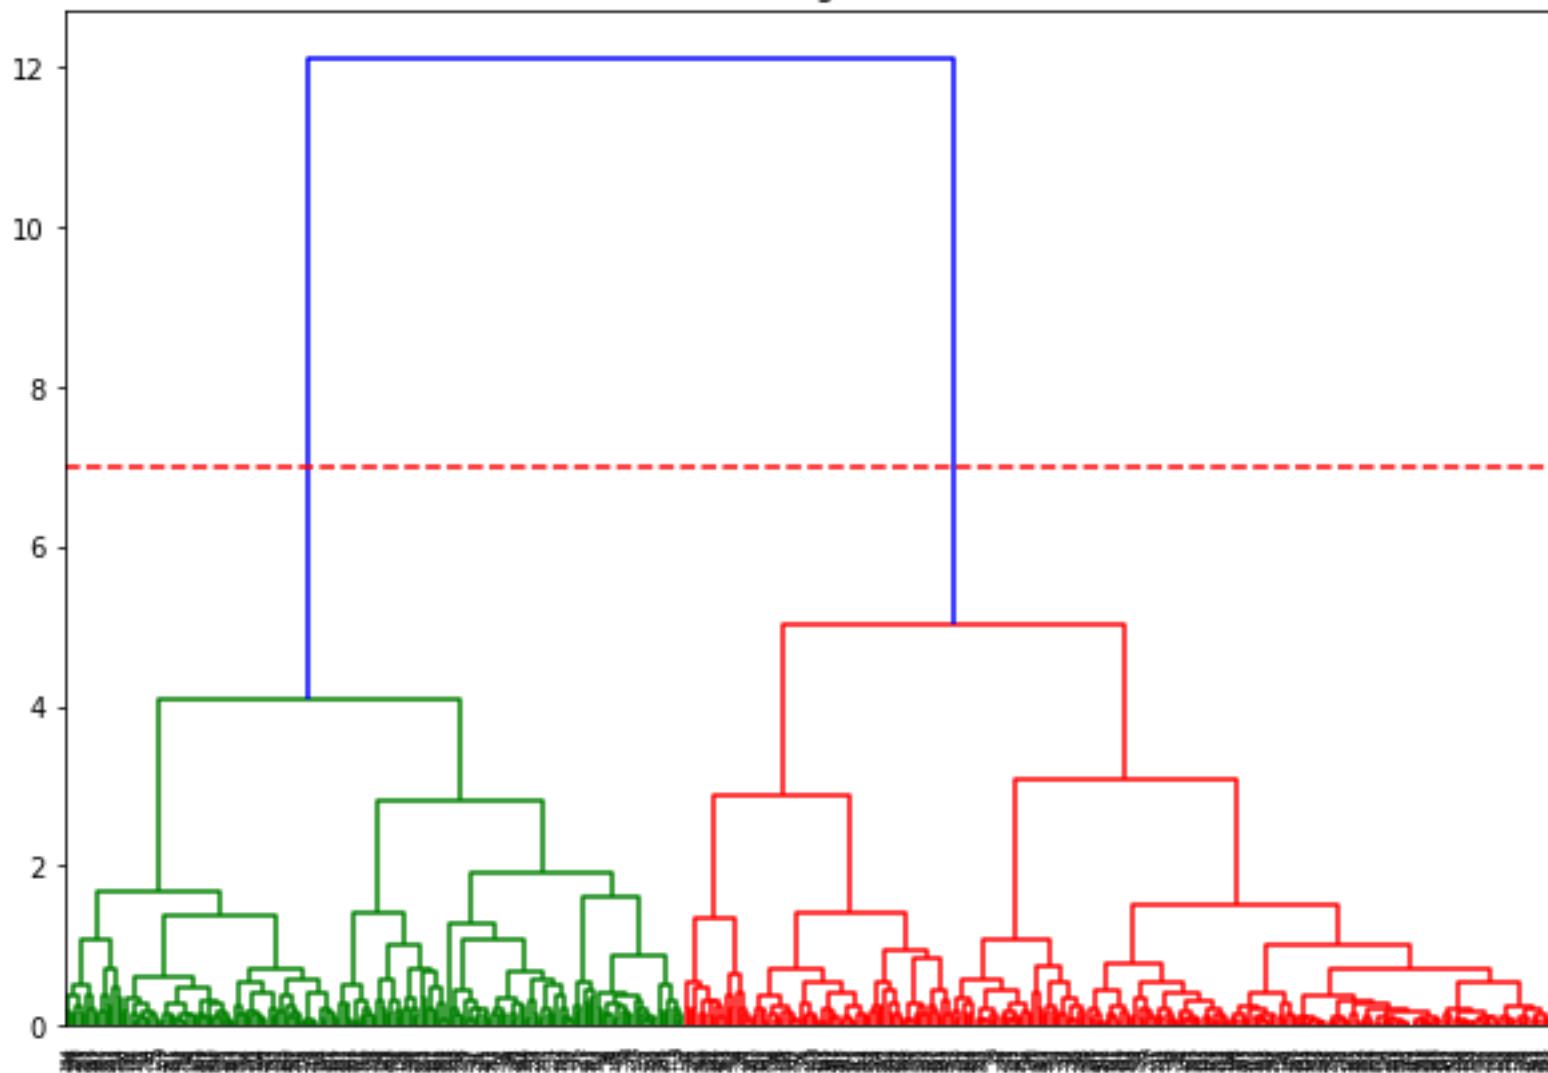


2 clusters

```
In [17]: plt.figure(figsize = (10,7))
.... plt.title("Dendograms")
.... dend = shc.dendrogram(shc.linkage(data_scaled, method = 'ward'))
.... plt.axhline(y=7, color = "r", linestyle = '--')
Out[17]: <matplotlib.lines.Line2D at 0x2242398ec50>
```



Dendrograms



Lets make clusters

```
In [18]: # _____ Lets apply Hierarchical clustering for 2 clusters

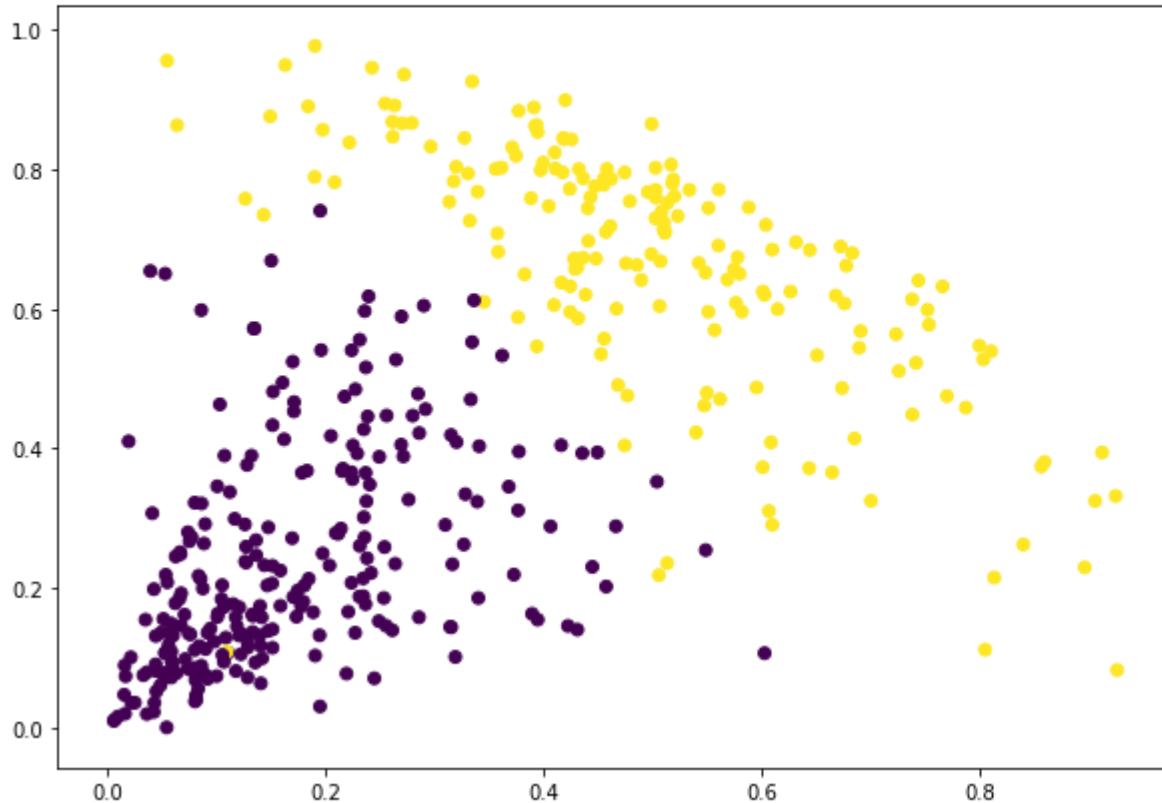
In [19]: from sklearn.cluster import AgglomerativeClustering

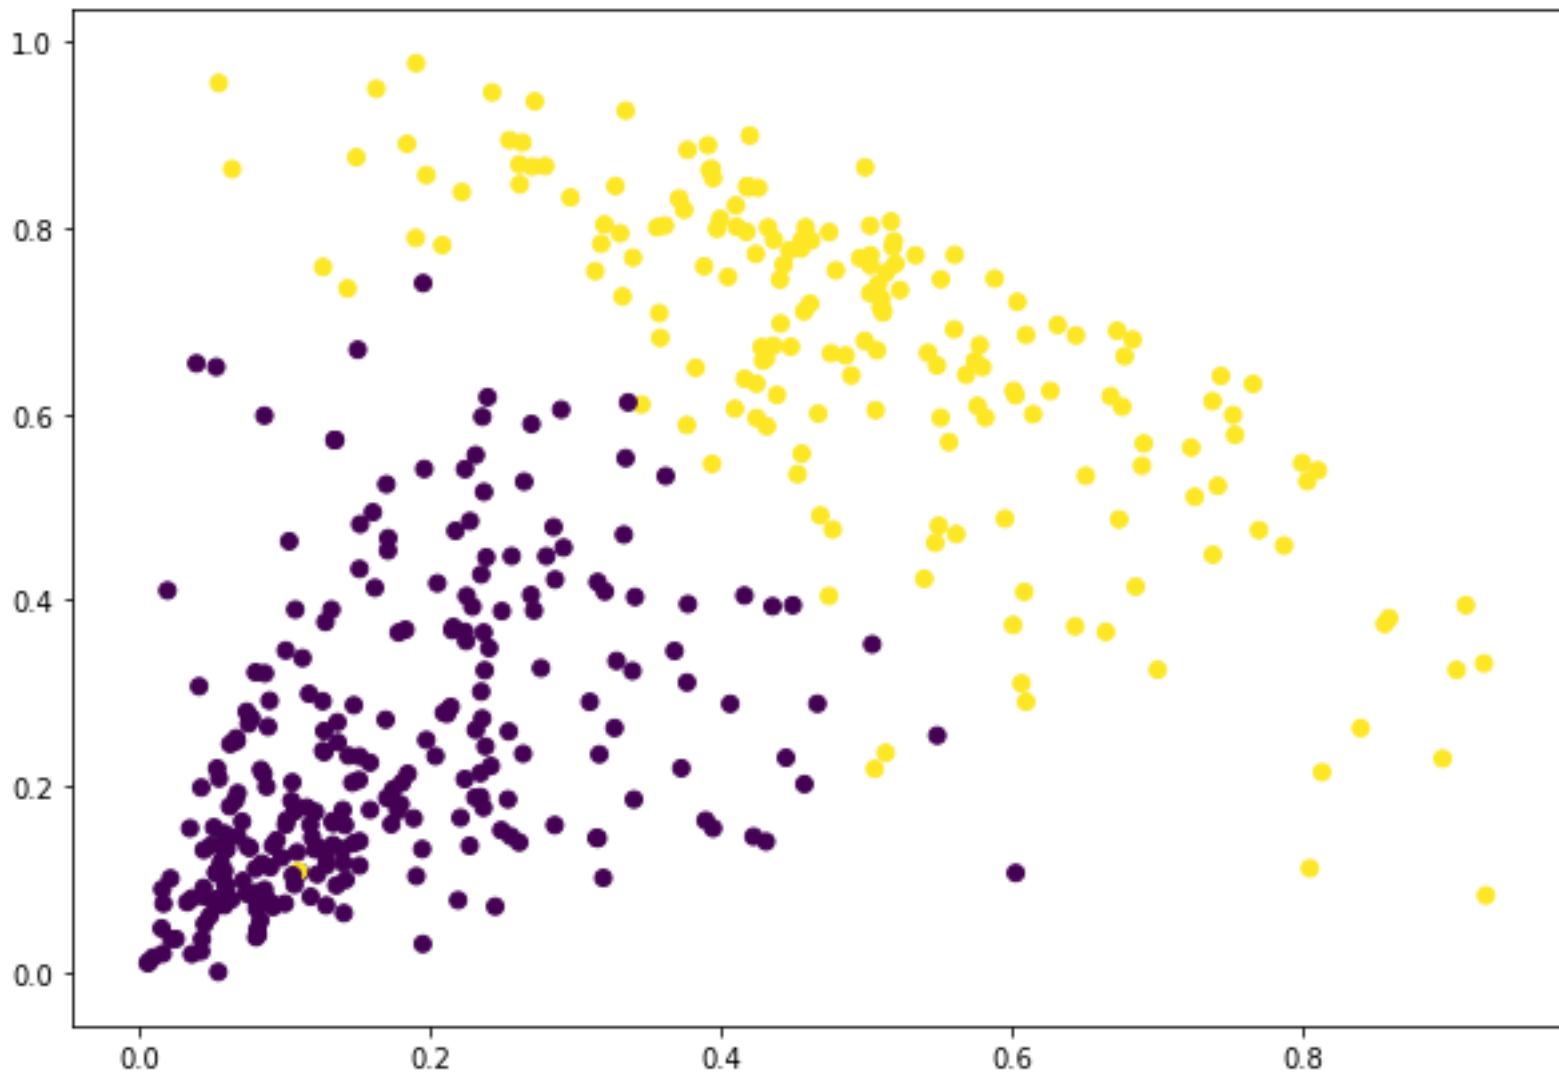
In [20]: cluster = AgglomerativeClustering(n_clusters = 2, affinity = 'euclidean',
linkage = 'ward')

In [21]: cluster.fit_predict(data_scaled)
Out[21]:
array([1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
       0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1,
       1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1,
       1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0,
       0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1,
       0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,
```

Plot

```
In [22]: plt.figure(figsize = (10,7))
....: plt.scatter(data_scaled['Milk'], data_scaled['Grocery'], c=cluster.labels_)
....:
Out[22]: <matplotlib.collections.PathCollection at 0x22423e25898>
```





K-means

Libraries

```
In [1]: # importing needed libraries  
  
In [2]: import pandas as pd  
  
In [3]: import matplotlib.pyplot as plt  
  
In [4]: from sklearn import preprocessing  
  
In [5]: from sklearn.preprocessing import scale  
  
In [6]: from sklearn.preprocessing import StandardScaler  
  
In [7]: from sklearn.cluster import KMeans  
  
In [8]: import seaborn as sns
```

```
In [9]: df = pd.read_csv("C:/Users/Dr Vinod/Desktop/DataSets1/USAArrests.csv")
```

```
In [10]: df = pd.DataFrame(df)
```

df - DataFrame

Index	Unnamed: 0	Murder	Assault	Ur
0	Alabama	13.2	236	58
1	Alaska	10	263	48
2	Arizona	8.1	294	80
3	Arkansas	8.8	190	50
4	California	9	276	91
5	Colorado	7.9	204	78
6	Connecticut	3.3	110	77
7	Delaware	5.9	238	72
8	Florida	15.4	335	80
9	Georgia	17.4	211	60
.....				

Format Resize Background Column min Save and Close

USArrests Data

df - DataFrame

Index	Unnamed: 0	Murder	Assault	Ur
0	Alabama	13.2	236	58
1	Alaska	10	263	48
2	Arizona	8.1	294	80
3	Arkansas	8.8	190	50
4	California	9	276	91
5	Colorado	7.9	204	78
6	Connecticut	3.3	110	77
7	Delaware	5.9	238	72
8	Florida	15.4	335	80
9	Georgia	17.4	211	60

Format Resize Background Column min Save and Close Close

Bring state's name as row names

```
In [11]: # Naming the 1st Unnamed column as States
```

```
In [12]: df.rename(columns={'Unnamed: 0':'States'}, inplace=True)
```

```
In [13]: # Setting States column as index
```

```
In [14]: df.set_index('States', inplace=True)
```

```
In [15]: df
```

```
Out[15]:
```

States	Murder	Assault	UrbanPop	Rape
Alabama	13.2	236	58	21.2
Alaska	10.0	263	48	44.5
Arizona	8.1	294	80	31.0
Arkansas	8.8	190	50	19.5
California	9.0	276	91	40.6

Scaling of variables

```
In [16]: # Creating X variable from df without States
```

```
In [17]: X = df[['Murder', 'Assault', 'Rape', 'UrbanPop']]
```

```
In [18]: # Scaling X variable
```

```
In [19]: scaler = StandardScaler()
```

```
In [20]: X_scaled = scaler.fit_transform( X )
```

Elbow Plot

```
In [21]: # Plotting for optimum nos of clusters
```

```
In [22]: plt.figure(figsize=(10, 8))
```

```
Out[22]: <Figure size 720x576 with 0 Axes><Figure size 720x576 with 0 Axes>
```

```
In [23]: wcss = []
```

```
In [24]: #____below in one block
```

```
In [25]: for i in range(1, 11):
```

```
....:     kmeans = KMeans(n_clusters = i, init = 'random', random_state = 42)
```

```
....:     kmeans.fit(X_scaled)
```

```
....:     wcss.append(kmeans.inertia_)
```

```
....:
```

```
....:
```

```
In [26]: plt.plot(range(1, 11), wcss, 'bx-')
```

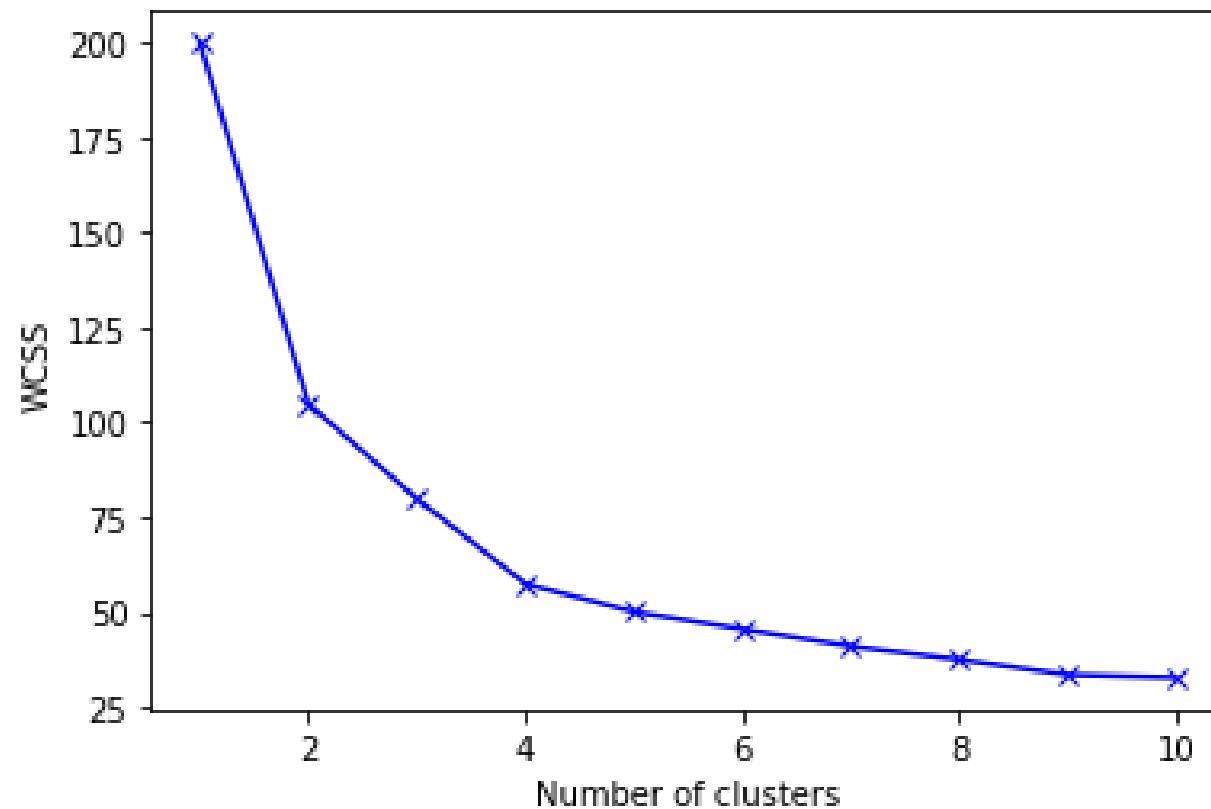
```
....: plt.title('The Elbow Method')
```

```
....: plt.xlabel('Number of clusters')
```

```
....: plt.ylabel('WCSS')
```

```
....: plt.show()
```

The Elbow Method



Build clusters

```
In [27]: # Running kmeans to our optimal number of clusters
```

```
In [28]: kmeans = KMeans(n_clusters= 4)
```

```
In [29]: clusters = kmeans.fit_predict(X_scaled)
```

```
In [30]: clusters
```

```
Out[30]:
```

```
array([2, 1, 1, 2, 1, 1, 0, 0, 1, 2, 0, 3, 1, 0, 3, 0, 3, 2, 3, 1, 0, 1,
       3, 2, 1, 3, 3, 1, 3, 0, 1, 1, 2, 3, 0, 0, 0, 0, 0, 2, 3, 2, 1, 0,
       3, 0, 0, 3, 3, 0])
```

Cluster Membership

```
In [31]: # Naming Clusters 1-4 instead of 0-3 and adding to dataframe
```

```
In [32]: y_kmeans1 = clusters + 1
```

```
In [33]: cluster = list(y_kmeans1)
```

```
In [34]: df['cluster'] = cluster
```

```
In [35]: df.head()
```

```
Out[35]:
```

	Murder	Assault	UrbanPop	Rape	cluster
States					
Alabama	13.2	236	58	21.2	3
Alaska	10.0	263	48	44.5	2
Arizona	8.1	294	80	31.0	2
Arkansas	8.8	190	50	19.5	3
California	9.0	276	91	40.6	2

In [36]: # States and counts in different clusters

In [37]: df[df['cluster']==1]

Out[37]:

States	Murder	Assault	UrbanPop	Rape	cluster
Connecticut	3.3	110	77	11.1	1
Delaware	5.9	238	72	15.8	1
Hawaii	5.3	46	83	20.2	1
Indiana	7.2	113	65	21.0	1
Kansas	6.0	115	66	18.0	1
Massachusetts	4.4	149	85	16.3	1
New Jersey	7.4	159	89	18.8	1
Ohio	7.3	120	75	21.4	1
Oklahoma	6.6	151	68	20.0	1
Oregon	4.9	159	67	29.3	1
Pennsylvania	6.3	106	72	14.9	1
Rhode Island	3.4	174	87	8.3	1
Utah	3.2	120	80	22.9	1
Virginia	8.5	156	63	20.7	1
Washington	4.0	145	73	26.2	1
Wyoming	6.8	161	60	15.6	1

In [38]: len(df[df['cluster']==1])

Out[38]: 16

Cluster-1

Cluster-2

```
In [39]: df[df['cluster']==2]
```

```
Out[39]:
```

States	Murder	Assault	UrbanPop	Rape	cluster
Alaska	10.0	263	48	44.5	2
Arizona	8.1	294	80	31.0	2
California	9.0	276	91	40.6	2
Colorado	7.9	204	78	38.7	2
Florida	15.4	335	80	31.9	2
Illinois	10.4	249	83	24.0	2
Maryland	11.3	300	67	27.8	2
Michigan	12.1	255	74	35.1	2
Missouri	9.0	178	70	28.2	2
Nevada	12.2	252	81	46.0	2
New Mexico	11.4	285	70	32.1	2
New York	11.1	254	86	26.1	2
Texas	12.7	201	80	25.5	2

```
In [40]: len(df[df['cluster']==2])
```

```
Out[40]: 13
```

Cluster-3

```
In [41]: df[df['cluster']==3]
```

```
Out[41]:
```

States	Murder	Assault	UrbanPop	Rape	cluster
Alabama	13.2	236	58	21.2	3
Arkansas	8.8	190	50	19.5	3
Georgia	17.4	211	60	25.8	3
Louisiana	15.4	249	66	22.2	3
Mississippi	16.1	259	44	17.1	3
North Carolina	13.0	337	45	16.1	3
South Carolina	14.4	279	48	22.5	3
Tennessee	13.2	188	59	26.9	3

```
In [42]: len(df[df['cluster']==3])
```

```
Out[42]: 8
```

Cluster-4

```
In [43]: df[df['cluster']==4]
```

```
Out[43]:
```

States	Murder	Assault	UrbanPop	Rape	cluster
Idaho	2.6	120	54	14.2	4
Iowa	2.2	56	57	11.3	4
Kentucky	9.7	109	52	16.3	4
Maine	2.1	83	51	7.8	4
Minnesota	2.7	72	66	14.9	4
Montana	6.0	109	53	16.4	4
Nebraska	4.3	102	62	16.5	4
New Hampshire	2.1	57	56	9.5	4
North Dakota	0.8	45	44	7.3	4
South Dakota	3.8	86	45	12.8	4
Vermont	2.2	48	32	11.2	4
West Virginia	5.7	81	39	9.3	4
Wisconsin	2.6	53	66	10.8	4

```
In [44]: len(df[df['cluster']==4])
```

```
Out[44]: 13
```

Cluster Profiling

```
In [45]: # Mean of clusters 1 to 4
```

```
In [46]: kmeans_mean_cluster = pd.DataFrame(round(df.groupby('cluster').mean(),1))
```

```
In [47]: kmeans_mean_cluster
```

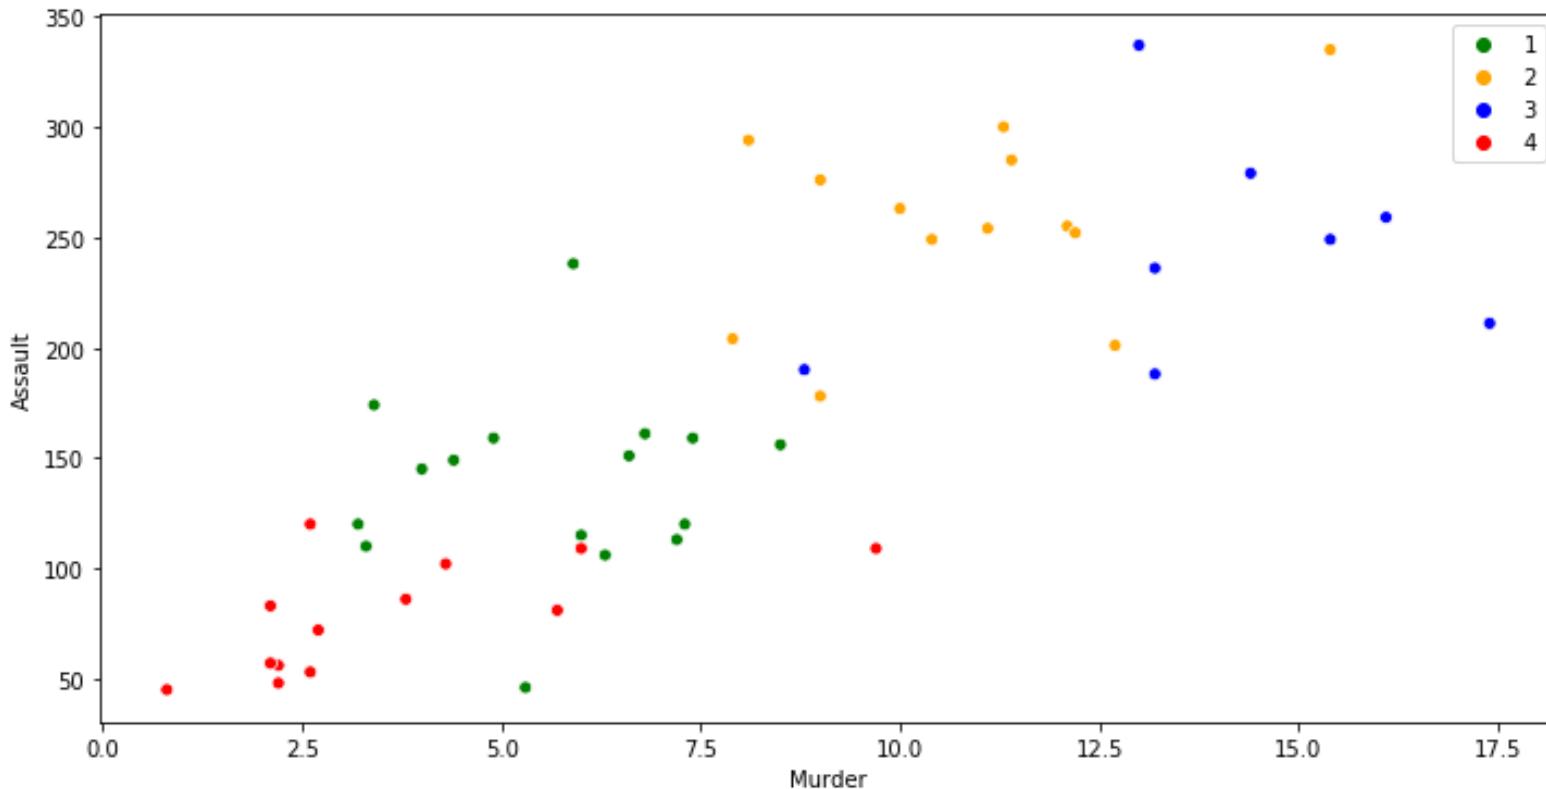
```
Out[47]:
```

cluster	Murder	Assault	UrbanPop	Rape
1	5.7	138.9	73.9	18.8
2	10.8	257.4	76.0	33.2
3	13.9	243.6	53.8	21.4
4	3.6	78.5	52.1	12.2

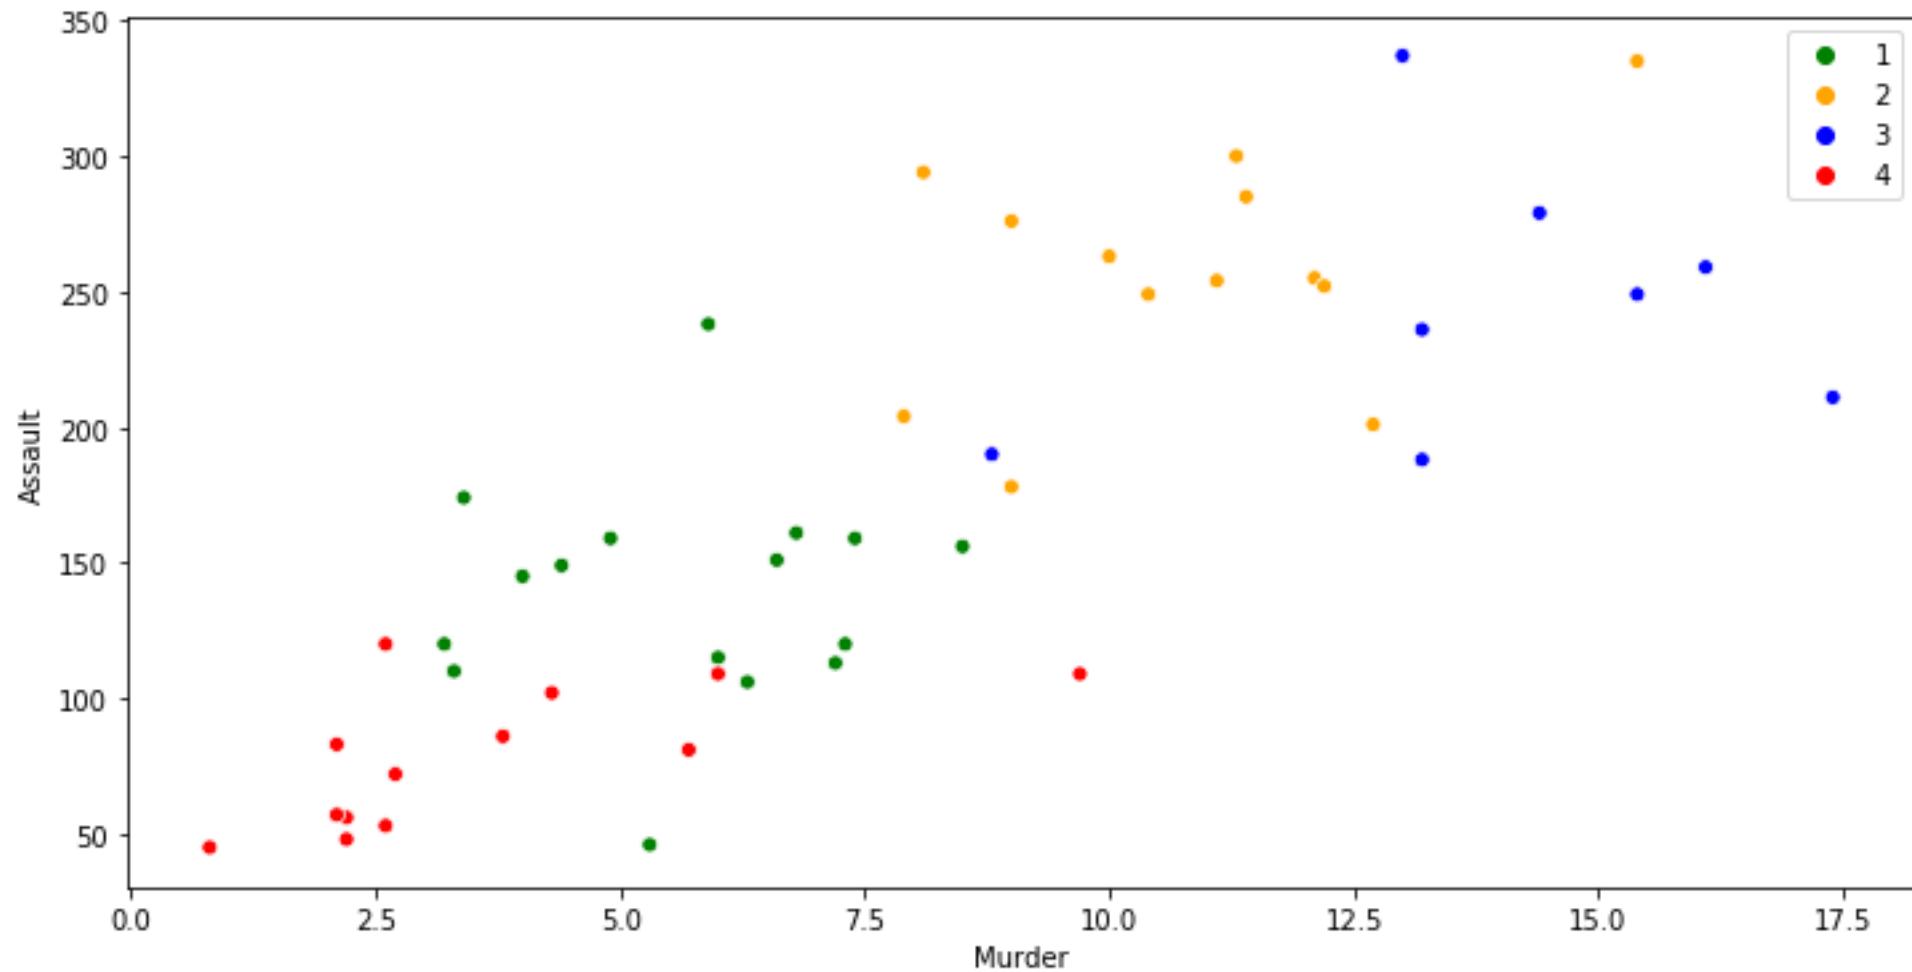
Plot

```
In [48]: x = df['Murder']
....: y = df['Assault']
....: plt.figure(figsize=(12,6))
....: sns.scatterplot(x, y, hue=y_kmeans1,
....:                  palette=['green','orange','blue','red'], legend='full')
....:
```

```
Out[48]: <matplotlib.axes._subplots.AxesSubplot at 0x23393fb3710>
```



Plot



Cophenetic Correlation

Average Linkage

	A	B	C	D	E
A	*	6.32 [7.49]	7.07 [7.49]	1.41 [1.41]	7.16 [7.49]
B		*	1.41 [1.74]	7.62 [7.49]	1.12 [1.12]
C			*	8.25 [7.49]	2.06 [1.74]
D				*	8.50 [7.49]
E					*

We merged b & e with (b, e) at a distance 1.12

Then we merged a & d with (a, d) at a distance 1.41

We merged c with ((b, e), c) at a distance 1.74

Then we merged (a, d) with ((b, e), c) at a distance 7.49

Average Linkage	
Dist	Merged at dist
6.32	7.49
7.07	7.49
1.41	1.41
7.16	7.49
1.41	1.74
7.62	7.49
1.12	1.12
8.25	7.49
2.06	1.74
8.5	7.49
Correl = 0.980423166	

Silhouette Analysis

Silhouette analysis can be used to determine the degree of separation between clusters. For each sample:

- Compute the average distance from all data points in the same cluster (a_i).
- Compute the average distance from all data points in the closest cluster (b_i).
- Compute the coefficient:

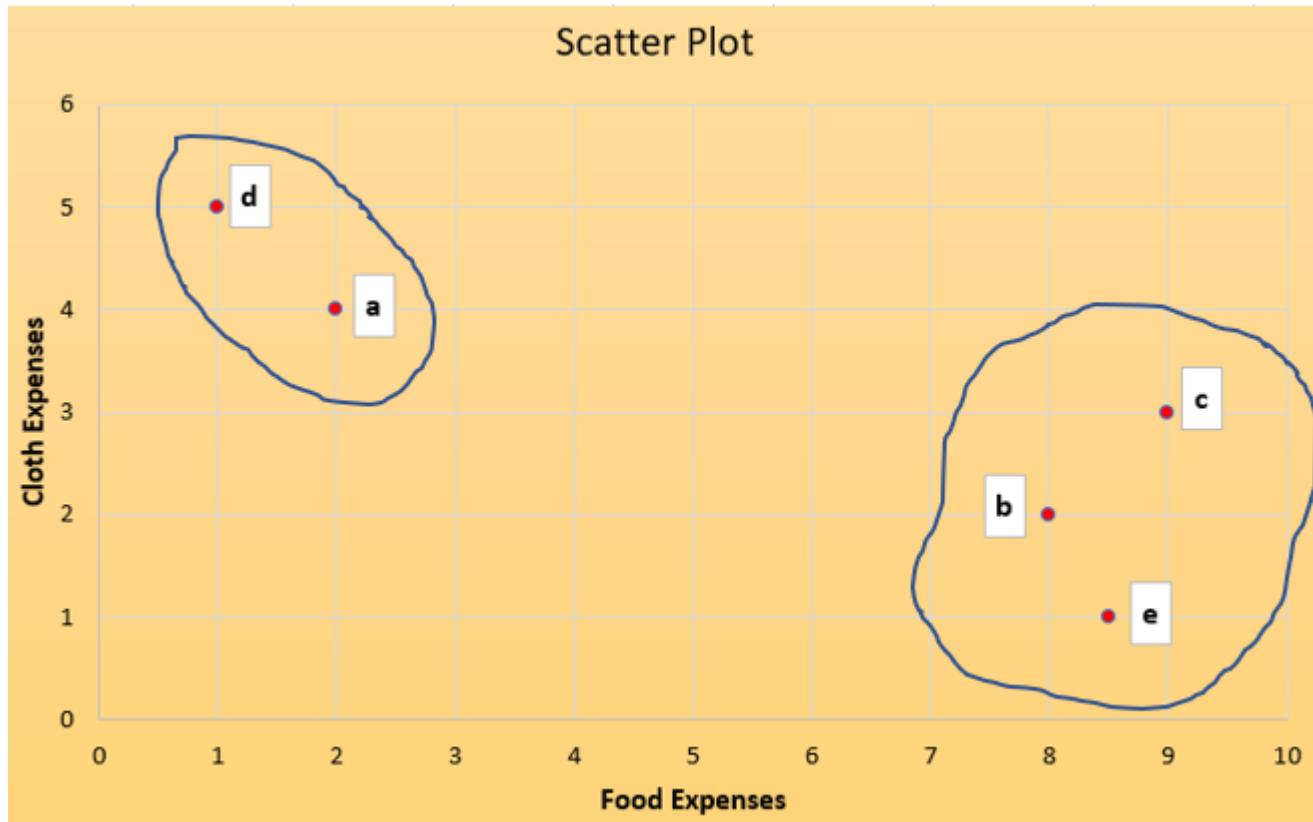


$$\frac{b^i - a^i}{\max(a^i, b^i)}$$

The coefficient can take values in the interval $[-1, 1]$.

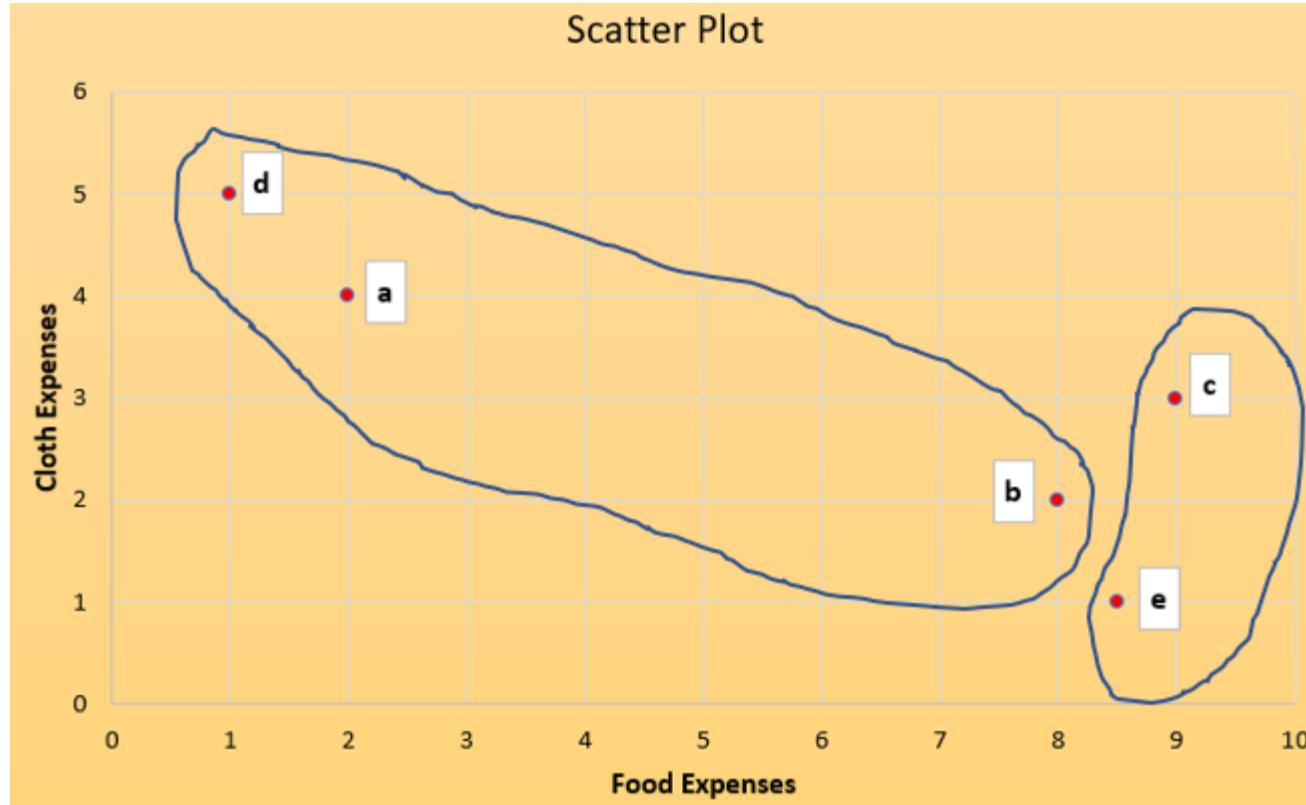
- If it is 0 \rightarrow the sample is very close to the neighbouring clusters.
- If it is 1 \rightarrow the sample is far away from the neighbouring clusters.
- If it is -1 \rightarrow the sample is assigned to the wrong clusters.

Silhouette_score (K Means)



Silhouette Score for a			
C1	ad	1.41	A
C2	ab	6.32	
C2	ac	7.07	
C2	ae	7.16	
AVG =		6.85	B
Sil_Sc(a)= 0.794161			
As it is close to 1, rightly allotted C1			

Silhouette_score (K Means)



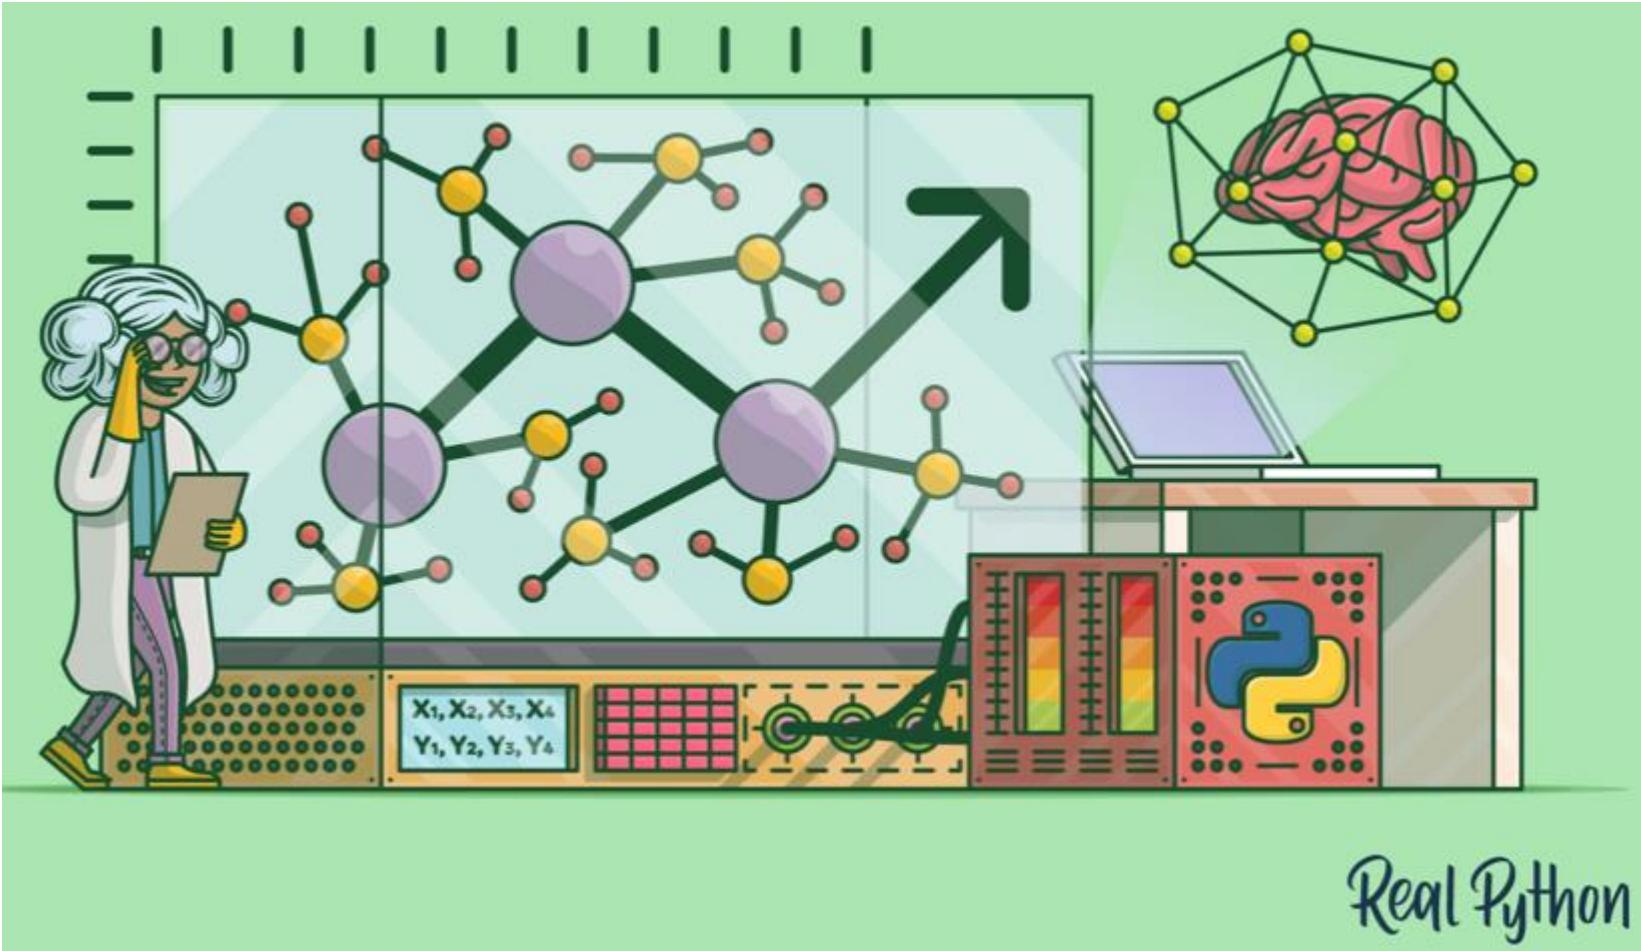
Silhouette Score for b		
	d	
C1	ba	6.32
C1	bd	7.62
	AVG=	6.97 A
C2	be	1.12
C2	bc	1.41
	AVG=	1.265 B
Sil_Sc(b)=		-0.81851
As it is close to -1, wrongly allotted to C1		

Artificial Neural Networks

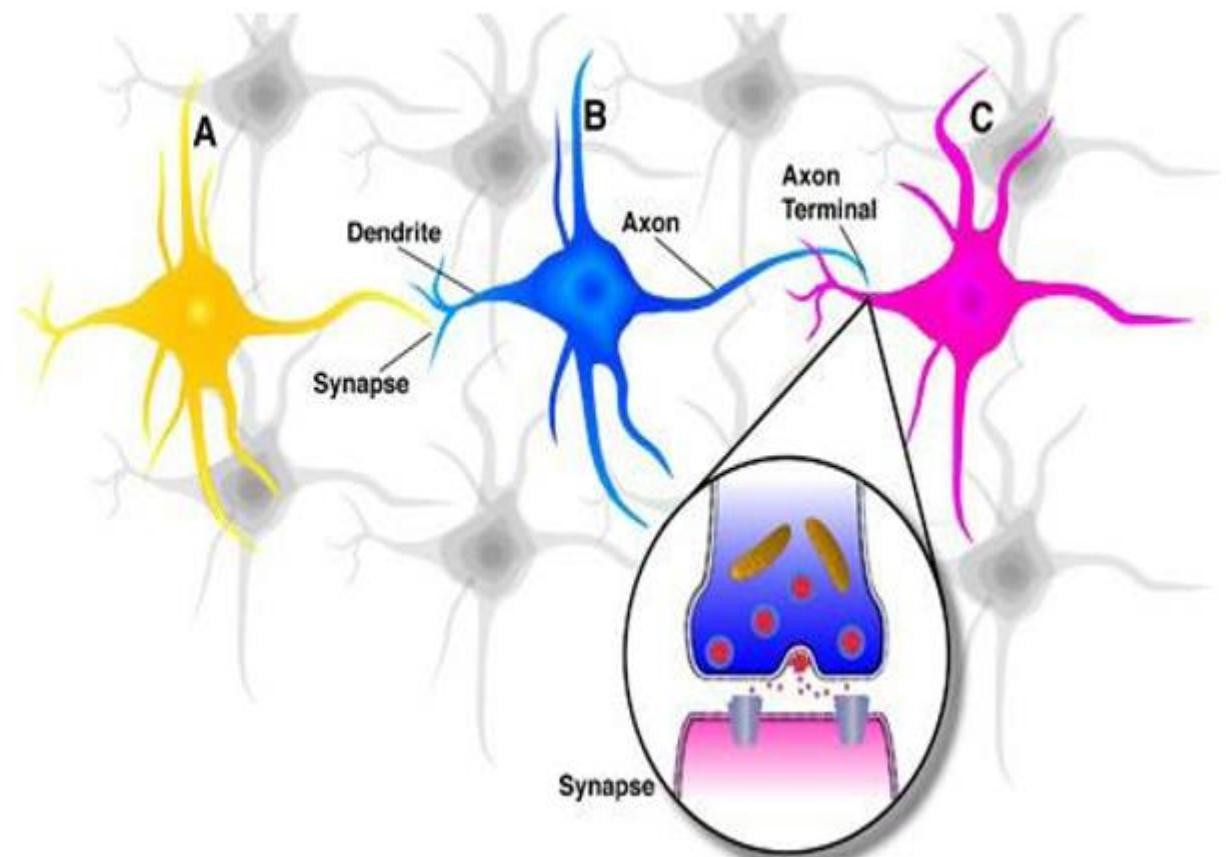
Human Brain!

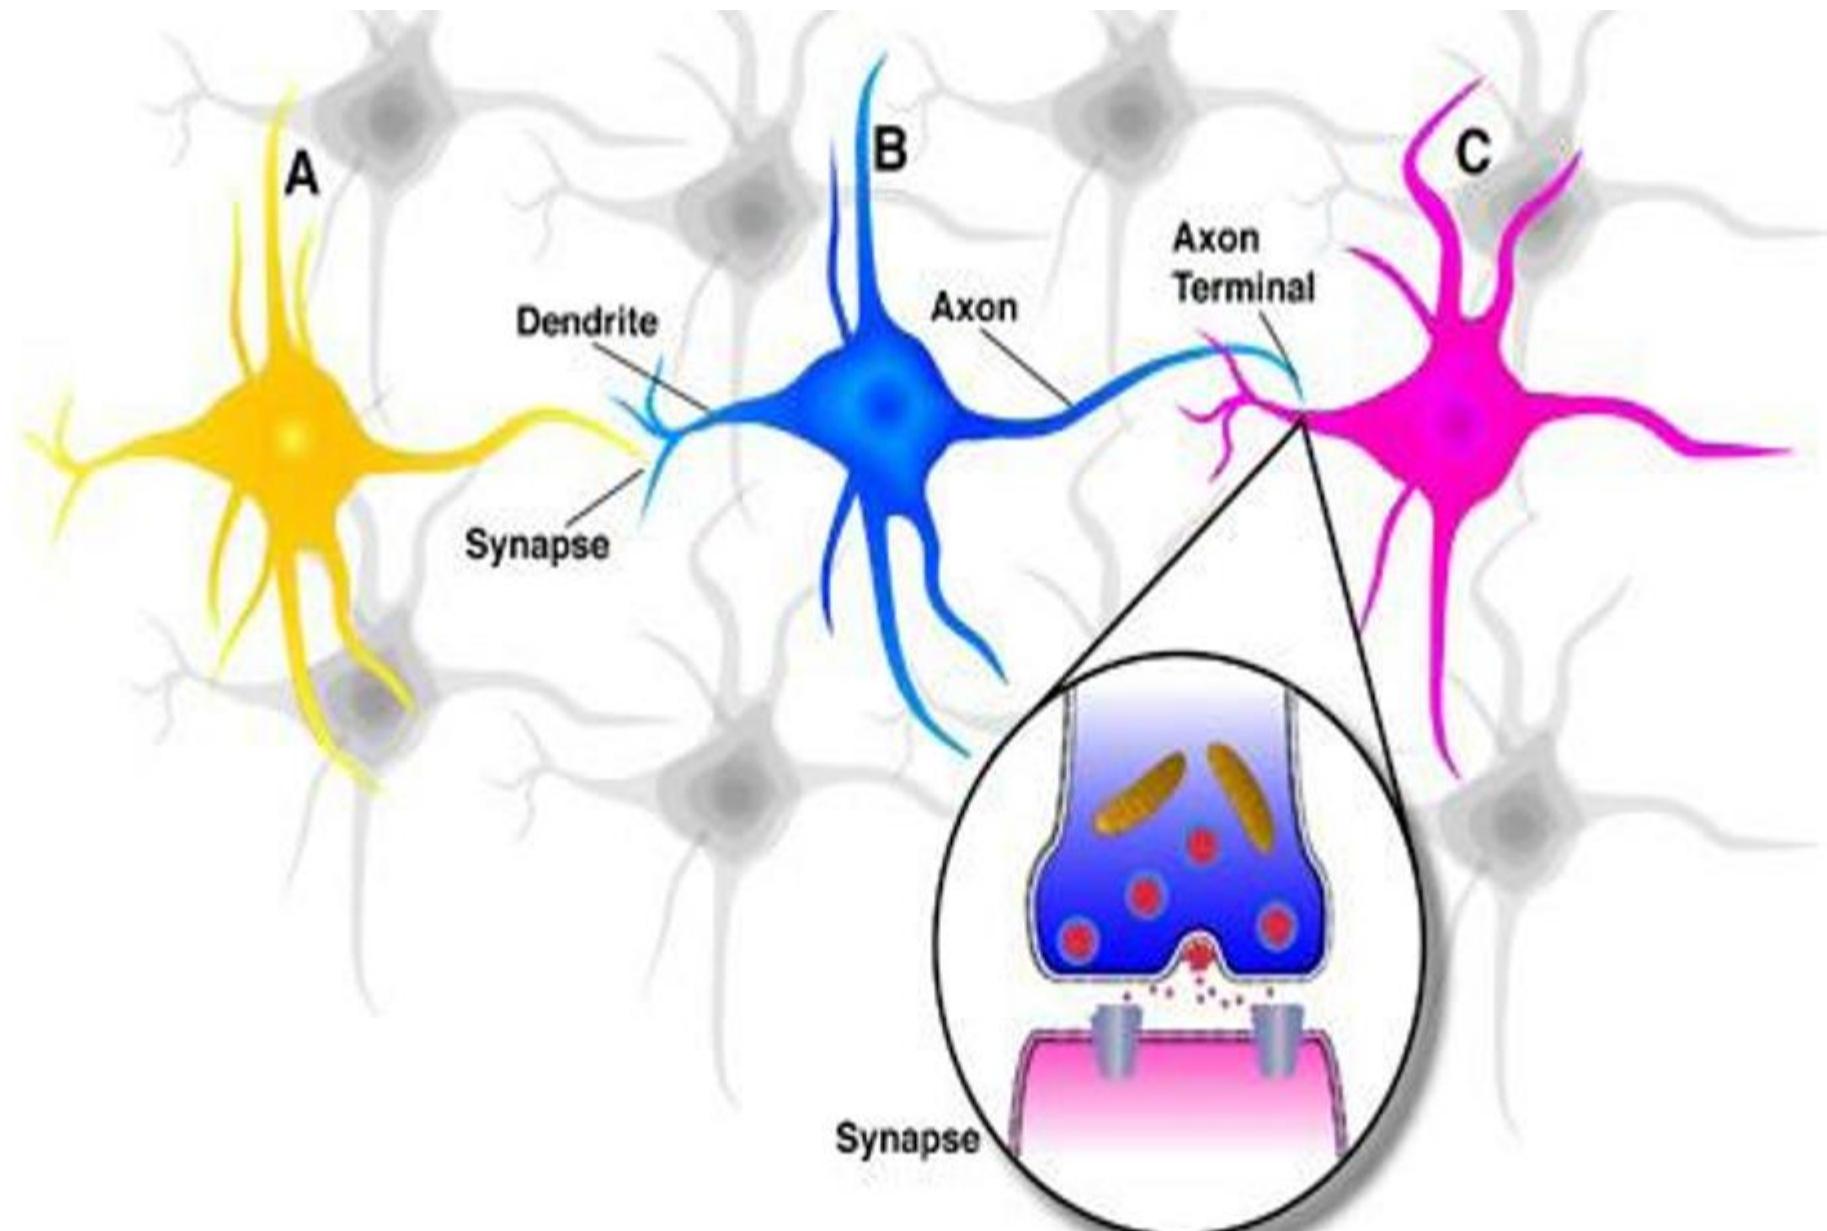


Why Brain!



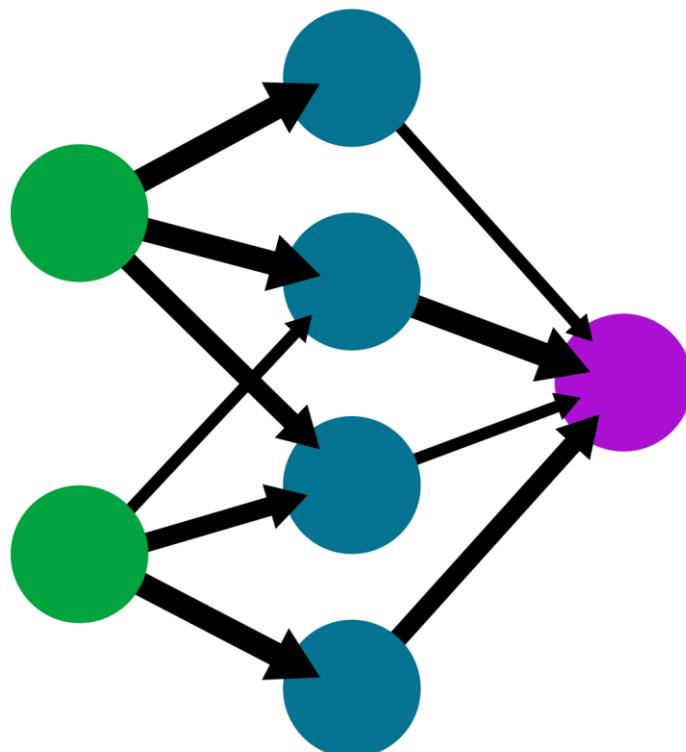
Biological Neural Network

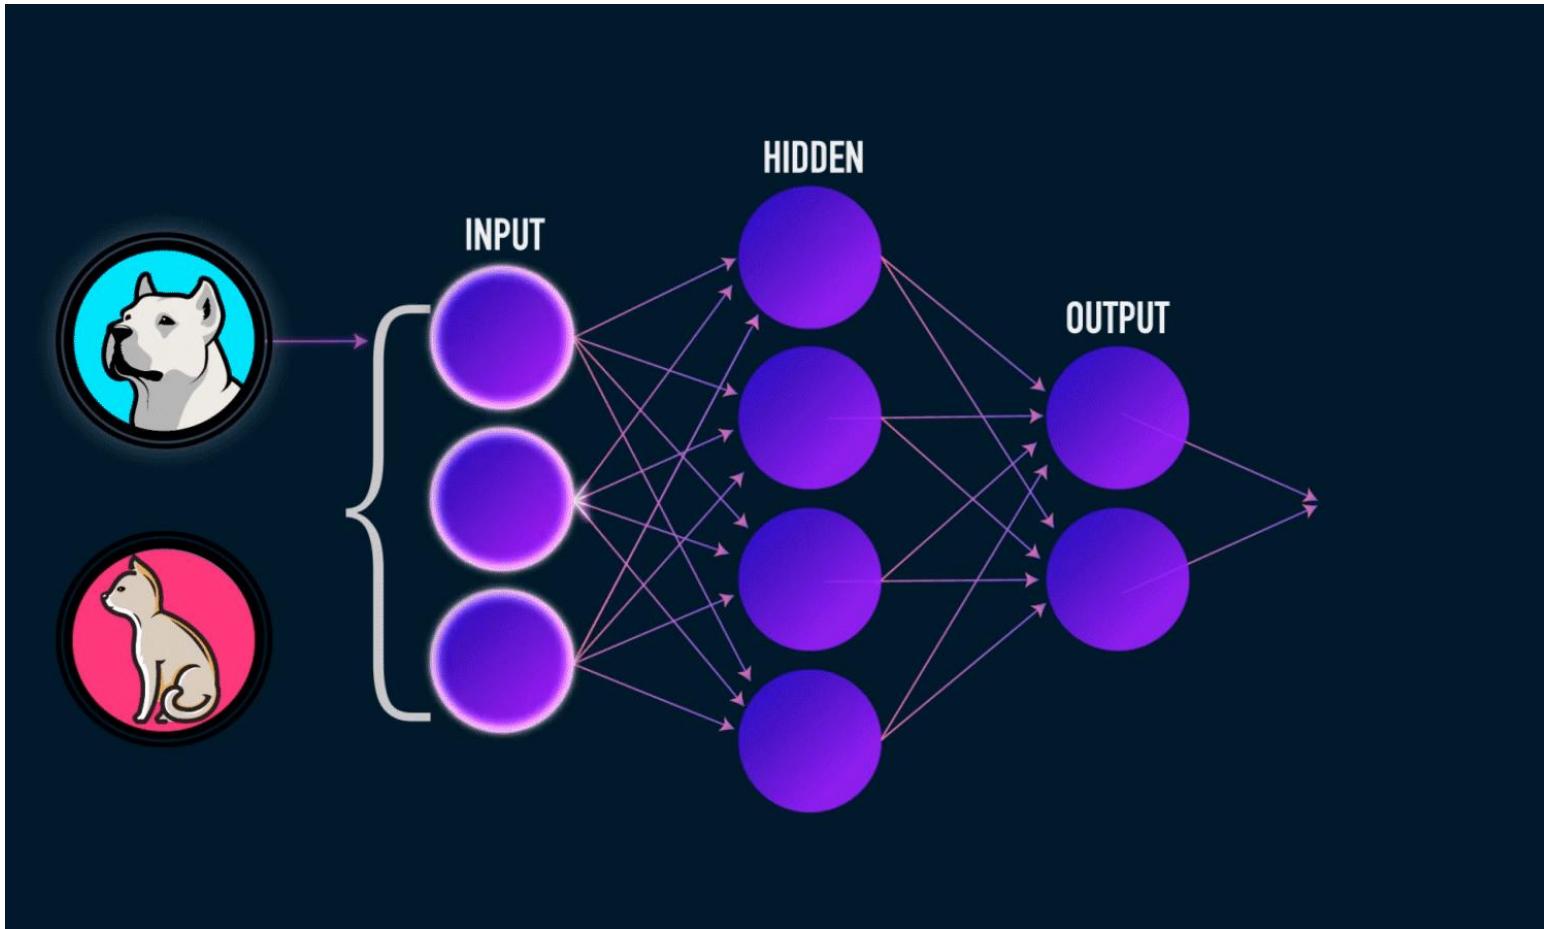




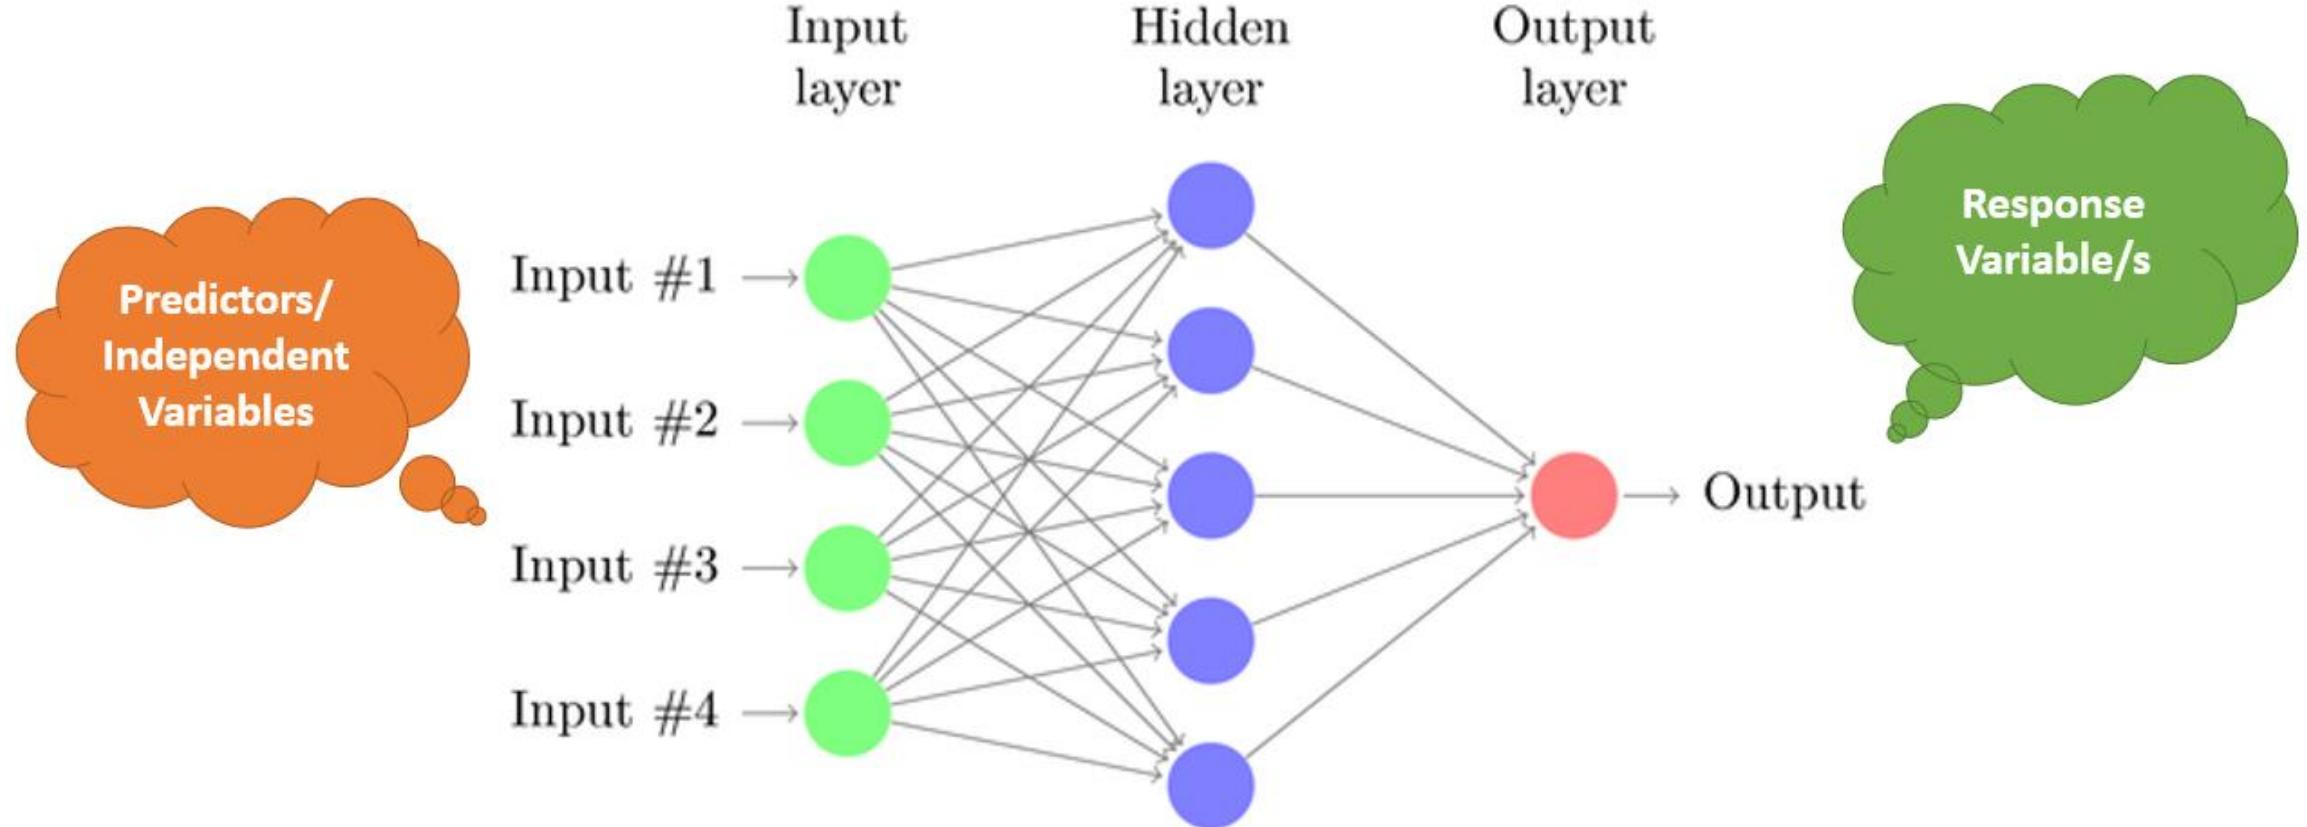
A simple neural network

input layer hidden layer output layer





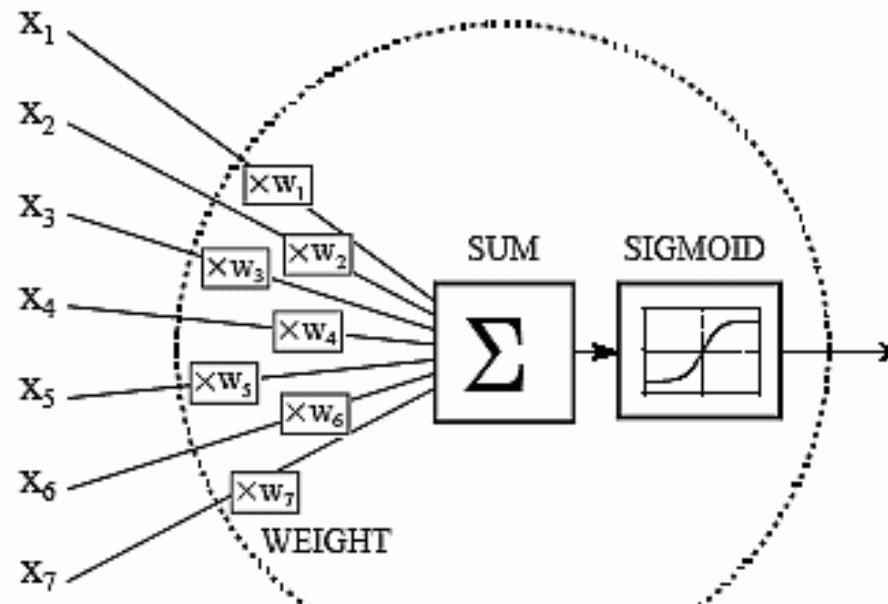
Network

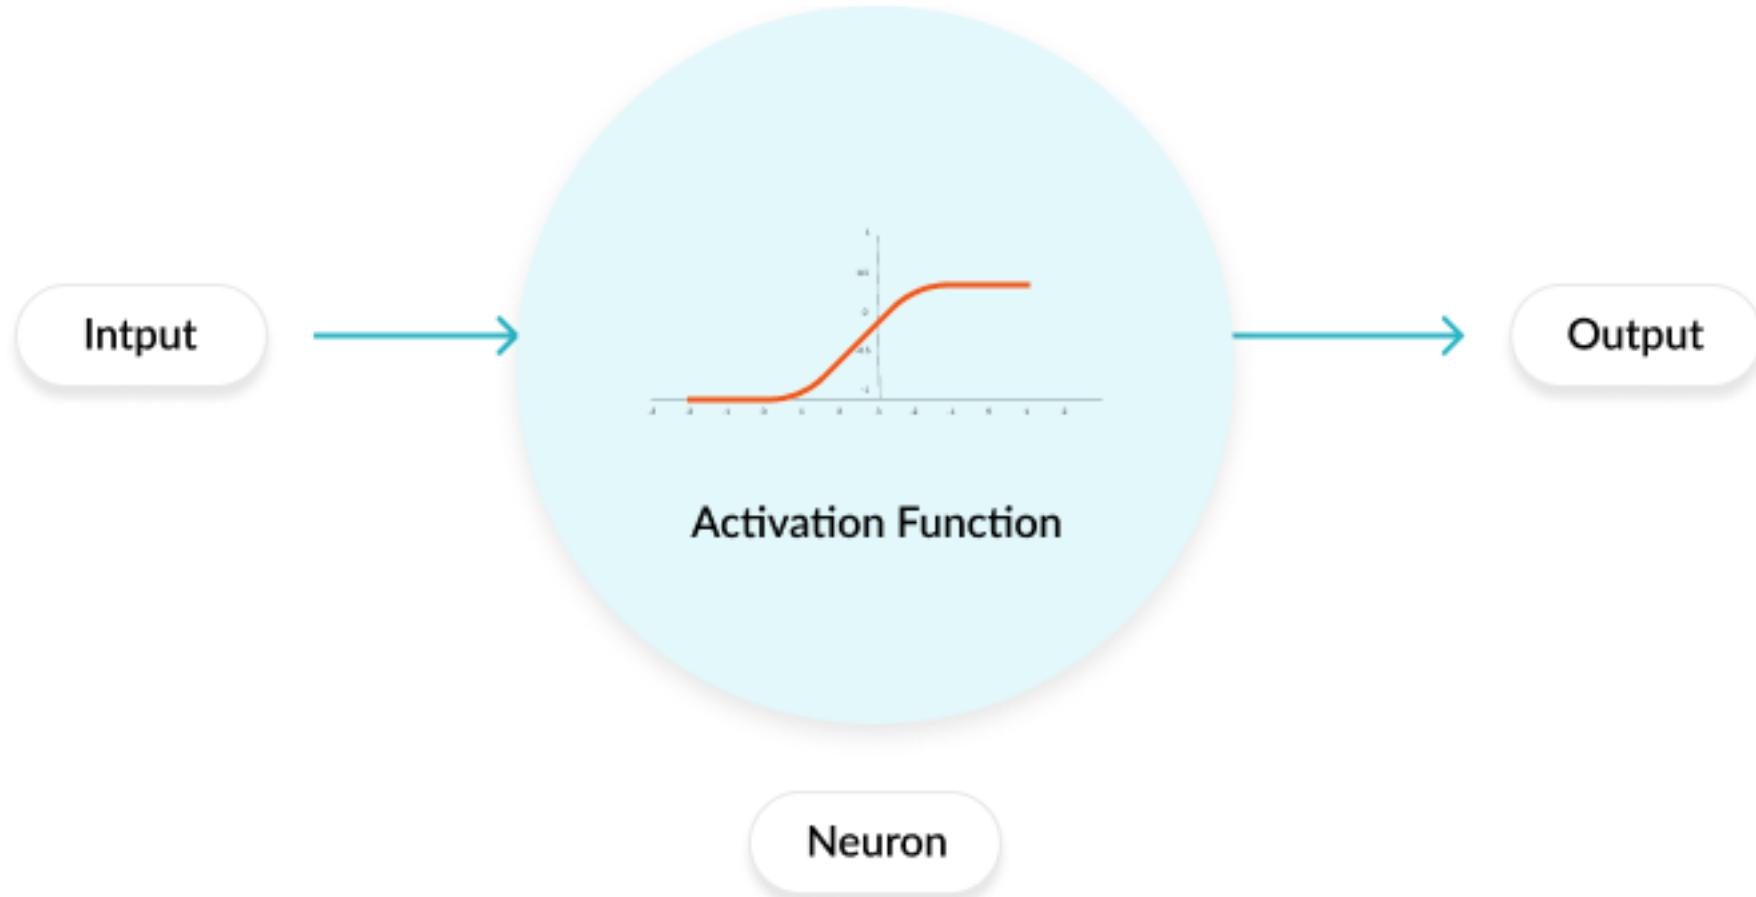


Activation Functions

FIGURE 26-6

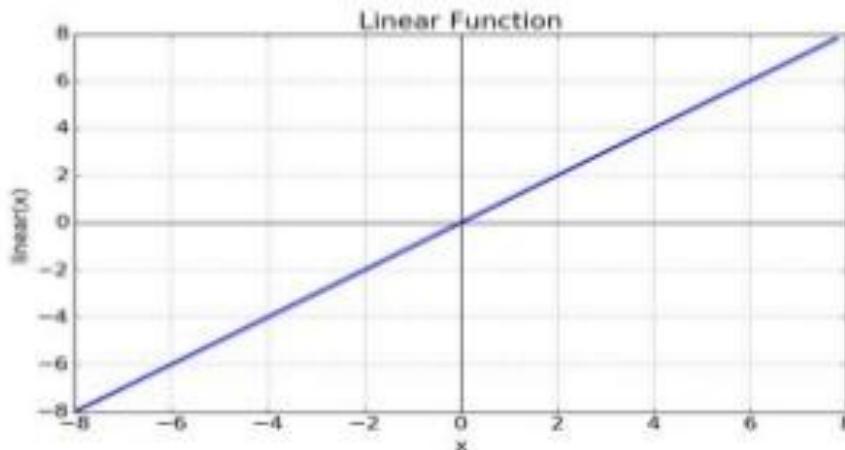
Neural network active node. This is a flow diagram of the active nodes used in the hidden and output layers of the neural network. Each input is multiplied by a weight (the w_N values), and then summed. This produces a single value that is passed through an "s" shaped nonlinear function called a *sigmoid*. The sigmoid function is shown in more detail in Fig. 26-7.



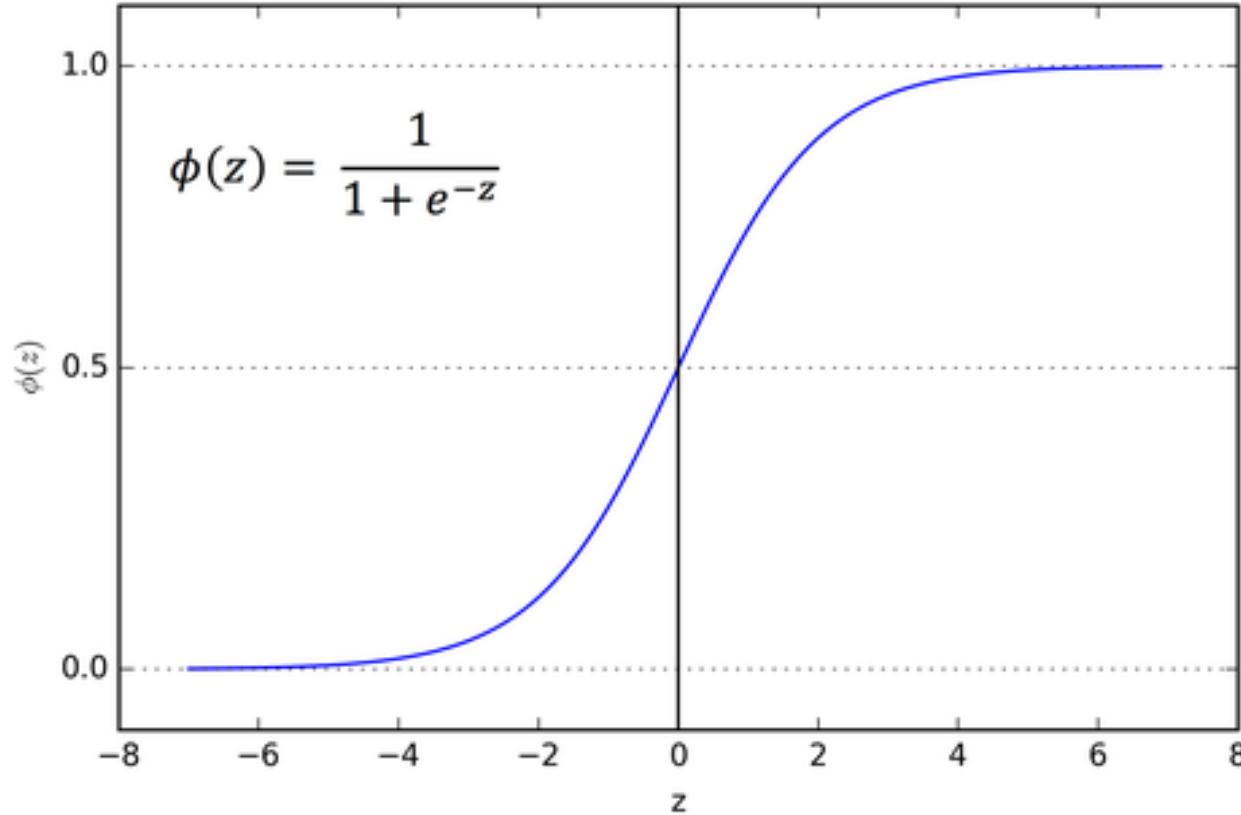


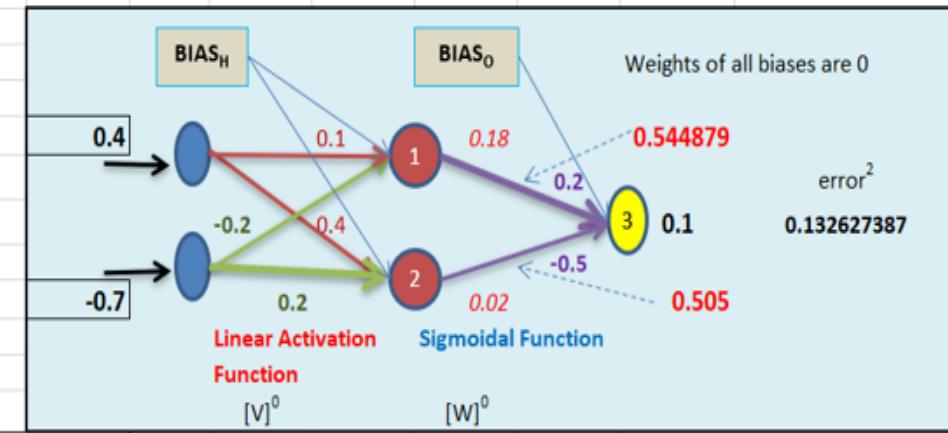
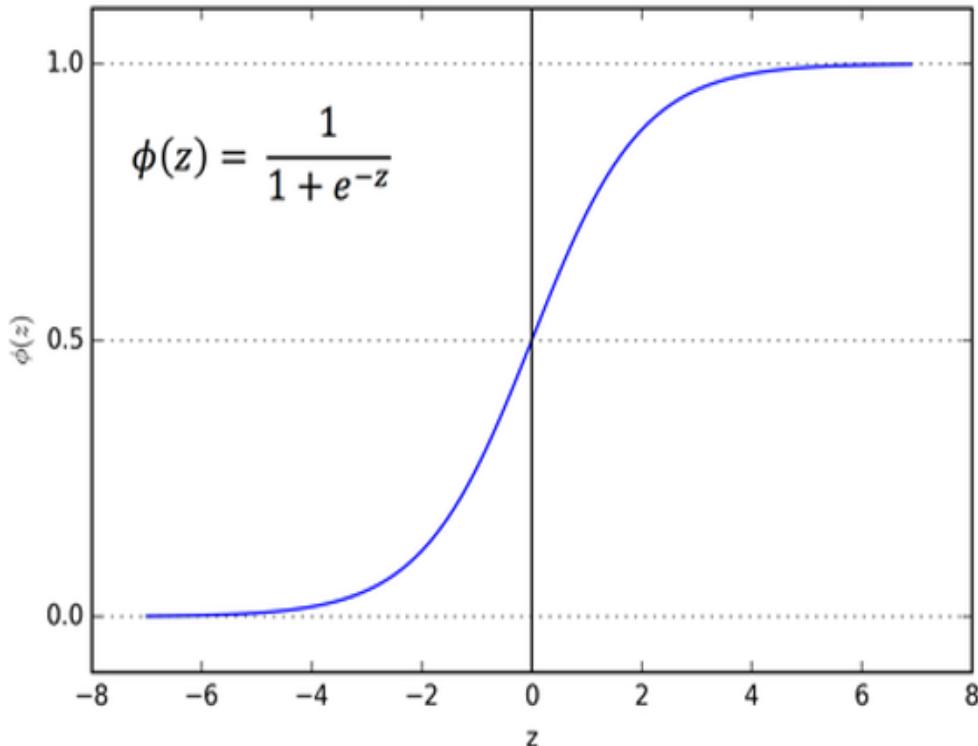
Linear or Identity Activation Function

- As you can see the function is a line or linear. Therefore, the output of the functions will not be confined between any range.
- Equation : $f(x) = x$
- Range : (-infinity to infinity)
- It doesn't help with the complexity or various parameters of usual data that is fed to the neural networks.



Sigmoid or logistic





At node 1

First input(0.4) multiplied by 0.1 + Second input (-0.7) multiplied by -0.2
but, 0.18 will not be taken forward!

Sigmoid function will translate 0.18 to

0.54488

For node 1,	0.18	-0.18
For node 2,	0.02	-0.02

for using here

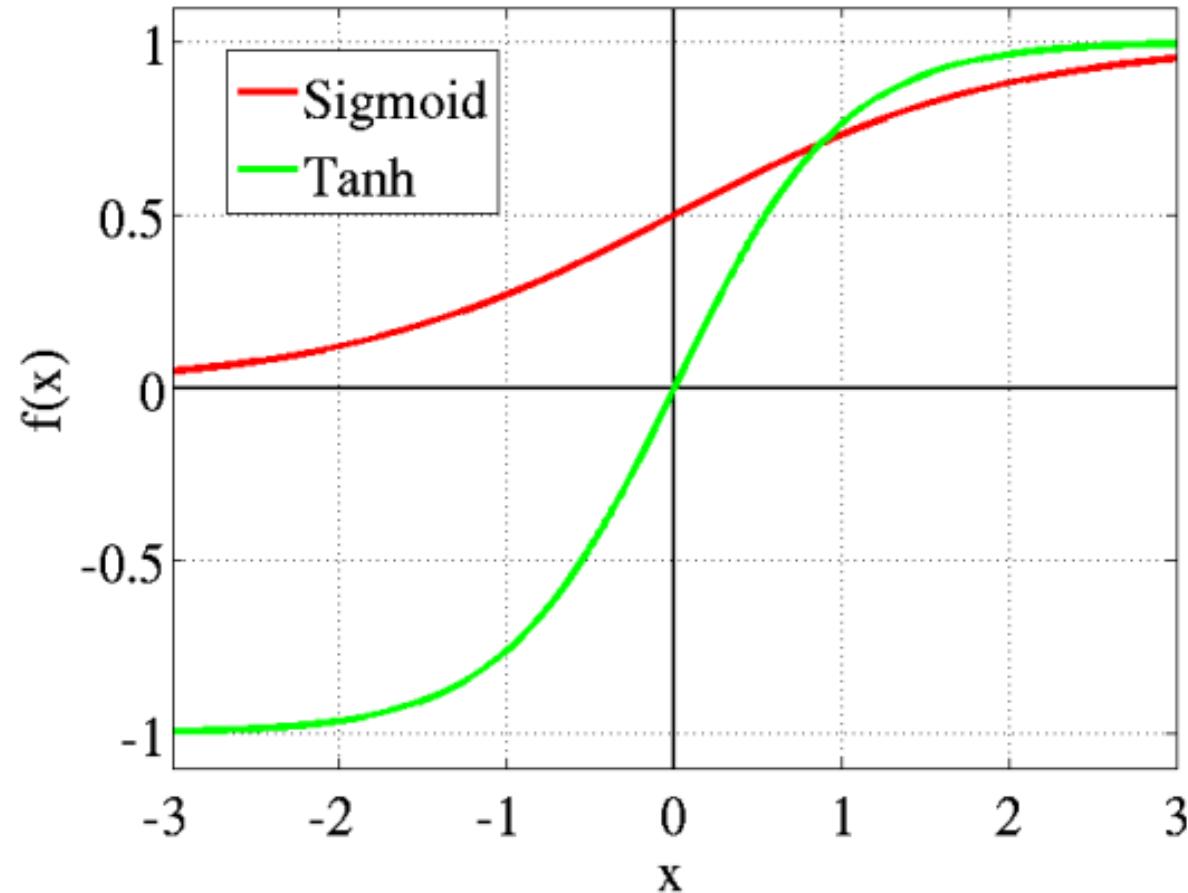
Step 4 Find output from hidden which will be input for output node.

Sigmoidal Function →

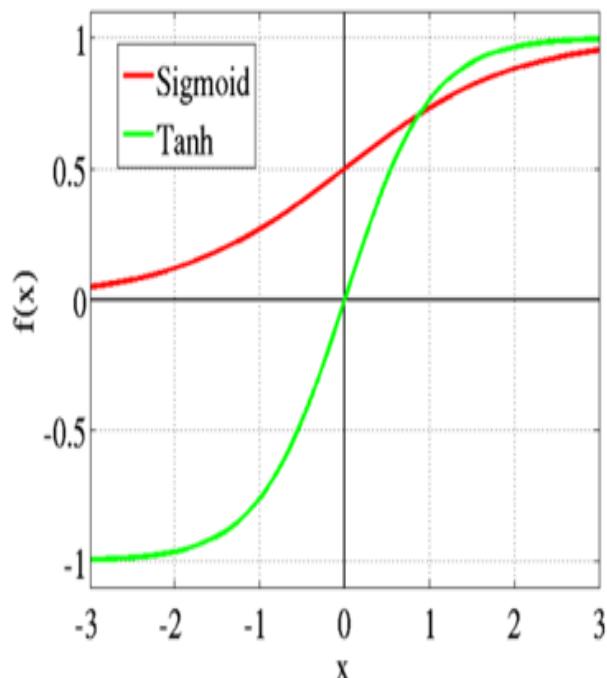
$1/(1+e^{-0.18})$ 0.544879
 $1/(1+e^{-0.02})$ 0.18 and 0.02

carefully note, output from hidden nodes is just not
rather these are converted by sigmoidal function!

tanh



tanh



TanH

The graph shows the tanh function $f(x) = \tanh(x) = \frac{2}{1+e^{-2x}} - 1$. The curve is a red sigmoid function centered at the origin.

At node 1

First input(0.4) multiplied by 0.1 + Second input (-0.7) multiplied by -0.2

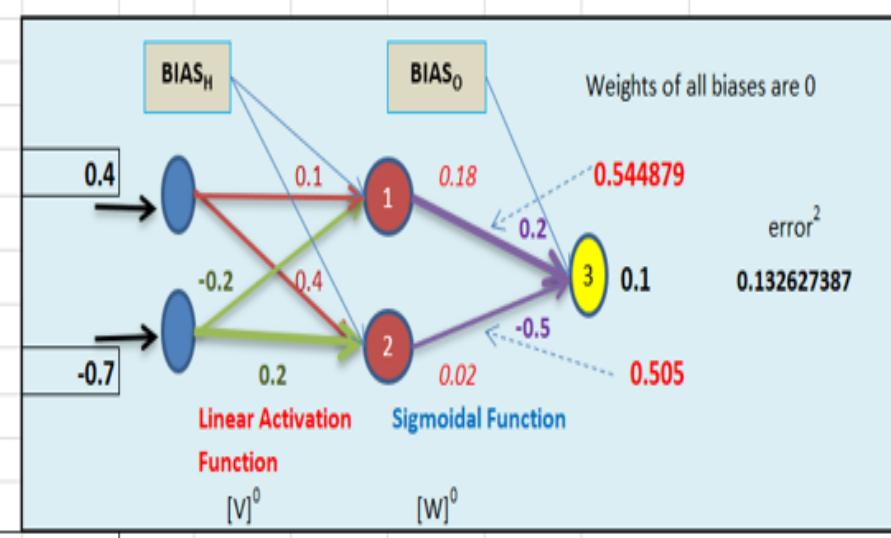
0.18

but, 0.18 will not be taken forward!

tanH function will translate 0.18 to

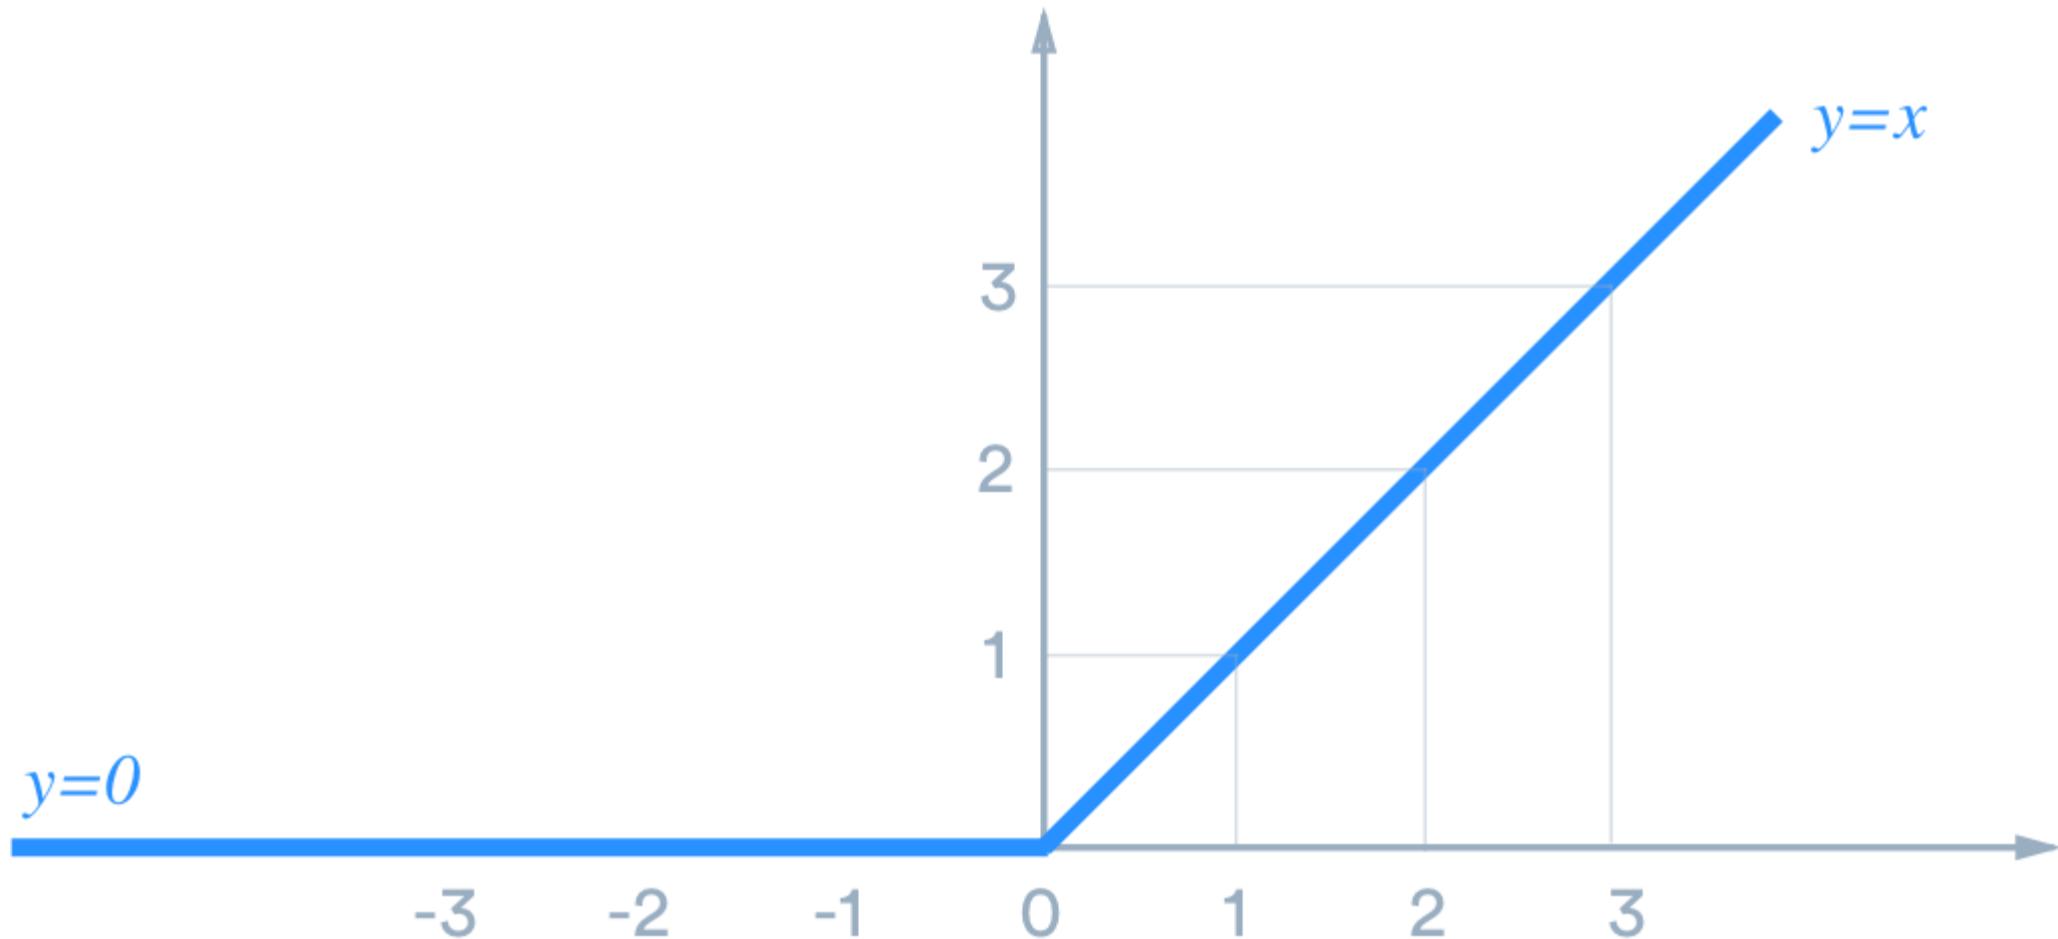
0.17808

0.178081

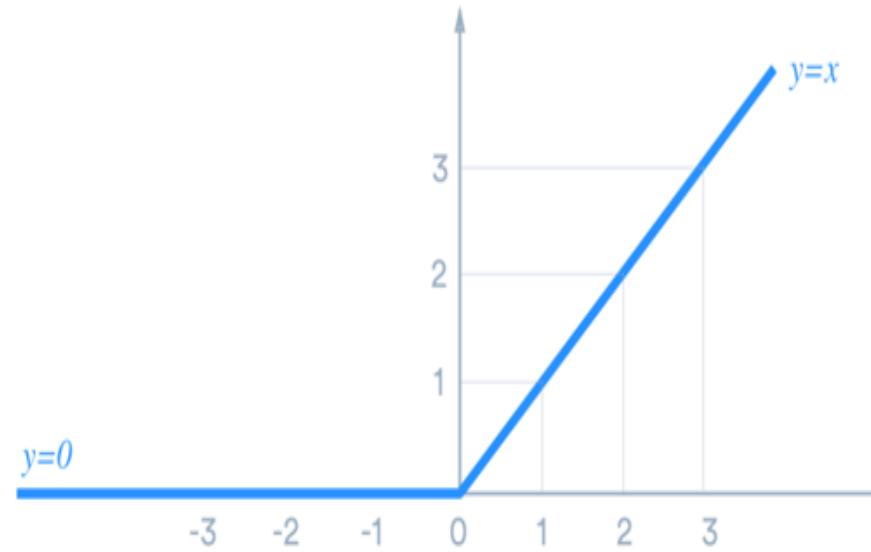


Denom	1.697676
2/Denom	1.178081
(2/Denom)-1	0.178081

relu



relu



Rectified
Linear Unit
(ReLU)



$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

At node 1

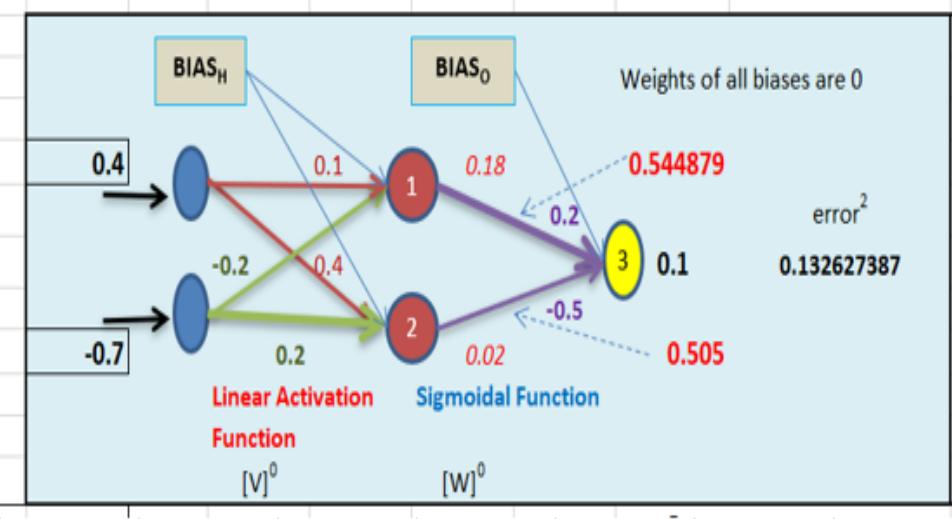
First input(0.4) multiplied by 0.1 + Second input (-0.7) multiplied by -0.2

0.18

but, 0.18 will be subjected to conditions as mentioned above, as 0.18 is ≥ 0

relu function will translate 0.18 to

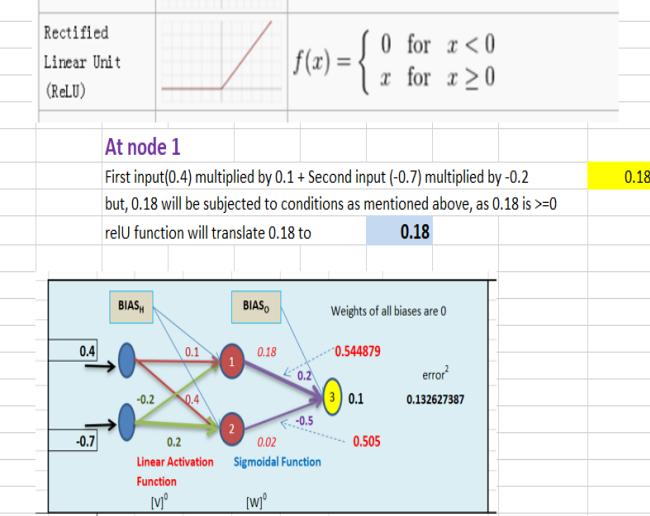
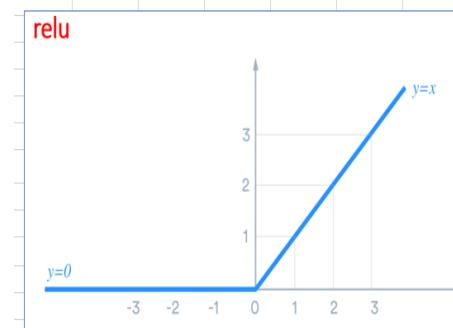
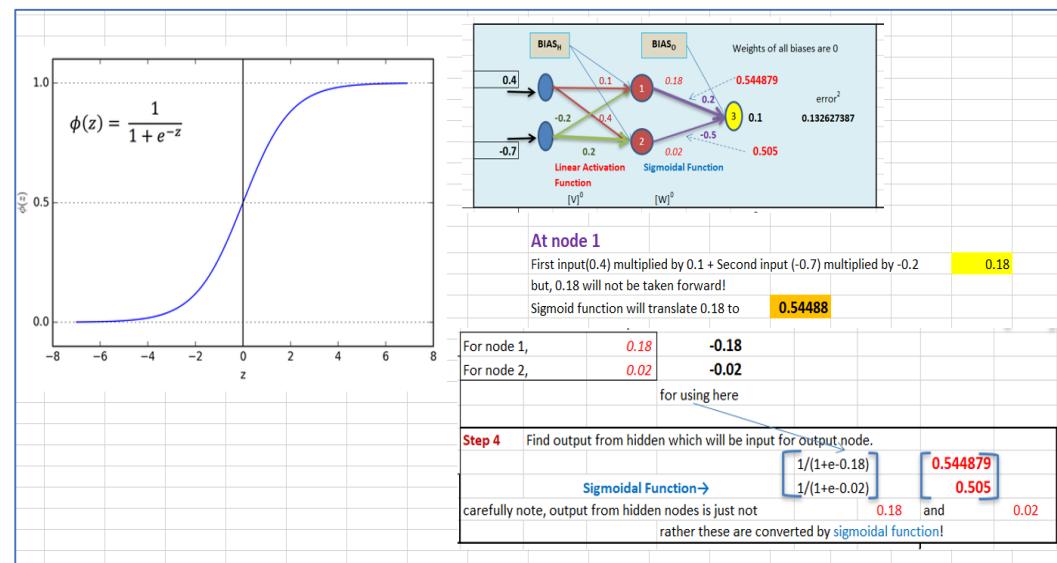
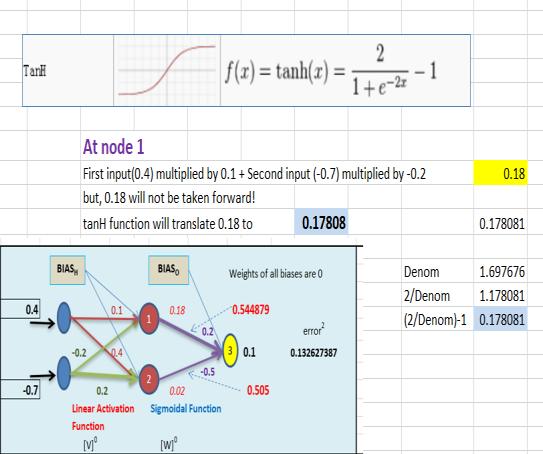
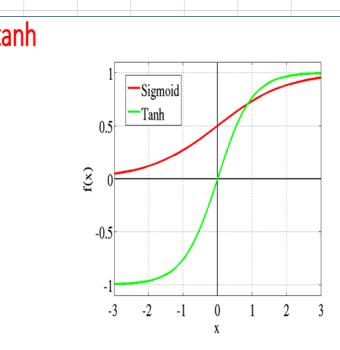
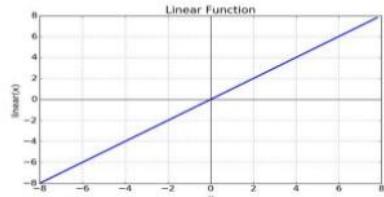
0.18



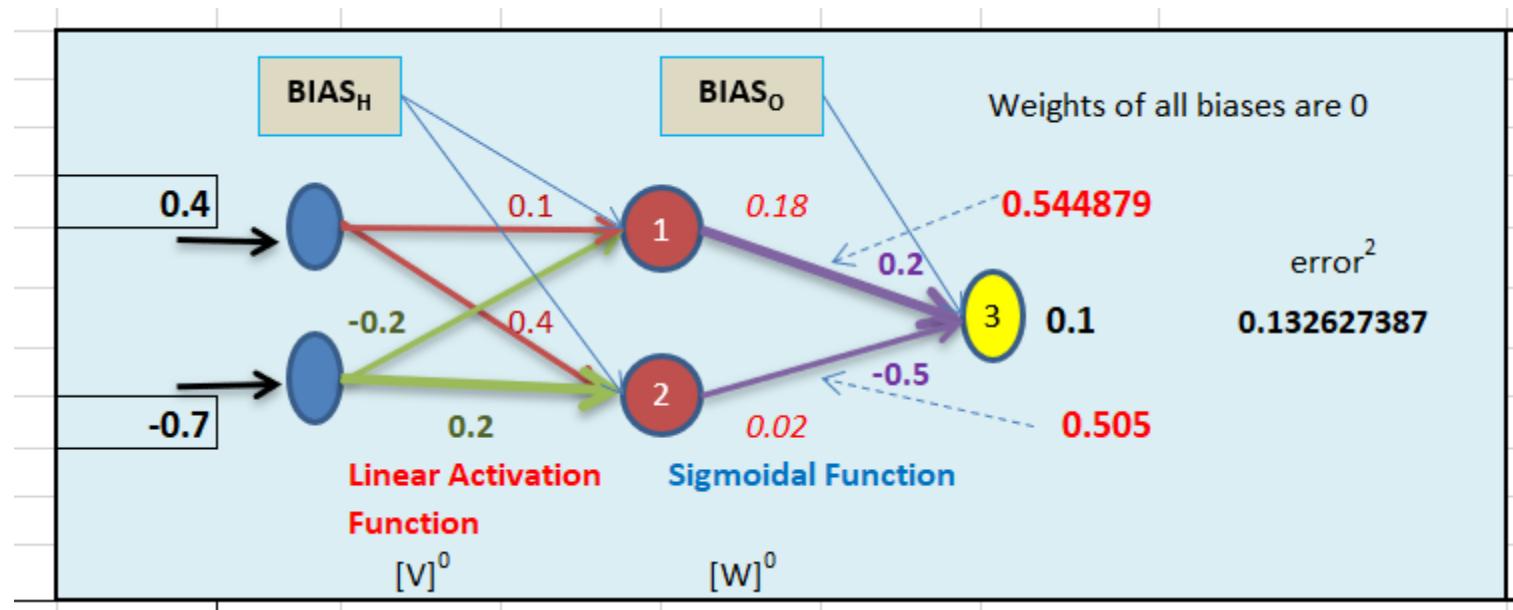
Lets see all together!

Linear or Identity Activation Function

- As you can see the function is a line or linear. Therefore, the output of the functions will not be confined between any range.
- Equation :** $f(x) = x$
- Range :** (-infinity to infinity)
- It doesn't help with the complexity or various parameters of usual data that is fed to the neural networks.



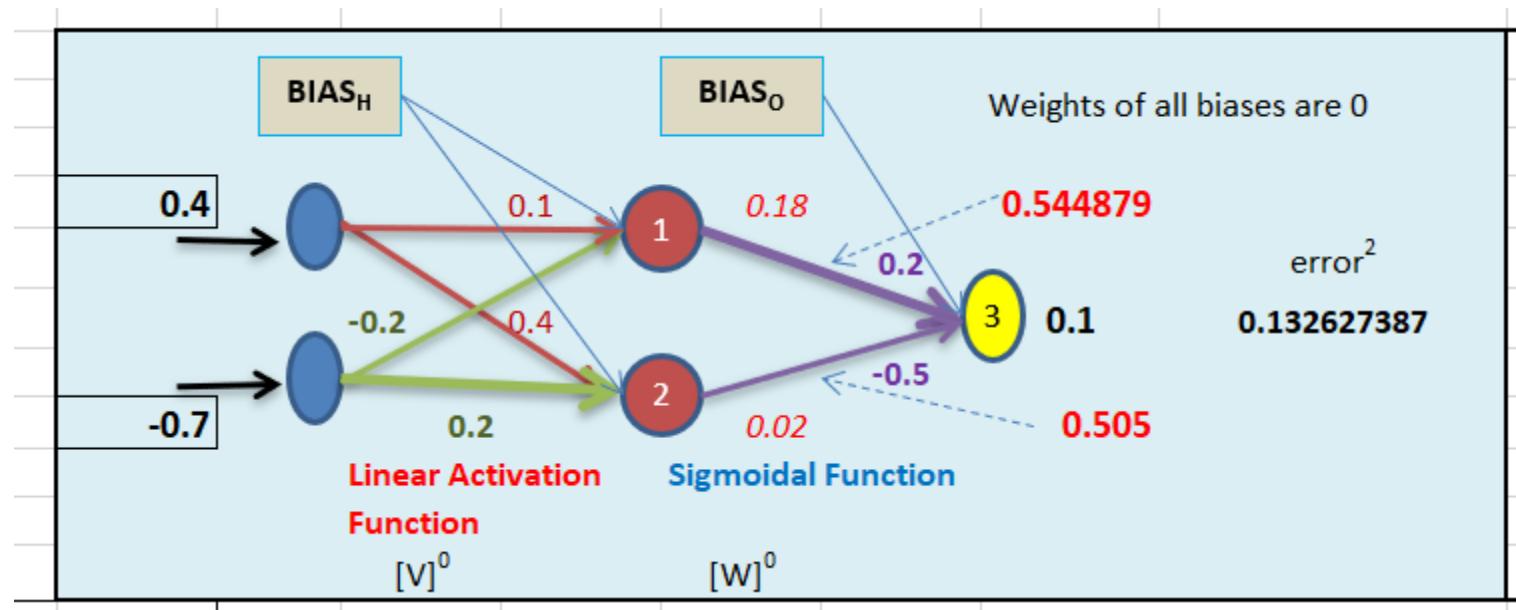
Gradient Descent



Gradient Descent

Sr No	I1	I2	O	
1	0.4	-0.7	0.1	$[O]_i = \{I\}_i =$
2	0.3	-0.5	0.05	output
3	0.6	0.1	0.3	of input layer
4	0.2	0.4	0.25	
5	0.1	-0.2	0.12	

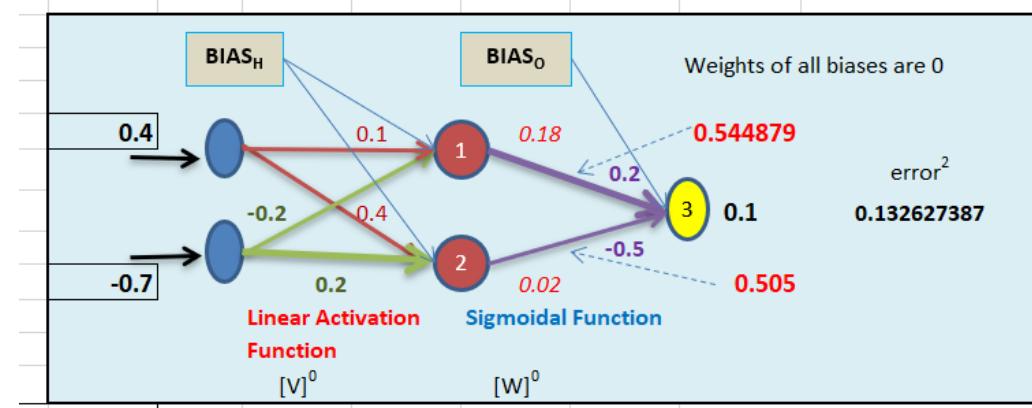
Gradient Descent



Gradient Descent

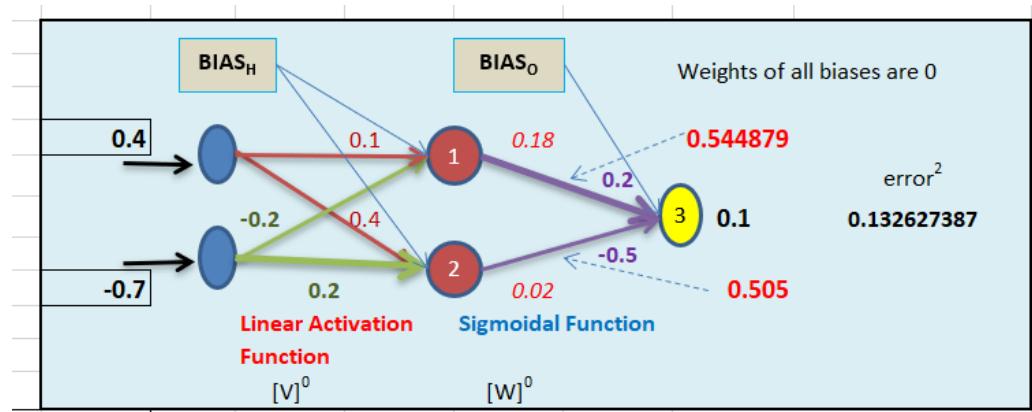
Step 1 All inputs are between -1 & +1 so no need to normalise.

We will take only 1st set I_1 as 0.4 and I_2 as -0.7 ; 0.4 & 0.7 are not autmtc



Gradient Descent

Step 2	Initialize wts as	$[V]^0$	$\begin{bmatrix} 0.1 & 0.4 \\ -0.2 & 0.2 \end{bmatrix}$	$[W]^0$	$\begin{bmatrix} 0.2 \\ -0.5 \end{bmatrix}$	$2*2$	$2*1$
rows * columns							

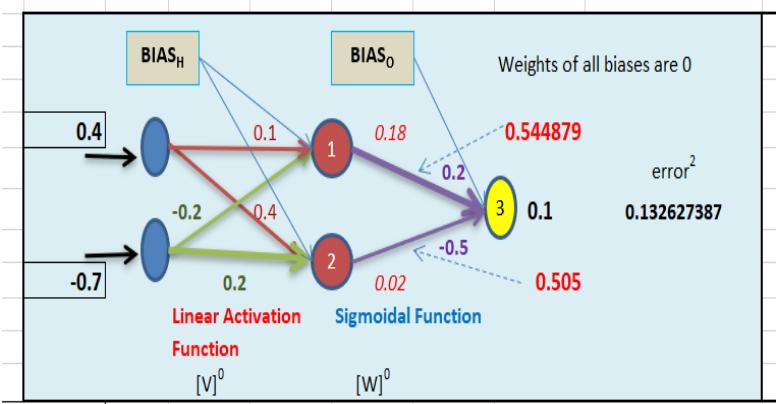


Gradient Descent

Step 3 Find $\{I\}_H = [V]^T * \{I\}_I$ or

I stands for input, to Hidden layer.

T stands for transpose.



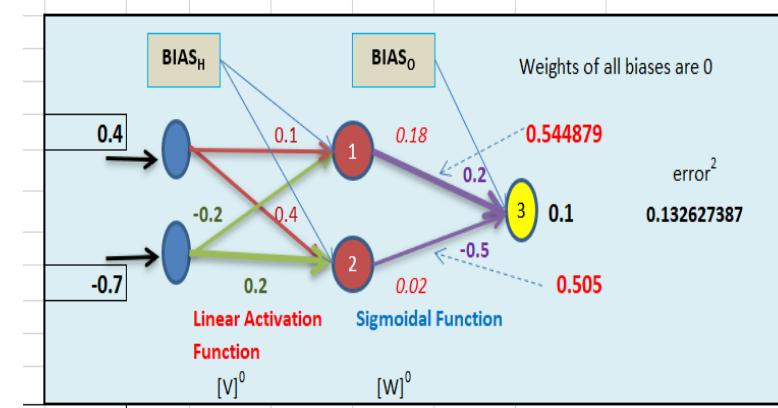
A handwritten diagram showing the multiplication of two matrices. The first matrix is a 2x2 matrix with columns 0.1, -0.2 and 0.4, 0.2. The second matrix is a 2x2 matrix with rows 0.4, 0.18 and -0.7, 0.18. The multiplication is performed row by column:

$$\begin{bmatrix} 0.1 & -0.2 \\ 0.4 & 0.2 \end{bmatrix} * \begin{bmatrix} 0.4 & 0.18 \\ -0.7 & 0.18 \end{bmatrix} = \begin{bmatrix} 0.02 \\ -0.7 \end{bmatrix}$$

Annotations include "Multiply row columnwise" and "This is matrix multiplication".

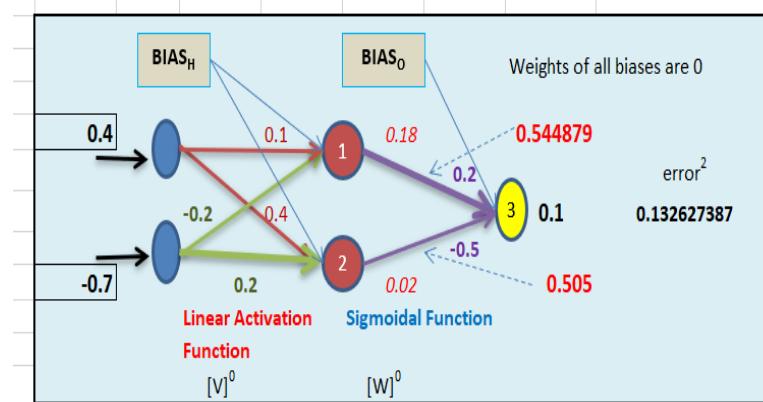
Gradient Descent

For node 1,	0.18	-0.18	
For node 2,	0.02	-0.02	
for using here			
Step 4 Find output from hidden which will be input for output node.			
Sigmoidal Function→	$1/(1+e^{-0.18})$	0.544879	
	$1/(1+e^{-0.02})$	0.505	
carefully note, output from hidden nodes is just not rather these are converted by sigmoidal function!	0.18	and	0.02



Gradient Descent

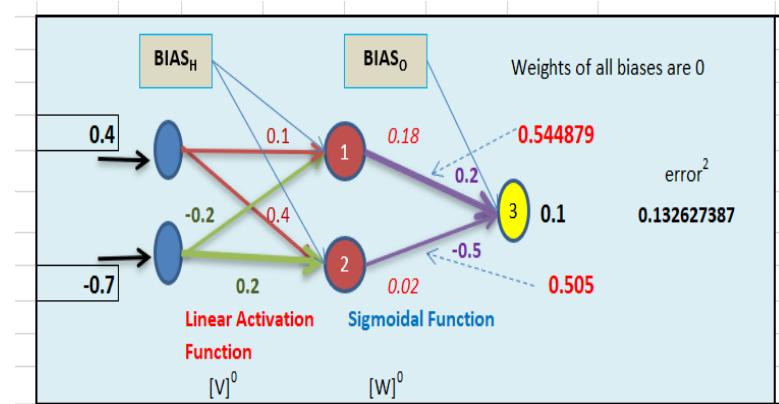
Step 5	Calculate output of output layer/node again by Simple matrix application.			
	$[I]_0 = [W]^T \cdot [O]_H = ($ 0.2 -0.5) Mutiply	0.544879 0.505	= -0.14352 And apply sigmoidal function and find output of output node/layer	0.143524



Gradient Descent

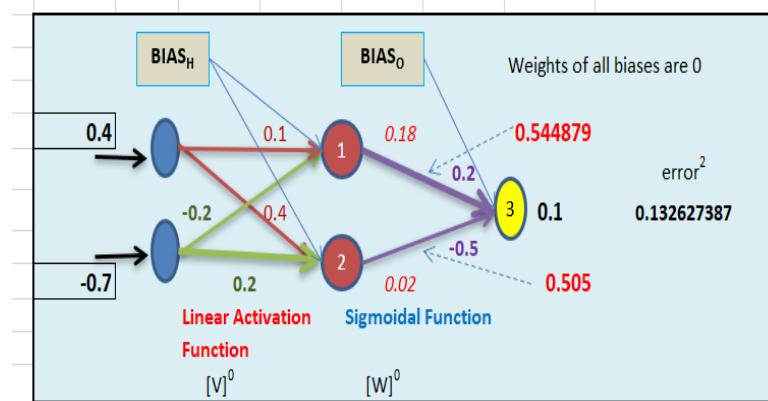
Step 6
 $1/[1+e^{-(-0.14352)}]$
0.46418

This is answer! But erroneous!!



Gradient Descent

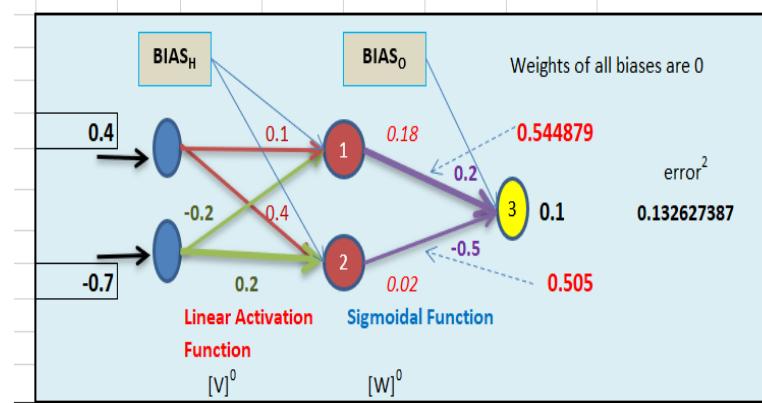
Step 7	Error = $(T_0 - O_0)^2 =$	0.132627
$T_0 =$	0.1	
$O_0 =$	0.46418	0.13262739



Gradient Descent

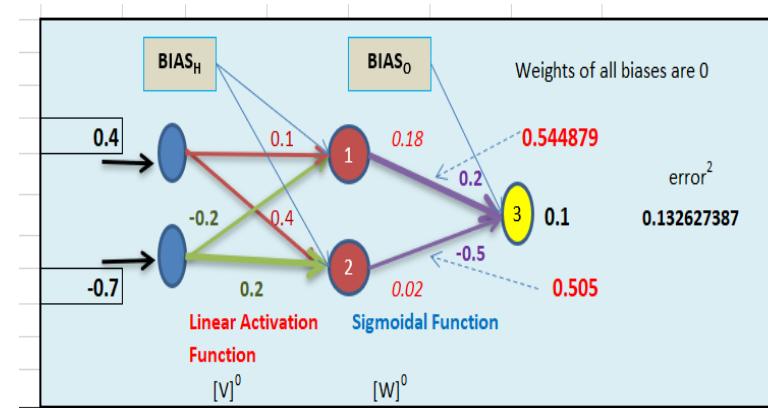
Step 8	Weight Adjustment
First find 'd' as =	$(T_0 - O_{01})(O_{01}(1-O_{01}))$ -0.0906
O_{01} = Output of output layer/node for 1st case.	

$$[Y] = [O]H <d> = \begin{bmatrix} 0.544879 \\ 0.505 \end{bmatrix} < -0.09058 > = \begin{bmatrix} -0.049354 \\ -0.045742 \end{bmatrix}$$



Gradient Descent

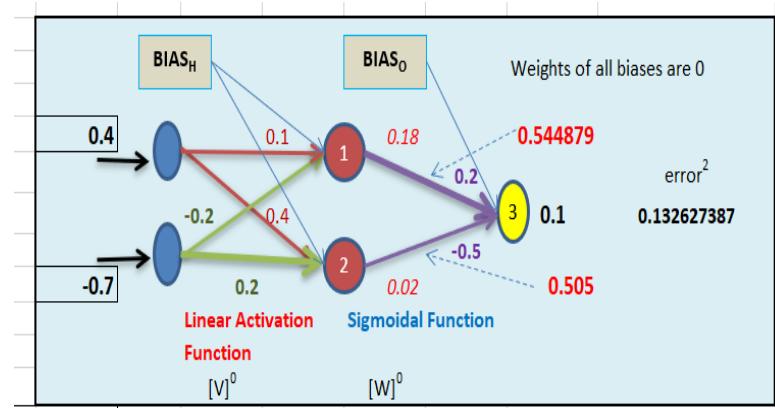
Step 9	Find $[\Delta W]^1$, as $= \alpha[\Delta W]^0 + \eta[Y]$
Assume η as	0.6
And initially, α as	0
α is Momentum coefficient	
η is Learning Rate	
So 1st term will be 0 as α is 0.	
$[\Delta W]^1 =$	$\begin{bmatrix} -0.049354 \\ -0.045742 \end{bmatrix} * 0.6$
$[\Delta W]^1 =$	$\begin{bmatrix} -0.029612 \\ -0.027445 \end{bmatrix}$



Gradient Descent

Step 10 $[e] = [W] * \{d\}$

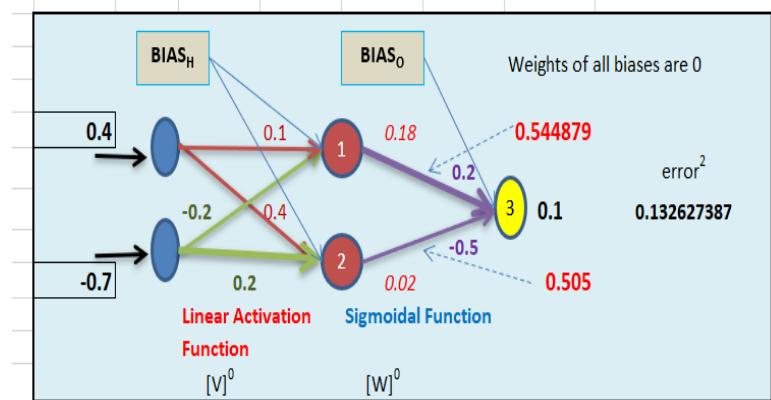
$$\begin{bmatrix} 0.2 \\ -0.5 \end{bmatrix} * -0.090577849 = \begin{bmatrix} -0.018 \\ 0.0453 \end{bmatrix}$$



Gradient Descent

Step 11 $[d^*] = e, \text{step 10's matrix} * \text{step 4's matrix} * (1 - \text{step 4's matrix})$

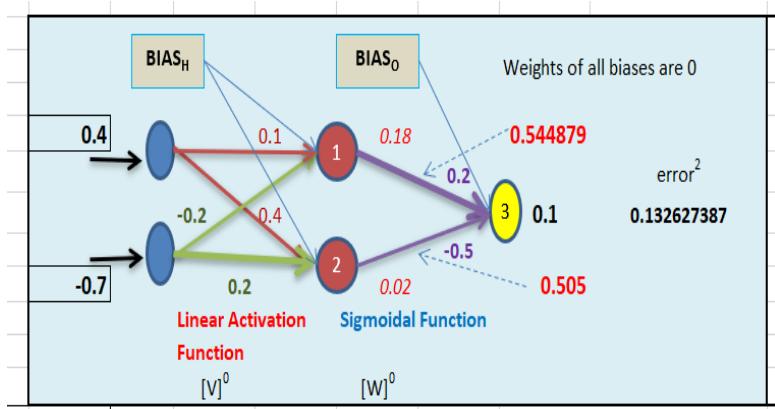
$$\begin{bmatrix} -0.00449241 \\ 0.0113211 \end{bmatrix}$$



Gradient Descent

Step 12 $[X] = \{O\}_{1 < d^*}$

$$\begin{bmatrix} 0.4 \\ -0.7 \end{bmatrix} < -0.004492406 \quad 0.011321 > \begin{bmatrix} -0.0018 \\ 0.004528 \\ 0.003145 \\ -0.00792 \end{bmatrix}$$



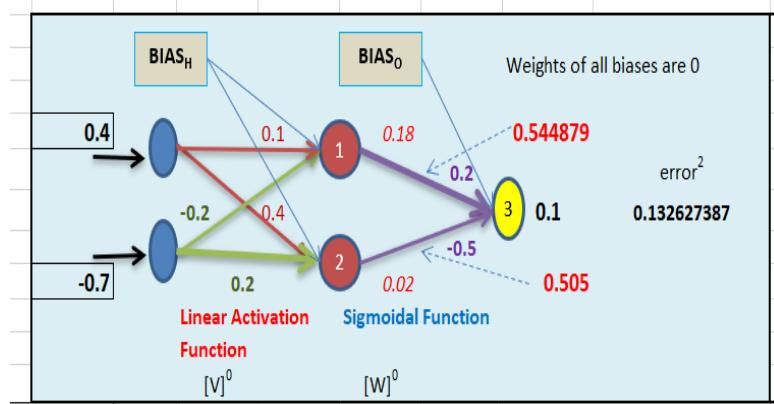
Gradient Descent

Step 13 Find $[\Delta V]^1$, as $= \alpha [\Delta V]^0 + \eta [X]$

Assume η as 0.6, AND α as zero.

$[\Delta V]^1$ is change in wts in Input Layer

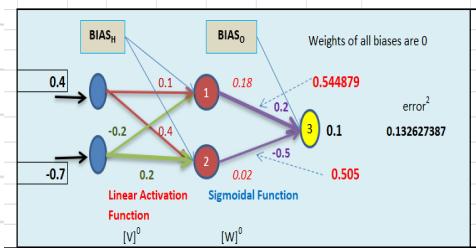
$$0.6 * \begin{bmatrix} -0.0018 & 0.004528 \\ 0.003145 & -0.00792 \end{bmatrix} = \begin{bmatrix} -0.00108 & 0.002717064 \\ 0.001887 & -0.004754862 \end{bmatrix}$$



Gradient Descent

Step 14 New wts in Input Layer
to Hidden Layer

$$[V]^1 = \text{Original wts} + \text{required/estimated change in wts in Input Layer}$$



$$\begin{bmatrix} 0.1 \\ -0.2 \end{bmatrix} + \begin{bmatrix} -0.001 \\ 0.002717 \\ 0.0019 \\ -0.00475 \end{bmatrix} = \begin{bmatrix} 0.0989 \\ -0.1981 \\ 0.402717064 \\ 0.195245138 \end{bmatrix}$$

$$[W]^1 = \text{Original wts} + \text{required/estimated change in wts [from hidden Layer to output layer].}$$

$$[W]^1 = \begin{bmatrix} -0.03 \\ -0.027 \end{bmatrix} + [W]^0 \begin{bmatrix} 0.2 \\ -0.5 \end{bmatrix} = \begin{bmatrix} 0.1704 \\ -0.5274 \end{bmatrix}$$

Lets try
classification

Classification

```
In [1]: # Jesus is my Saviour!  
  
In [2]: import os  
  
In [3]: os.chdir('C:\\\\Users\\\\Dr Vinod\\\\Desktop\\\\WD_python')  
  
In [4]: # our exported file will appear here  
  
In [5]: import pandas as pd  
  
In [6]: import numpy as np  
  
In [7]: import matplotlib.pyplot as plt  
  
In [8]: from sklearn.preprocessing import normalize
```



Import data

```
In [9]: data = pd.read_csv("C:/Users/Dr Vinod/Desktop/DataSets1/BreastCancer.csv")
```

```
In [10]: data = pd.DataFrame(data)
```

```
In [11]: data = data.dropna() # important step
```

```
In [12]: data.head(3)
```

```
Out[12]:
```

	Unnamed: 0	Id	C1.thickness	...	Normal.nucleoli	Mitoses	Class
0	1	1000025		5	...	1	1 benign
1	2	1002945		5	...	2	1 benign
2	3	1015425		3	...	1	1 benign

```
[3 rows x 12 columns]
```

```
In [13]: data.shape
```

```
Out[13]: (683, 12)
```



See data

```
In [14]: data.info()
<class 'pandas.core.frame.DataFrame'>
Int64Index: 683 entries, 0 to 698
Data columns (total 12 columns):
Unnamed: 0      683 non-null int64
Id              683 non-null int64
Cl.thickness    683 non-null int64
Cell.size       683 non-null int64
Cell.shape      683 non-null int64
Marg.adhesion   683 non-null int64
Epith.c.size    683 non-null int64
Bare.nuclei    683 non-null float64
Bl.cromatin     683 non-null int64
Normal.nucleoli 683 non-null int64
Mitoses         683 non-null int64
Class            683 non-null object
dtypes: float64(1), int64(10), object(1)
memory usage: 69.4+ KB
```



Create Predictors

```
In [15]: X = data.iloc[:, 2:10]
```

```
In [16]: X.shape  
Out[16]: (683, 8)
```

```
In [17]: X.info()  
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 683 entries, 0 to 698  
Data columns (total 8 columns):  
Cl.thickness      683 non-null int64  
Cell.size         683 non-null int64  
Cell.shape        683 non-null int64  
Marg.adhesion     683 non-null int64  
Epith.c.size      683 non-null int64  
Bare.nuclei       683 non-null float64  
Bl.cromatin       683 non-null int64  
Normal.nucleoli   683 non-null int64  
dtypes: float64(1), int64(7)  
memory usage: 48.0 KB
```



Create Response Variable

```
In [18]: y = data.iloc[:, 11]
```

```
In [19]: y.shape  
Out[19]: (683,)
```

```
In [20]: y.dtype # what is 'o', object  
Out[20]: dtype('O')
```



Train & Test Data and Scaling

```
In [21]: #_____ train and test data
```

```
In [22]: from sklearn.model_selection import train_test_split
```

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
In [24]: #_____ pre processing of data
```

```
In [25]: from sklearn.preprocessing import StandardScaler
```

```
In [26]: scaler = StandardScaler()
```

```
In [27]: # Fit only to the training data
```

```
In [28]: scaler.fit(X_train)
```

```
Out[28]: StandardScaler(copy=True, with_mean=True, with_std=True)
```



```
In [29]: # Now apply the transformations to the data:
```

```
In [30]: X_train = scaler.transform(X_train)
```

```
In [31]: X_train
```

```
Out[31]:
```

```
array([[ 1.56880671,  0.27189965,  0.58119997, ...,  1.72488735,
        0.22810623,  1.7032463 ],
       [ 0.16578656, -0.72484941, -0.76681536, ..., -0.72450584,
        -0.1890023 , -0.61606781],
       [-0.53572351, -0.72484941, -0.76681536, ..., -0.72450584,
        0.64521476,  1.7032463 ],
       ...,
       [ 1.91956174,  1.93314807,  1.25520764, ..., -0.45235104,
        1.47943182,  1.37191571],
       [-1.23723359, -0.72484941, -0.76681536, ..., -0.72450584,
        -1.02321937, -0.61606781],
       [ 0.16578656, -0.72484941, -0.42981153, ..., -0.72450584,
        -0.1890023 , -0.61606781]])
```

```
In [32]: X_train.shape # Look at Variable explorer!
```

```
Out[32]: (512, 8)
```



```
In [33]: X_test = scaler.transform(X_test)
```

```
In [34]: X_test
```

Out[34]:

```
array([[-1.23723359, -0.72484941, -0.76681536, ..., -0.72450584,
       -1.02321937, -0.61606781],
      [ 0.5165416 , -0.39259973, -0.0928077 , ..., -0.72450584,
       -1.02321937, -0.61606781],
      [ 0.86729663,  1.60089839, -0.0928077 , ...,  0.36411335,
       1.47943182,  1.7032463 ],
      ...,
      [-0.88647855, -0.72484941, -0.0928077 , ..., -0.72450584,
       -0.60611083, -0.61606781],
      [-0.18496848,  0.93639902,  0.91820381, ...,  0.63626815,
       1.47943182,  1.37191571],
      [-0.18496848,  2.26539776,  0.24419614, ...,  1.72488735,
       2.31364888,  2.36590747]])
```

```
In [35]: X_test.shape
```

Out[35]: (171, 8)



```
In [36]: from sklearn.neural_network import MLPClassifier
```

```
In [37]: mlp = MLPClassifier(hidden_layer_sizes=(30,30,30))
```

```
In [38]: mlp.fit(X_train,y_train)
```

```
C:\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:566:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the  
optimization hasn't converged yet.
```

```
    % self.max_iter, ConvergenceWarning)
```

```
Out[38]:
```

```
MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,  
              beta_2=0.999, early_stopping=False, epsilon=1e-08,  
              hidden_layer_sizes=(30, 30, 30), learning_rate='constant',  
              learning_rate_init=0.001, max_iter=200, momentum=0.9,  
              n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,  
              random_state=None, shuffle=True, solver='adam', tol=0.0001,  
              validation_fraction=0.1, verbose=False, warm_start=False)
```



Predict

```
In [39]: # _____ Let's predict our model
```

```
In [40]: predictions = mlp.predict(X_test)
```

```
In [41]: predictions # too large output!
```

```
Out[41]:
```

```
array(['benign', 'benign', 'malignant', 'benign', 'benign', 'benign',
       'benign', 'benign', 'benign', 'benign', 'malignant', 'benign',
       'benign', 'benign', 'benign', 'malignant', 'malignant', 'benign',
       'benign', 'benign', 'benign', 'malignant', 'benign', 'malignant',
       'malignant', 'malignant', 'benign', 'benign', 'benign', 'benign',
       'benign', 'benign', 'benign', 'benign', 'benign', 'benign',
       'benign', 'malignant', 'malignant', 'malignant', 'malignant',
       'malignant', 'benign', 'malignant', 'malignant', 'benign', 'benign'],
      dtype='|S10')
```



Results

```
In [42]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [43]: print(confusion_matrix(y_test,predictions))
```

```
[[115  2]
 [ 3  51]]
```

```
In [44]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
benign	0.97	0.98	0.98	117
malignant	0.96	0.94	0.95	54
accuracy			0.97	171
macro avg	0.97	0.96	0.97	171
weighted avg	0.97	0.97	0.97	171



```
In [45]: # Curve

In [46]: from sklearn import metrics

In [47]: from sklearn.metrics import roc_curve

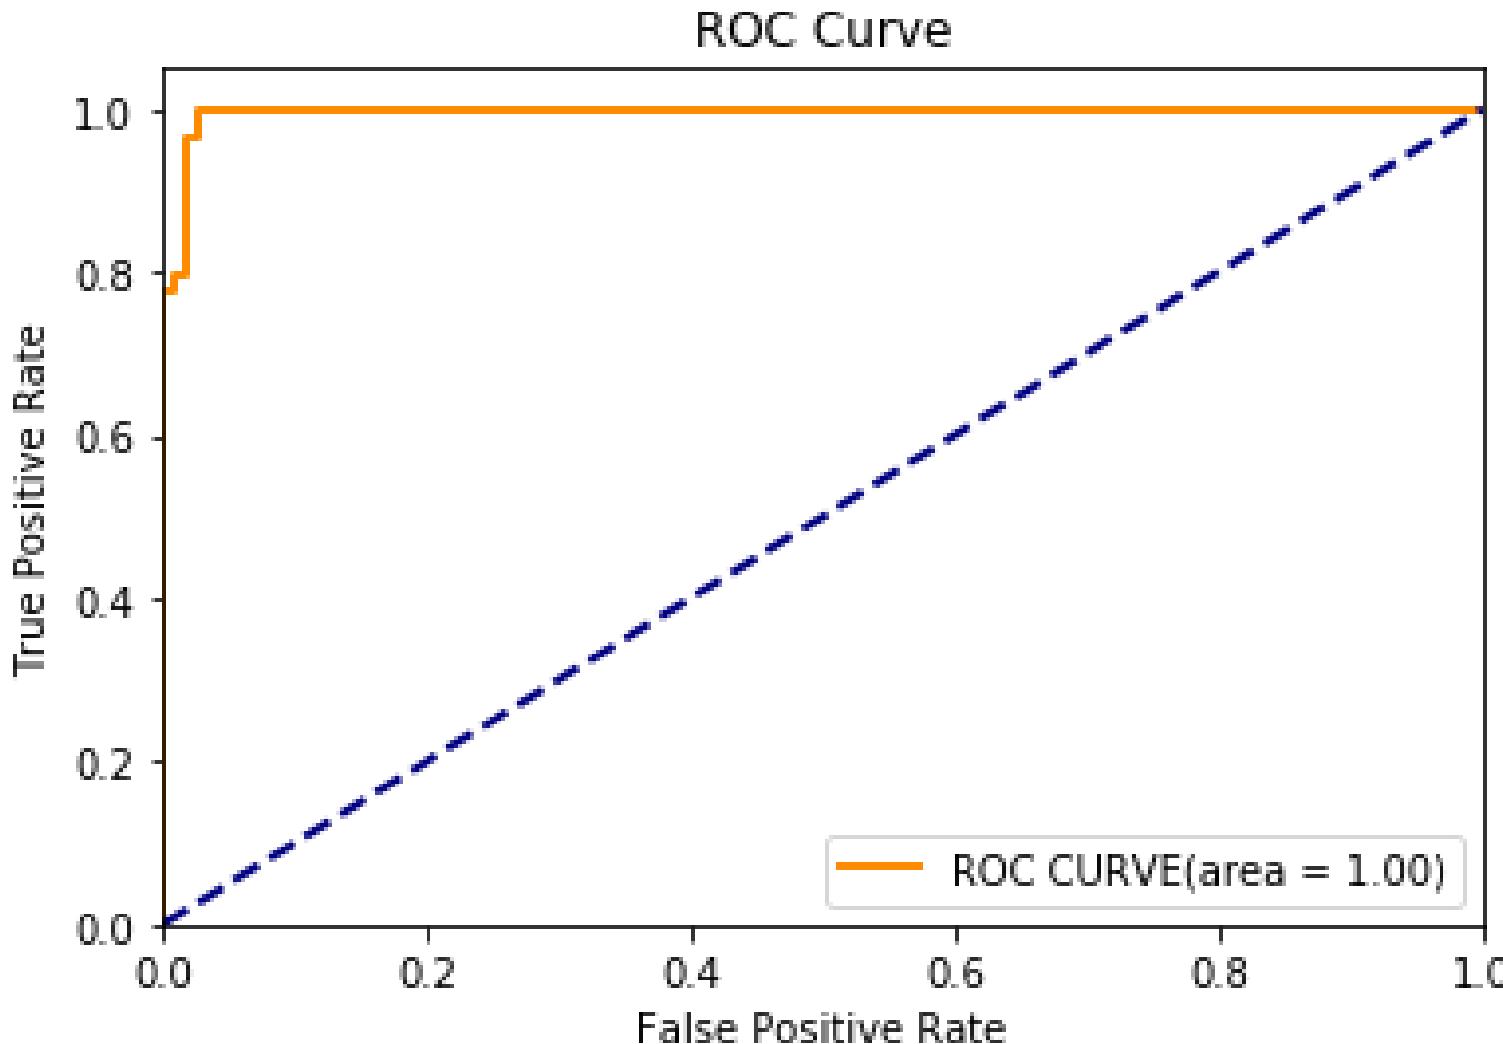
In [48]: y_pred_prob = mlp.predict_proba(X_test)[:, 1]

In [49]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob, pos_label =
'malignant')

In [50]: auc = metrics.roc_auc_score(y_test, y_pred_prob)

In [51]: plt.figure()
....: lw = 2
....: plt.plot(fpr, tpr, color = 'darkorange', lw = lw, label = 'ROC CURVE(area
= %0.2f)' % auc)
....: plt.plot([0,1], [0,1], color = 'navy', lw = lw, linestyle = '--')
....: plt.xlim([0.0, 1.0])
....: plt.ylim([0.0, 1.05])
....: plt.xlabel('False Positive Rate')
....: plt.ylabel('True Positive Rate')
....: plt.title('ROC Curve')
....: plt.legend(loc = 'lower right')
....: plt.show()
....: ####
```







Lets try
Prediction

```
In [1]: # Jesus is my Saviour!  
  
In [2]: import os  
  
In [3]: os.chdir('C:\\\\Users\\\\Dr Vinod\\\\Desktop\\\\WD_python')  
  
In [4]: ## our exported file will appear here  
  
In [5]: import pandas as pd  
  
In [6]: import numpy as np  
  
In [7]: import matplotlib.pyplot as plt  
  
In [8]: import seaborn as sns  
  
In [9]: from sklearn.preprocessing import normalize
```



```
In [10]: # _____ import data
```

```
In [11]: data = pd.read_csv("C:/Users/Dr Vinod/Desktop/DataSets1/insurance.csv")
```

```
In [12]: data = pd.DataFrame(data)
```

```
In [13]: data = data.dropna() # important step
```

```
In [14]: data['sex'].head()
```

```
Out[14]:
```

```
0    female  
1     male  
2     male  
3     male  
4     male  
Name: sex, dtype: object
```

```
In [15]: data['sex'].dtype # object
```

```
Out[15]: dtype('O')
```

```
In [16]: data['sex'].to_frame()
```

```
Out[16]:
```

	sex
0	female
1	male
2	male
3	male
4	male



Convert words into numbers

```
In [17]: data['sex']=data.get('sex').replace('male',1)

In [18]: data['sex']=data.get('sex').replace('female',2)

In [19]: c_sex = data['sex'].value_counts()

In [20]: c_sex
Out[20]:
1    676
2    662
Name: sex, dtype: int64
```



Convert words into numbers

```
In [21]: data['smoker']=data.get('smoker').replace('yes',1)

In [22]: data['smoker']=data.get('smoker').replace('no',2)

In [23]: c_smoker = data['smoker'].value_counts()

In [24]: c_smoker
Out[24]:
2    1064
1    274
Name: smoker, dtype: int64
```



```
In [25]: data.head(3)
```

```
Out[25]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.90	0	1	southwest	16884.9240
1	18	1	33.77	1	2	southeast	1725.5523
2	28	1	33.00	3	2	southeast	4449.4620

```
In [26]: data.shape
```

```
Out[26]: (1338, 7)
```

```
In [27]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1338 entries, 0 to 1337
Data columns (total 7 columns):
age            1338 non-null int64
sex            1338 non-null int64
bmi            1338 non-null float64
children       1338 non-null int64
smoker         1338 non-null int64
region          1338 non-null object
charges        1338 non-null float64
dtypes: float64(2), int64(4), object(1)
memory usage: 83.6+ KB
```



```
In [28]: # we dont think that region can be a predictor, so drop this
```

```
In [29]: del data['region']
```

```
In [30]: data.head(3)
```

```
Out[30]:
```

```
   age  sex    bmi  children  smoker      charges
0   19    2  27.90        0       1  16884.9240
1   18    1  33.77        1       2  1725.5523
2   28    1  33.00        3       2  4449.4620
```

```
In [31]: data.shape
```

```
Out[31]: (1338, 6)
```

```
In [32]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 1338 entries, 0 to 1337
```

```
Data columns (total 6 columns):
```

```
age          1338 non-null int64
```

```
sex          1338 non-null int64
```

```
bmi          1338 non-null float64
```

```
children     1338 non-null int64
```

```
smoker       1338 non-null int64
```

```
charges      1338 non-null float64
```

```
dtypes: float64(2), int64(4)
```

```
memory usage: 73.2 KB
```



```
In [33]: # _____ lets create Predictors and RV
```

```
In [34]: X = data.iloc[:, 0:5]
```

```
In [35]: X.shape
```

```
Out[35]: (1338, 5)
```

```
In [36]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1338 entries, 0 to 1337
Data columns (total 5 columns):
age            1338 non-null int64
sex            1338 non-null int64
bmi            1338 non-null float64
children       1338 non-null int64
smoker         1338 non-null int64
dtypes: float64(1), int64(4)
memory usage: 62.7 KB
```

```
In [37]: y = data.iloc[:, 5]
```

```
In [38]: y.shape
```

```
Out[38]: (1338, )
```

```
In [39]: y.dtype # float64
```

```
Out[39]: dtype('float64')
```



```
In [40]: #_____train and test data  
  
In [41]: from sklearn.model_selection import train_test_split  
  
In [42]: X_train, X_test, y_train, y_test = train_test_split(X, y)  
  
In [43]: #_____pre processing of data  
  
In [44]: from sklearn.preprocessing import StandardScaler  
  
In [45]: scaler = StandardScaler()
```



```
In [48]: # Now apply the transformations to the data:
```

```
In [49]: X_train = scaler.transform(X_train)
```

```
In [50]: X_train
```

```
Out[50]:
```

```
array([[-0.09586918, -0.98909239,  1.24993093,  1.61515654, -1.98516667],
       [ 1.18243375, -0.98909239,  1.55984391, -0.90037453,  0.50373604],
       [ 1.5375179 ,  1.0110279 ,  2.16019888, -0.90037453,  0.50373604],
       ...,
       [-0.16688601,  1.0110279 , -0.6955437 , -0.90037453, -1.98516667],
       [ 1.46650107, -0.98909239, -0.80263404, -0.90037453,  0.50373604],
       [-1.37417211,  1.0110279 ,  0.12548229, -0.90037453,  0.50373604]])
```

```
In [51]: X_train.shape # Look at Variable explorer!
```

```
Out[51]: (1003, 5)
```



```
In [52]: X_test = scaler.transform(X_test)

In [53]: X_test
Out[53]:
array([[ 1.18243375e+00,  1.01102790e+00, -3.85630724e-01,
       -9.00374532e-01,  5.03736042e-01],
       [ 3.30231797e-01, -9.89092386e-01,  5.43556550e-04,
       7.76646184e-01,  5.03736042e-01],
       [-1.01908796e+00, -9.89092386e-01, -8.02634043e-01,
       -9.00374532e-01,  5.03736042e-01],
       ...,
       [-2.48523501e-02, -9.89092386e-01, -3.87253305e-01,
       -6.18641740e-02, -1.98516667e+00],
       [ 1.11141692e+00, -9.89092386e-01,  1.14040672e+00,
       1.61515654e+00,  5.03736042e-01],
       [-6.64003815e-01,  1.01102790e+00,  2.78004909e-01,
       -6.18641740e-02,  5.03736042e-01]])
```

```
In [54]: X_test.shape
Out[54]: (335, 5)
```



```
In [55]: from sklearn.neural_network import MLPRegressor  
  
In [56]: mlp = MLPRegressor(hidden_layer_sizes=(30,30,30))  
  
In [57]: mlp  
Out[57]:  
MLPRegressor(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,  
             beta_2=0.999, early_stopping=False, epsilon=1e-08,  
             hidden_layer_sizes=(30, 30, 30), learning_rate='constant',  
             learning_rate_init=0.001, max_iter=200, momentum=0.9,  
             n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,  
             random_state=None, shuffle=True, solver='adam', tol=0.0001,  
             validation_fraction=0.1, verbose=False, warm_start=False)
```



```
In [58]: mlp.fit(X_train,y_train)
C:\Anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:566:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the
optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)
Out[58]:
MLPRegressor(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
             beta_2=0.999, early_stopping=False, epsilon=1e-08,
             hidden_layer_sizes=(30, 30, 30), learning_rate='constant',
             learning_rate_init=0.001, max_iter=200, momentum=0.9,
             n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
             random_state=None, shuffle=True, solver='adam', tol=0.0001,
             validation_fraction=0.1, verbose=False, warm_start=False)
```



```
In [59]: #_____Let's predict our model  
  
In [60]: predictions = mlp.predict(X_test)  
  
In [61]: error = y_test - predictions  
  
In [62]: error  
Out[62]:  
837      118.166783  
184     -792.811360  
192     -1565.733495  
1184    -8087.493868  
868      1424.136765  
46     -4242.032138
```



```
In [63]: len(error) # 335  
Out[63]: 335
```

```
In [64]: squared_errors = error*error

In [65]: squared_errors
Out[65]:
837    1.396339e+04
184    6.285499e+05
192    2.451521e+06
1184   6.540756e+07
868    2.028166e+06
```



```
In [66]: sum_squared_errors = sum(squared_errors)

In [67]: sum_squared_errors_by_n = sum_squared_errors/335

In [68]: import math

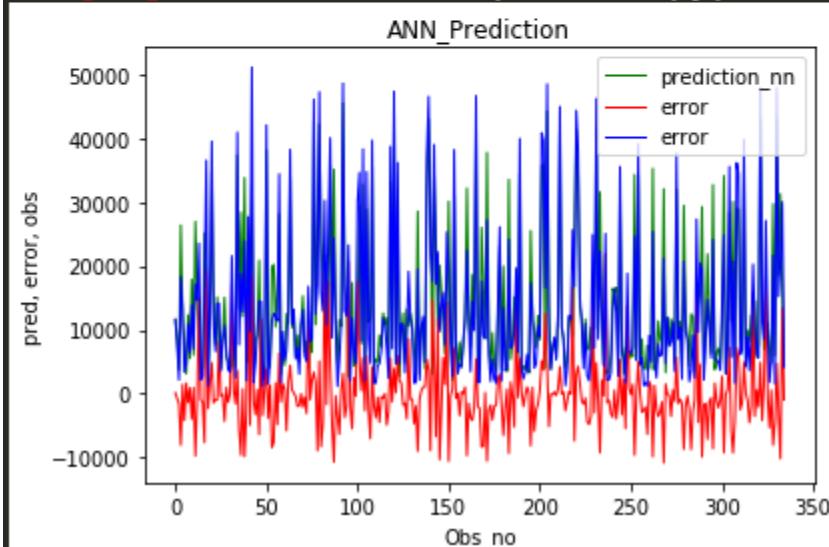
In [69]: RMSE = math.sqrt(sum_squared_errors_by_n)

In [70]: RMSE
Out[70]: 6046.392664728476
```

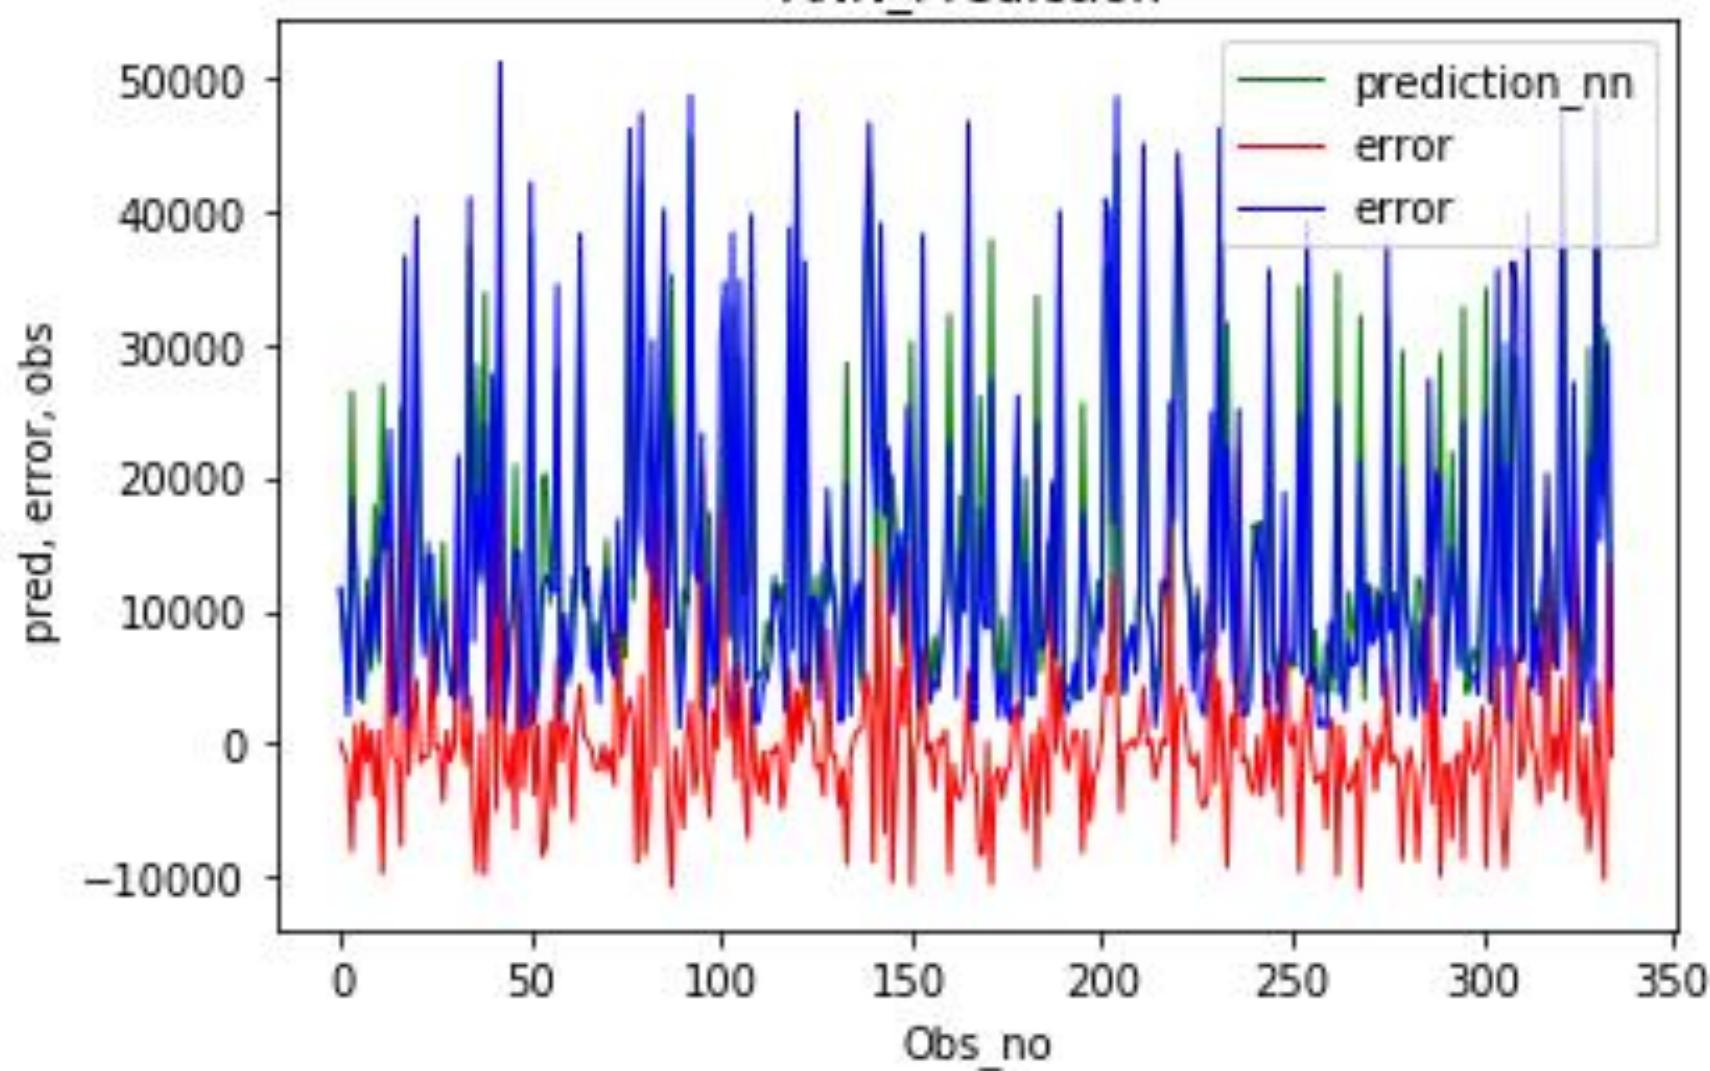
```
In [71]: x = np.arange(0, 335, 1)
```

```
In [72]: x
.... plt.plot(x, predictions, 'g', label = 'prediction_nn', linewidth = 1)
.... plt.plot(x, error, 'r', label = 'error', linewidth = 1)
.... plt.plot(x, y_test, 'b', label = 'error', linewidth = 1)
.... plt.title('ANN_Prediction')
.... plt.xlabel("Obs_no")
.... plt.ylabel("pred, error, obs")
.... plt.legend()
.... plt.show
```

```
Out[72]: <function matplotlib.pyplot.show(*args, **kw)>
```

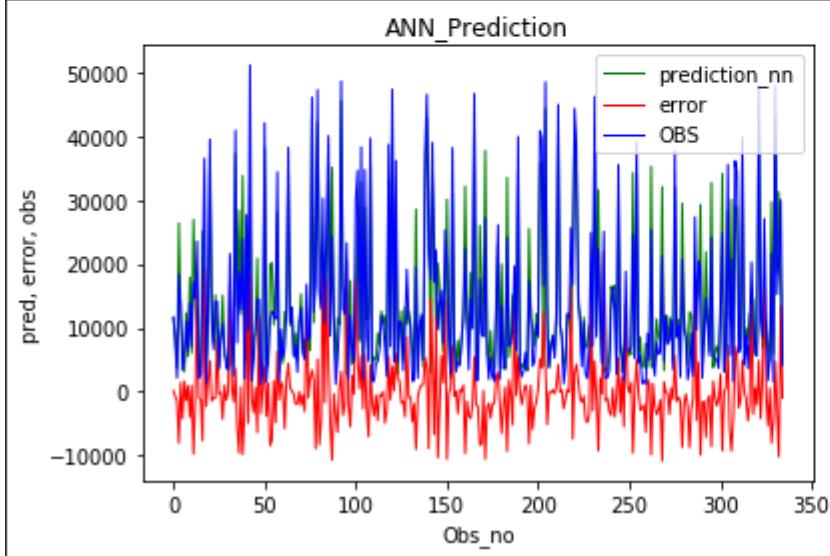


ANN_Prediction

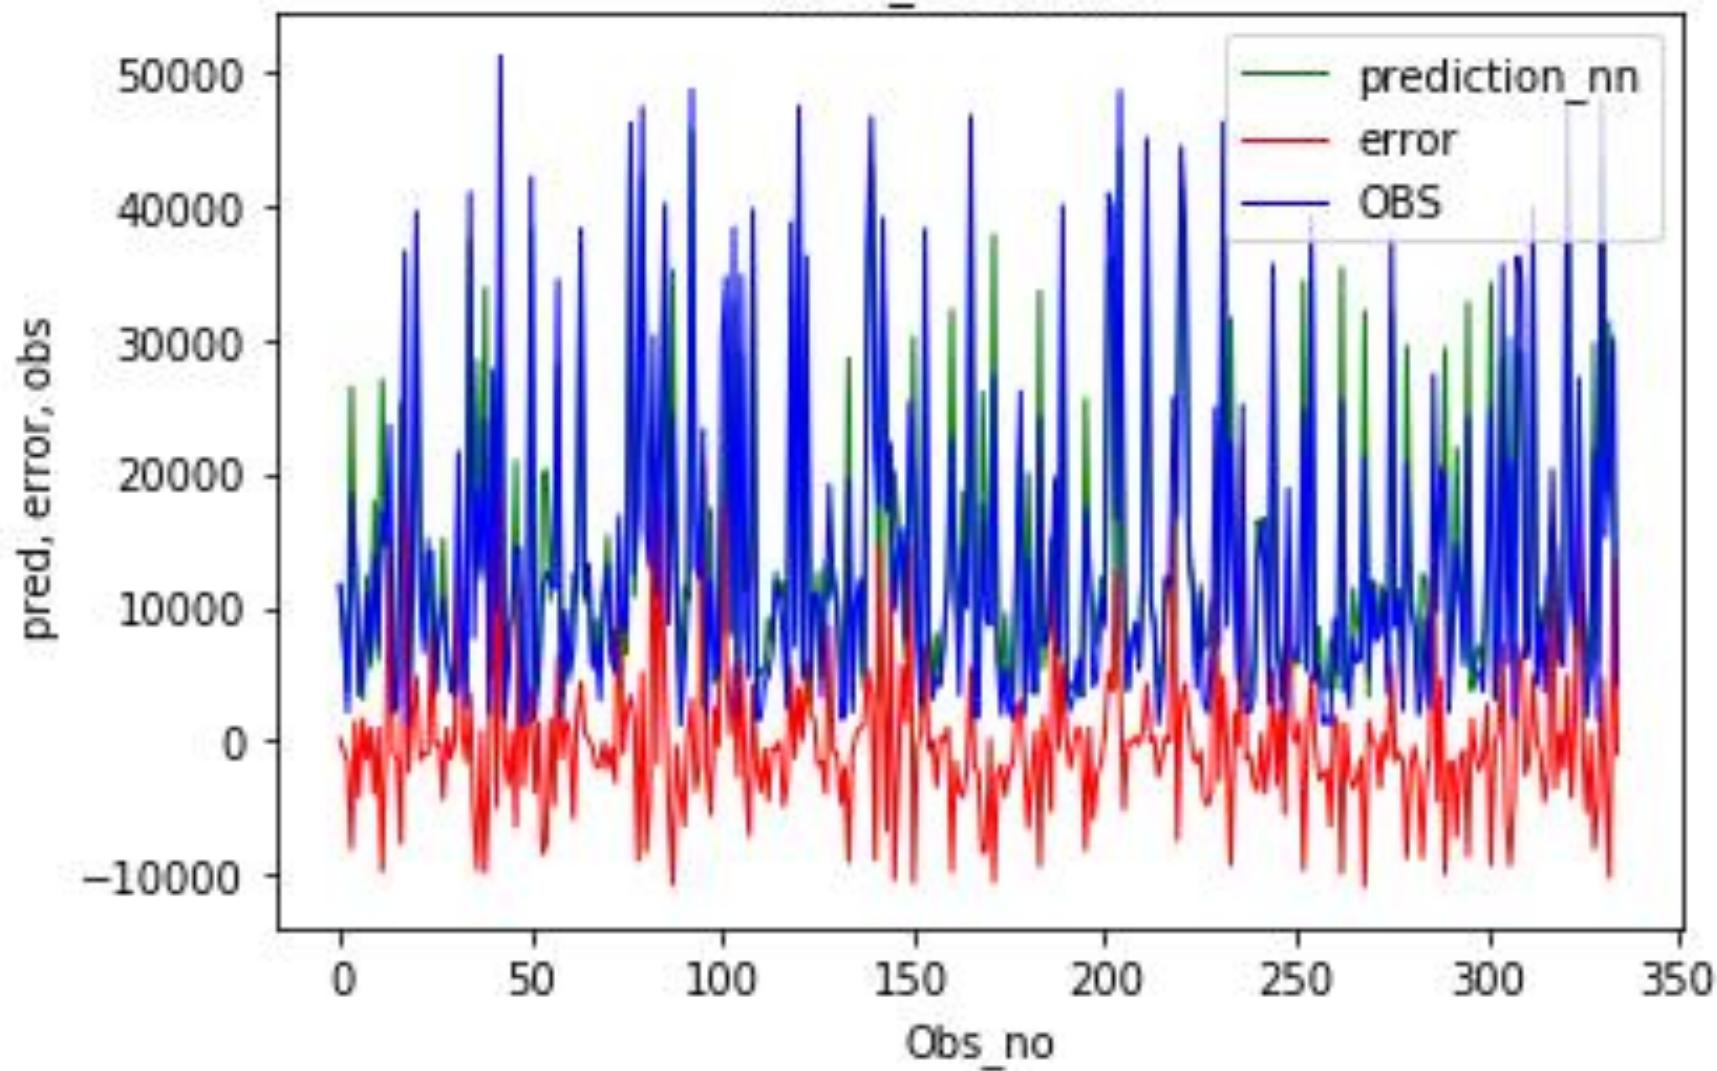


```
In [73]: x
.... plt.plot(x, predictions, 'g', label = 'prediction_nn', linewidth = 1)
.... plt.plot(x, error, 'r', label = 'error', linewidth = 1)
.... plt.plot(x, y_test, 'b', label = 'OBS', linewidth = 1)
.... plt.title('ANN_Prediction')
.... plt.xlabel("Obs_no")
.... plt.ylabel("pred, error, obs")
.... plt.legend()
.... plt.show
```

Out[73]: <function matplotlib.pyplot.show(*args, **kw)>



ANN_Prediction



Lets compare with Linear Regression

```
In [74]: #_____Linear Regression

In [75]: from sklearn.linear_model import LinearRegression

In [76]: from sklearn import metrics

In [77]: regressor = LinearRegression()

In [78]: lin_reg = regressor.fit(X_train, y_train) #training the algorithm

In [79]: predictions1 = lin_reg.predict(X_test)

In [80]: error1 = y_test - predictions1

In [81]: error1
Out[81]:
837      327.820881
184     -2237.620100
192     -342.472798
```

```
In [82]: len(error1) # 335
Out[82]: 335

In [83]: squared_errors1 = error1*error1

In [84]: squared_errors1
Out[84]:
837    1.074665e+05
184    5.006944e+06
192    1.172876e+05
1184   8.092924e+07
```



```
In [85]: sum_squared_errors1 = sum(squared_errors1)

In [86]: sum_squared_errors_by_n1 = sum_squared_errors1/335

In [87]: import math

In [88]: RMSE1 = math.sqrt(sum_squared_errors_by_n1)

In [89]: RMSE1
Out[89]: 6325.716772924798
```

